

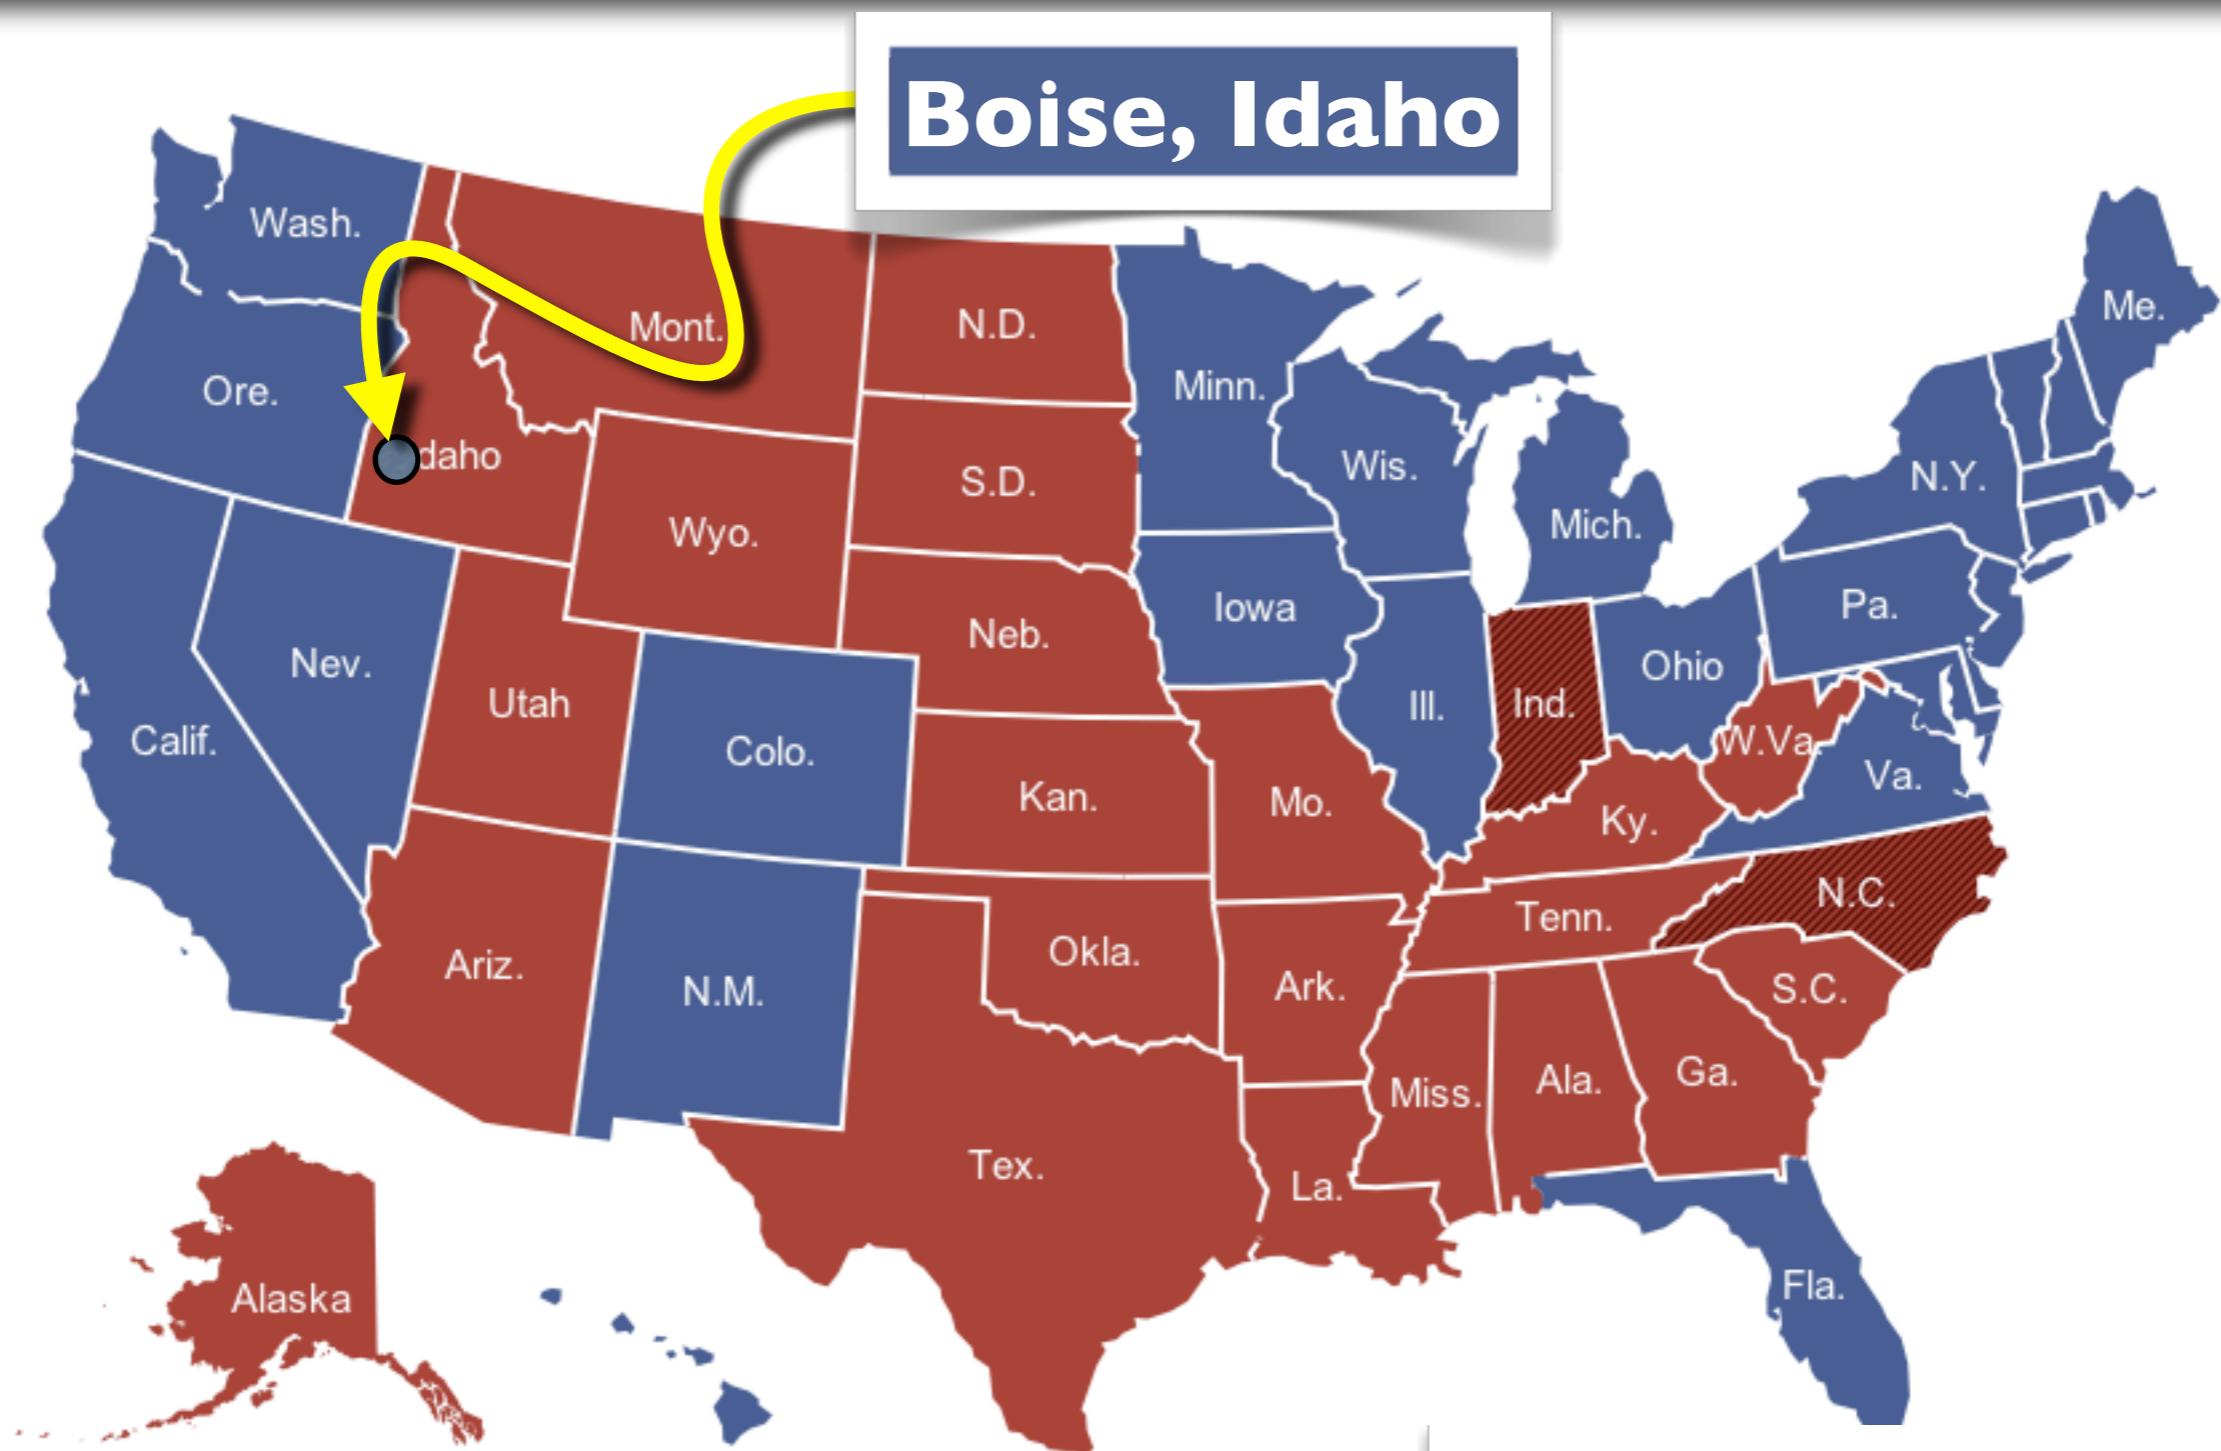
# A hybrid adaptive mesh framework for finite volume schemes on a forest of locally refined Cartesian meshes

**Donna Calhoun** (Boise State University)

Carsten Burstedde (University of Bonn, Germany)

SIAM Geosciences  
Padua, Italy  
June 17-20, 2013

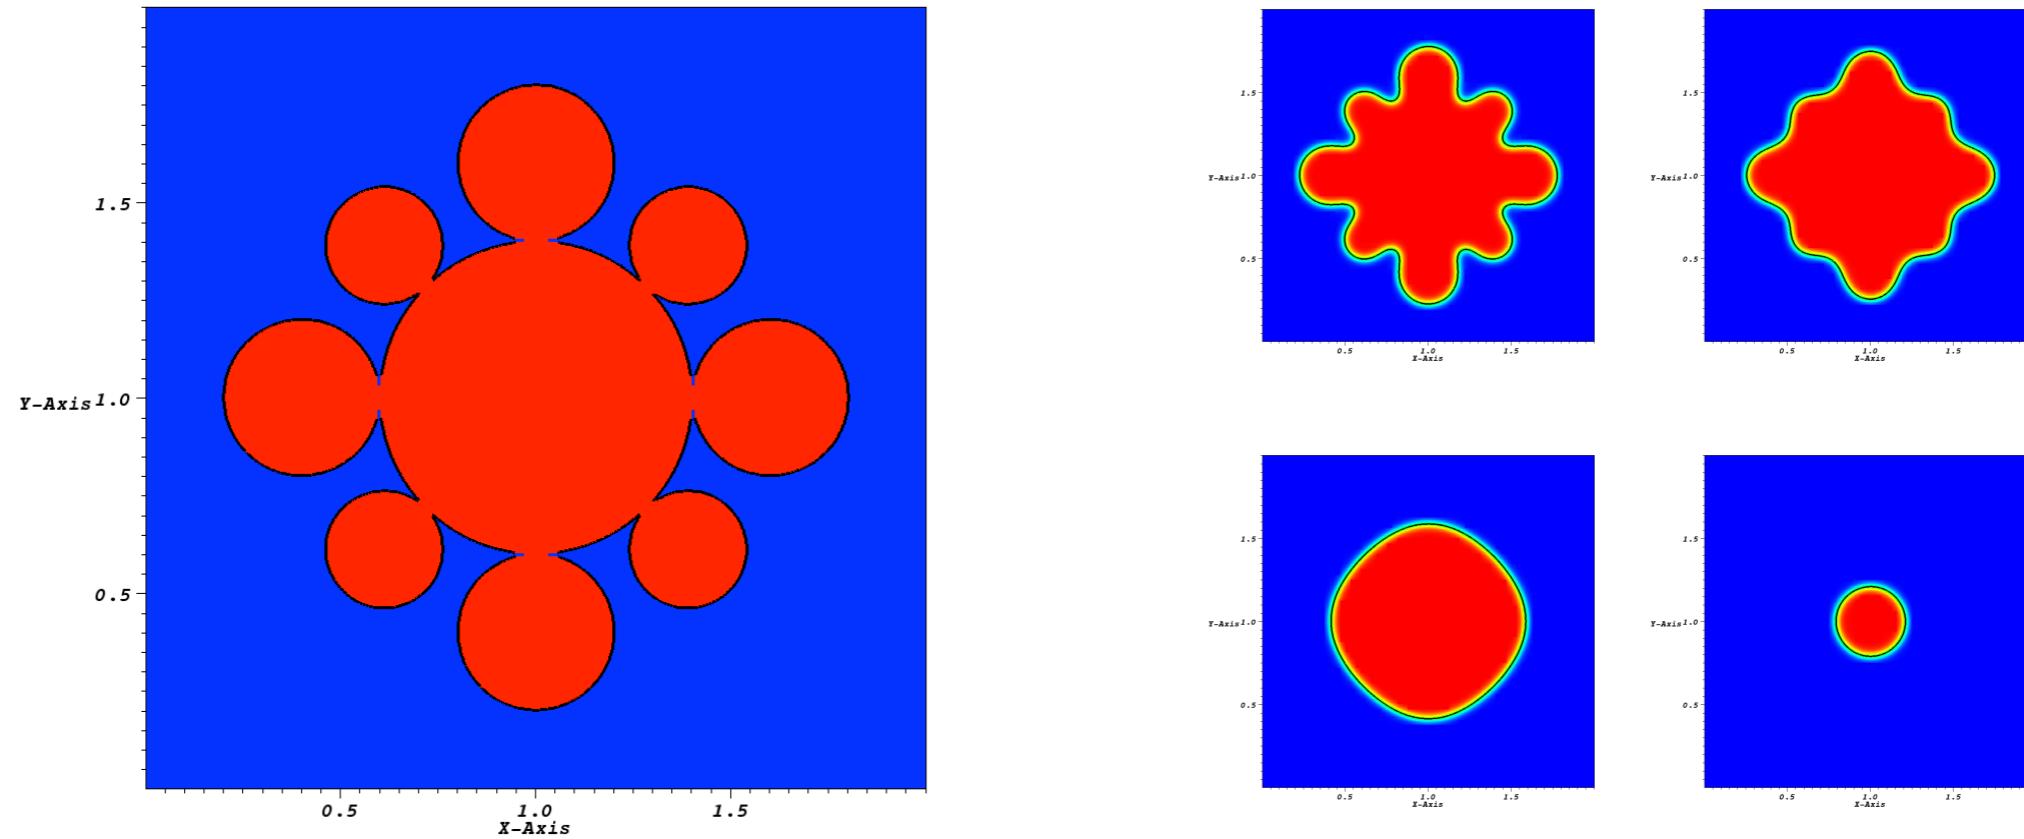
# Where is Boise?



*Boise comes from boisé which is French for ‘wooded’ or ‘forested’.*

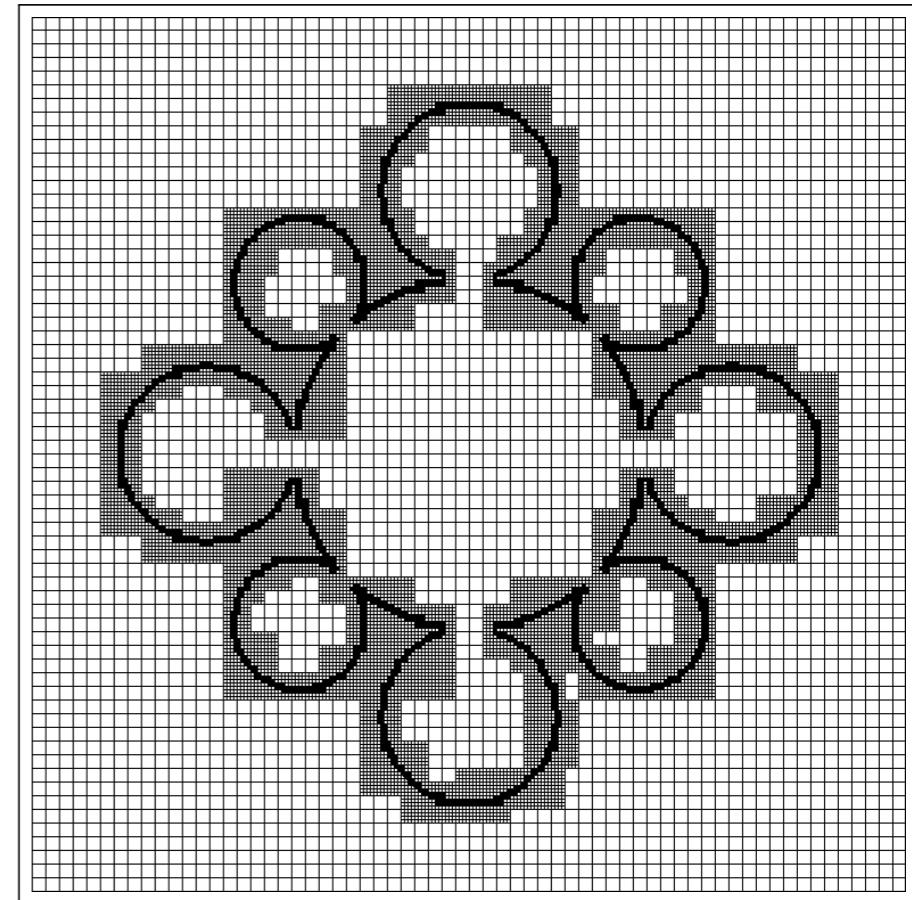
# Adaptive Mesh Refinement

When solving PDEs using mesh based methods, it is generally recognized that many problems could benefit enormously from a multi-resolution grid, or spatial adaptivity.

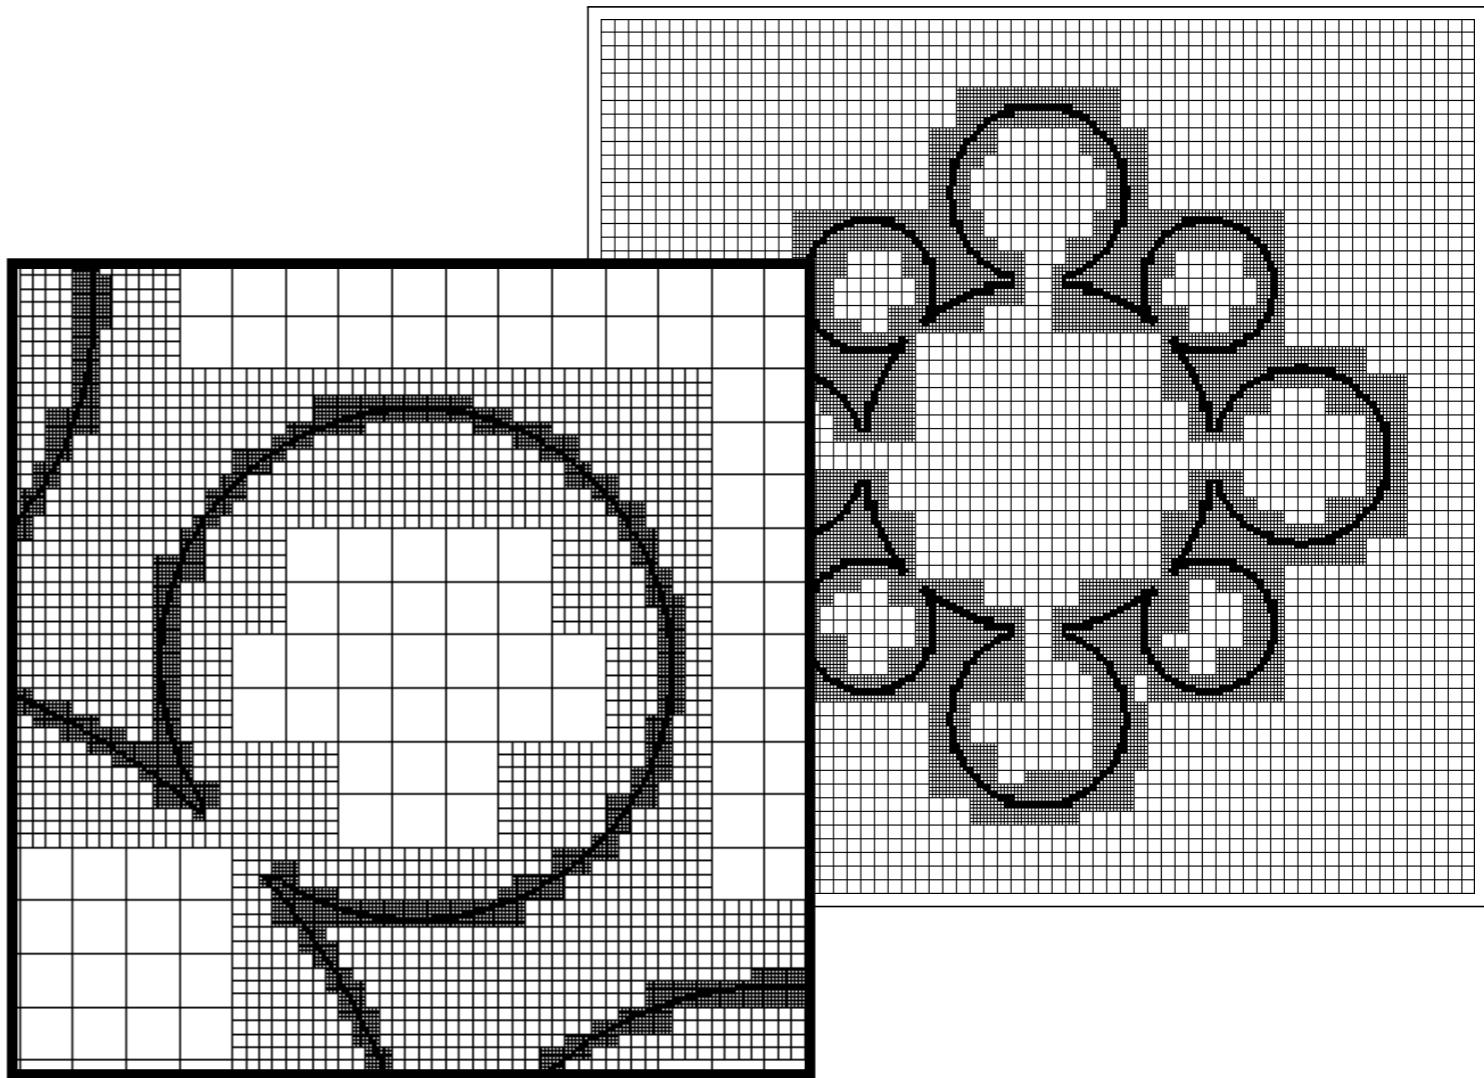


Allen Cahn equation - Flow by mean curvature

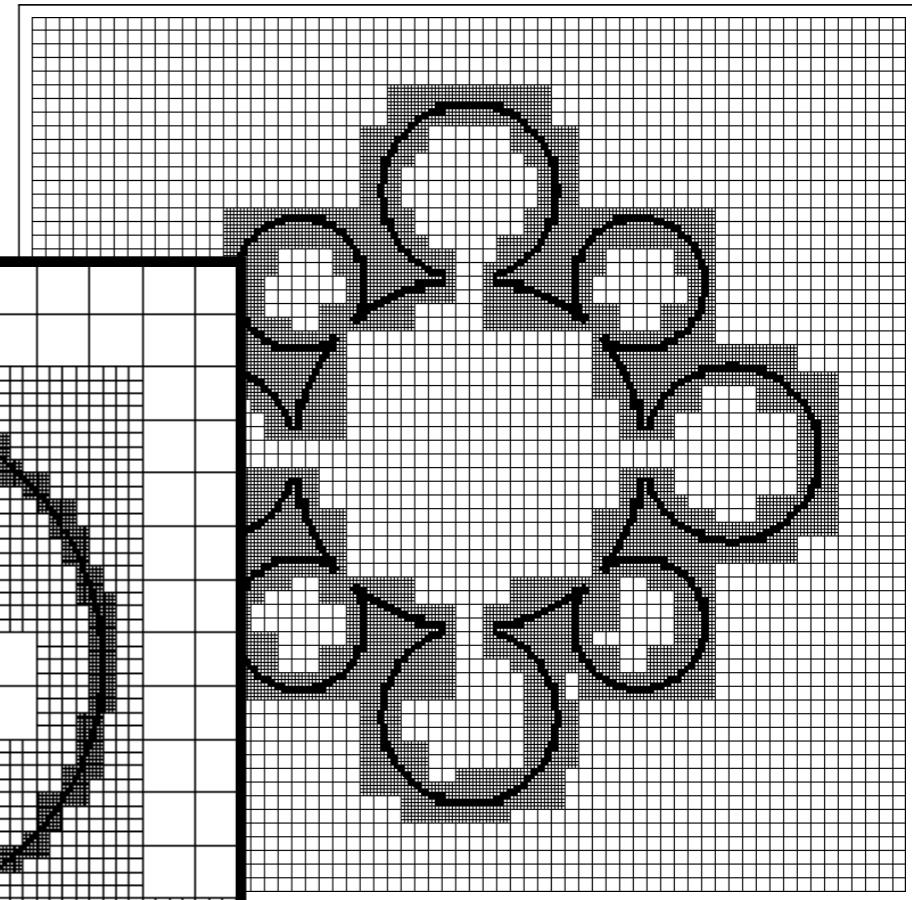
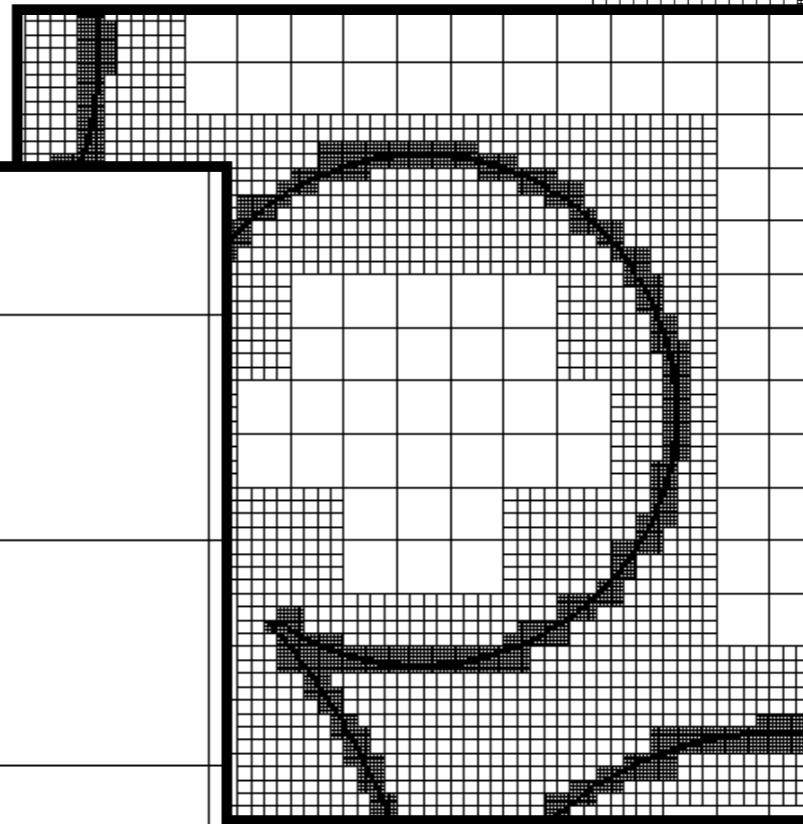
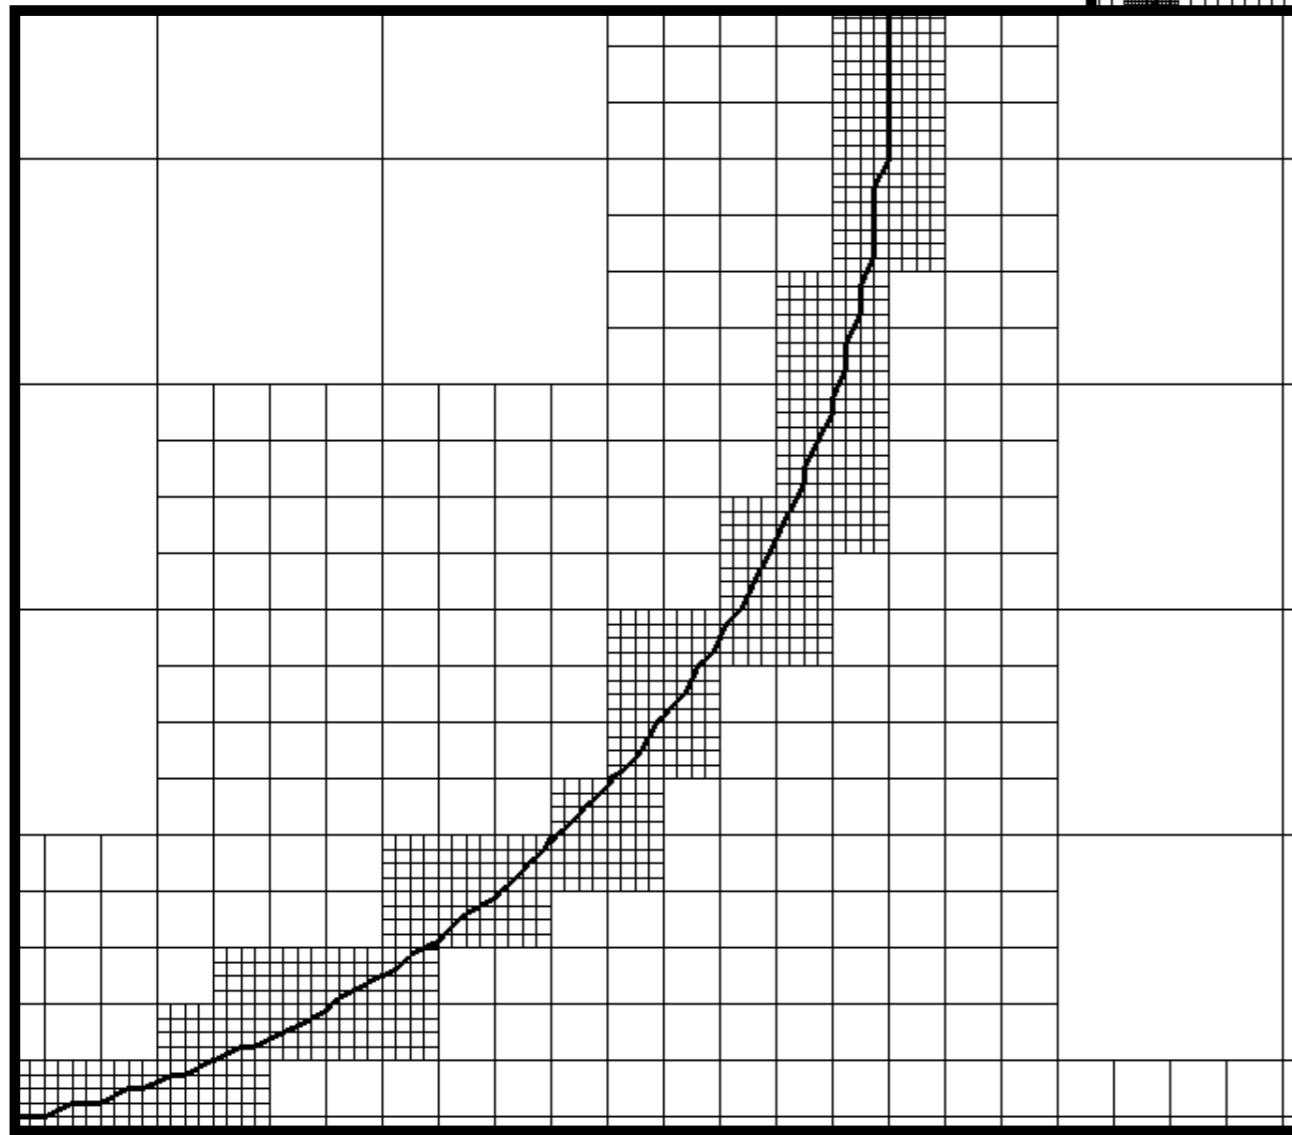
# Example : Flow by mean curvature



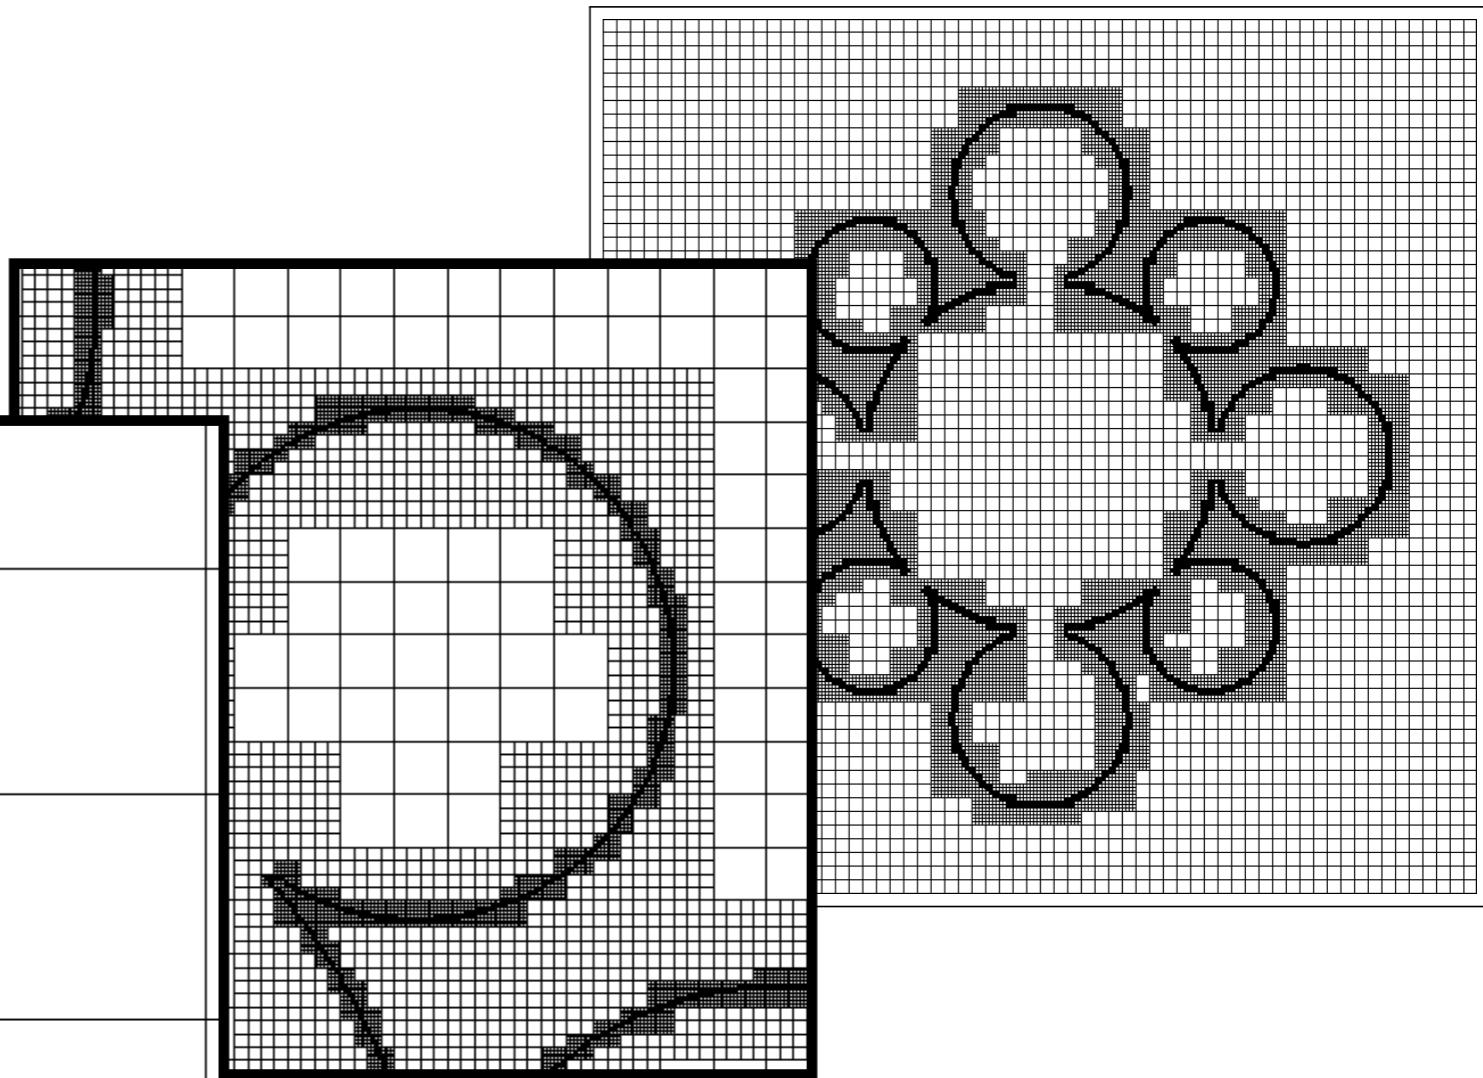
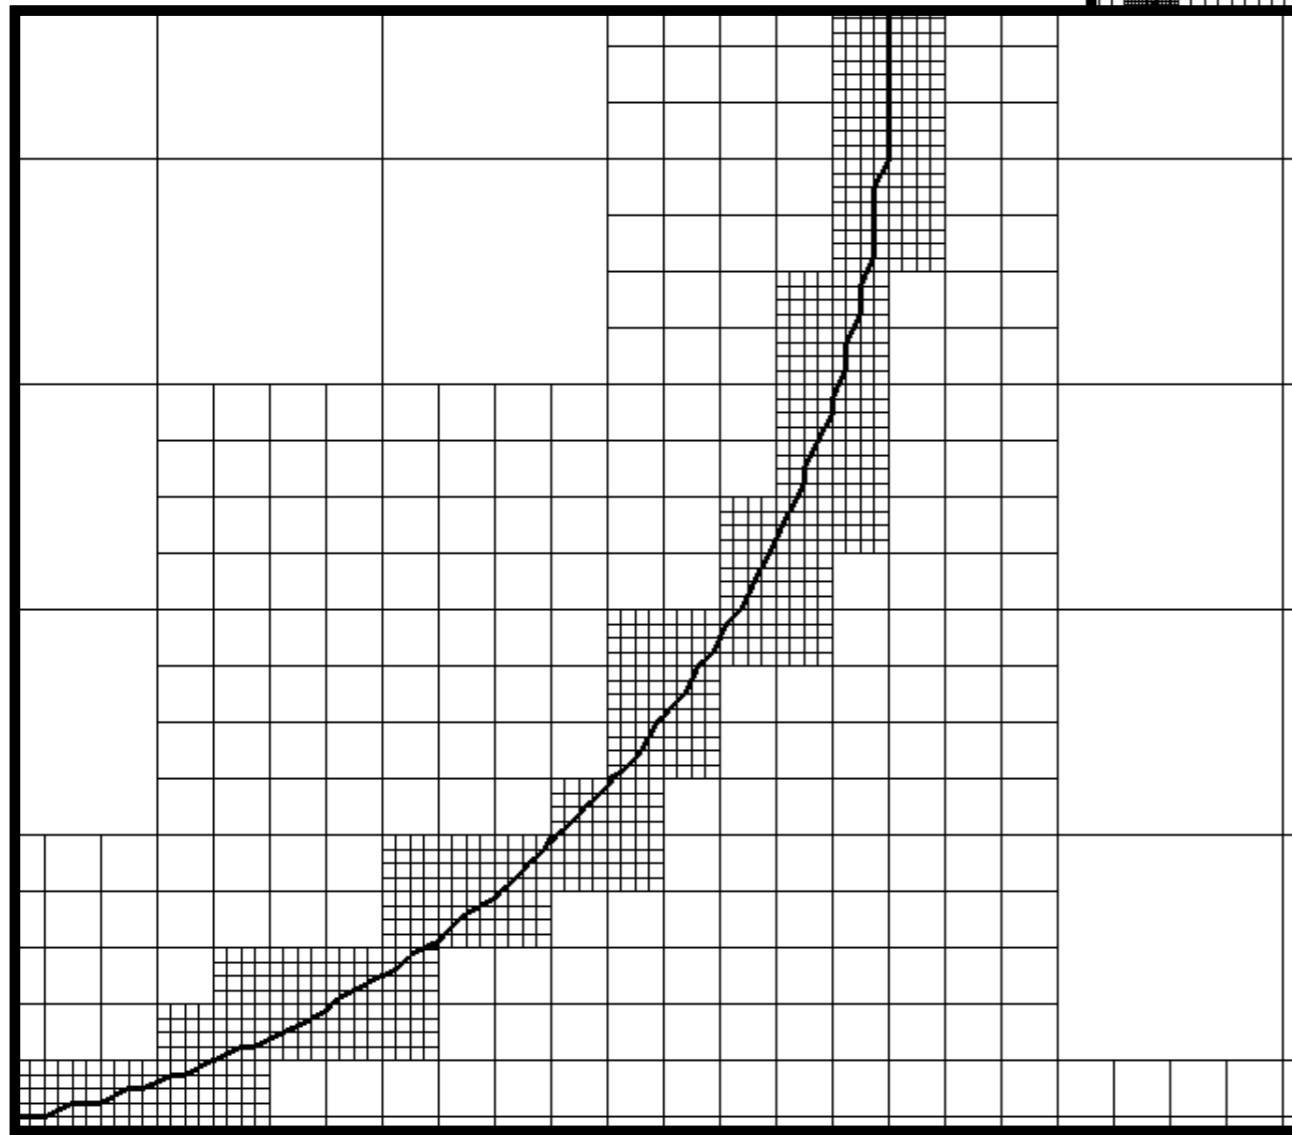
# Example : Flow by mean curvature



# Example : Flow by mean curvature

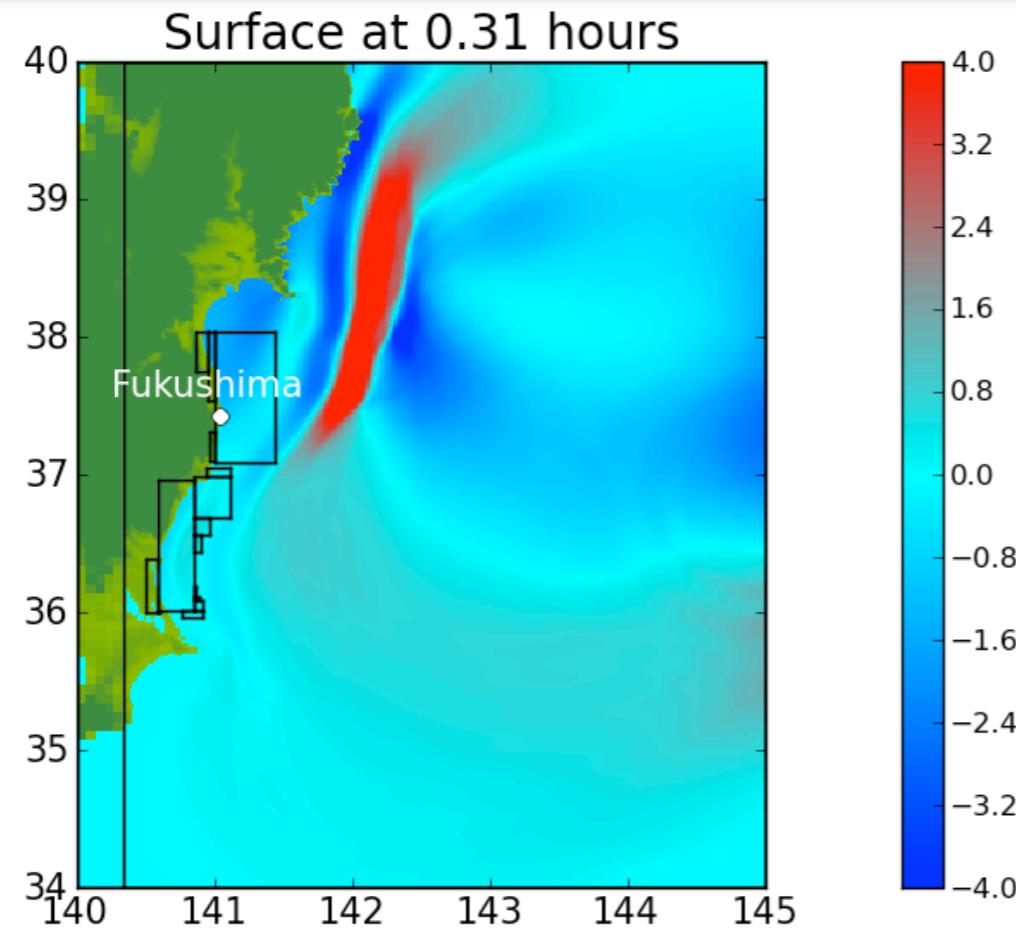


# Example : Flow by mean curvature

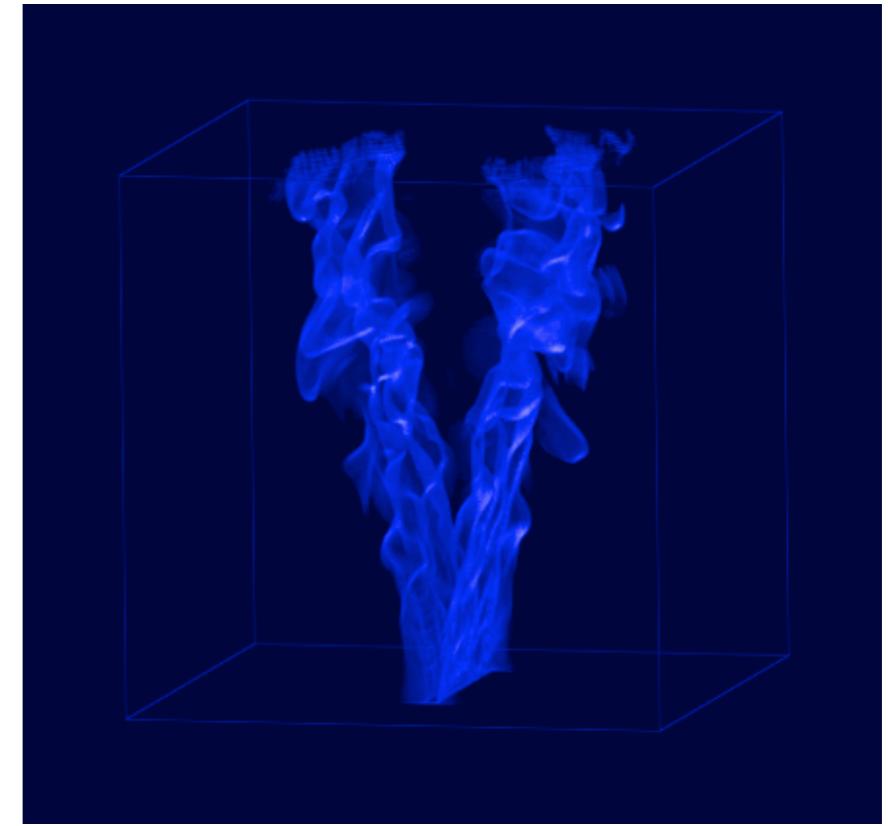


*Put more grid resolution where needed*

# Applications for AMR



Tsunami modeling (R. LeVeque, D. George, M. Berger)



Rod stabilized V-flame (J. B. Bell,  
Lawrence Berkeley Lab)

- Tracer transport in the atmosphere
- Astrophysics
- Shock capturing for aerodynamic applications
- Regional weather forecasting, hurricanes, ...

# Isaac Newton Institute, Fall 2012

“Multi-scale numerics for the ocean and atmosphere”

# Isaac Newton Institute, Fall 2012

“Multi-scale numerics for the ocean and atmosphere”

What I learned from the atmospheric science community :

# Isaac Newton Institute, Fall 2012

“Multi-scale numerics for the ocean and atmosphere”

What I learned from the atmospheric science community :

- There remains skepticism about how effective adaptive mesh refinement (AMR) can be in weather and climate modeling

# Isaac Newton Institute, Fall 2012

“Multi-scale numerics for the ocean and atmosphere”

What I learned from the atmospheric science community :

- There remains skepticism about how effective adaptive mesh refinement (AMR) can be in weather and climate modeling
- Coarse/fine boundaries with abrupt resolution changes are a source of much numerical hand-wringing,

“Multi-scale numerics for the ocean and atmosphere”

What I learned from the atmospheric science community :

- There remains skepticism about how effective adaptive mesh refinement (AMR) can be in weather and climate modeling
- Coarse/fine boundaries with abrupt resolution changes are a source of much numerical hand-wringing,
- Multirate time stepping is not always used,

“Multi-scale numerics for the ocean and atmosphere”

What I learned from the atmospheric science community :

- There remains skepticism about how effective adaptive mesh refinement (AMR) can be in weather and climate modeling
- Coarse/fine boundaries with abrupt resolution changes are a source of much numerical hand-wringing,
- Multirate time stepping is not always used,
- Lack of good refinement criteria dampens enthusiasm for trying out AMR,

“Multi-scale numerics for the ocean and atmosphere”

What I learned from the atmospheric science community :

- There remains skepticism about how effective adaptive mesh refinement (AMR) can be in weather and climate modeling
- Coarse/fine boundaries with abrupt resolution changes are a source of much numerical hand-wringing,
- Multirate time stepping is not always used,
- Lack of good refinement criteria dampens enthusiasm for trying out AMR,
- When AMR *is* used in numerical weather prediction, the goals are often modest : “Do no harm!”

“Multi-scale numerics for the ocean and atmosphere”

What I learned from the atmospheric science community :

- There remains skepticism about how effective adaptive mesh refinement (AMR) can be in weather and climate modeling
- Coarse/fine boundaries with abrupt resolution changes are a source of much numerical hand-wringing,
- Multirate time stepping is not always used,
- Lack of good refinement criteria dampens enthusiasm for trying out AMR,
- When AMR *is* used in numerical weather prediction, the goals are often modest : “Do no harm!”
- Grids are often only static; not dynamically refined.

# Work with AMR

*Mesh generation and mesh adaptation for large-scale Earth-system modelling, (Phil. Trans. Roy. Soc, 2009)* compiled and edited by N. Nikiforakis

- M. Berger, et al (AMRClaw)
- C. Jablonowski et al
- H. Weller (INI organizer)
- C. Gatti-Bono and P. Colella (Chombo)
- R. Klein, N. Nikiforakis et al
- J. Behrens et al
- G. Pau, A. Almgren, J. Bell (LBL, California)
- C. Castro et al
- M. Piggot et al

# Finite volume schemes

# Finite volume schemes

- Second order finite volume schemes on logically Cartesian meshes,

# Finite volume schemes

- Second order finite volume schemes on logically Cartesian meshes,
- PDEs are typically solved in “conservation” form, and so discrete conservation of mass, momentum and so on is assured,

# Finite volume schemes

- Second order finite volume schemes on logically Cartesian meshes,
- PDEs are typically solved in “conservation” form, and so discrete conservation of mass, momentum and so on is assured,
- Performance efficiency is obtained by using locally solution adapted nested grids (adaptive mesh refinement, or AMR)

# Finite volume schemes

- Second order finite volume schemes on logically Cartesian meshes,
- PDEs are typically solved in “conservation” form, and so discrete conservation of mass, momentum and so on is assured,
- Performance efficiency is obtained by using locally solution adapted nested grids (adaptive mesh refinement, or AMR)
- Geometry handled using grid mappings, or cut-cell methods, in conjunction with multi-block domains.

# Finite volume schemes

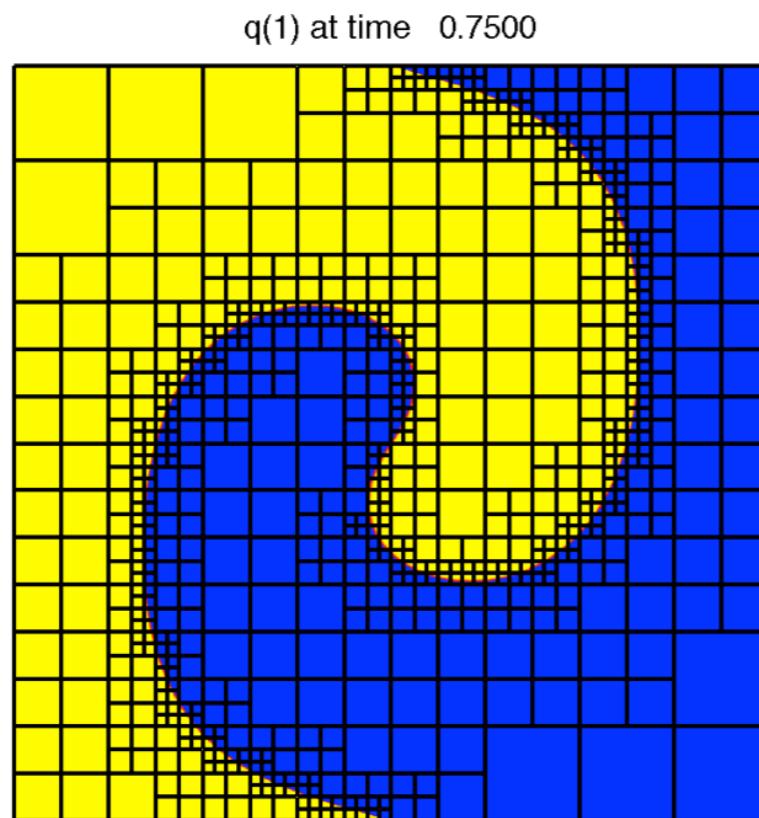
- Second order finite volume schemes on logically Cartesian meshes,
- PDEs are typically solved in “conservation” form, and so discrete conservation of mass, momentum and so on is assured,
- Performance efficiency is obtained by using locally solution adapted nested grids (adaptive mesh refinement, or AMR)
- Geometry handled using grid mappings, or cut-cell methods, in conjunction with multi-block domains.
- Widespread use in science and engineering,

# Finite volume schemes

- Second order finite volume schemes on logically Cartesian meshes,
- PDEs are typically solved in “conservation” form, and so discrete conservation of mass, momentum and so on is assured,
- Performance efficiency is obtained by using locally solution adapted nested grids (adaptive mesh refinement, or AMR)
- Geometry handled using grid mappings, or cut-cell methods, in conjunction with multi-block domains.
- Widespread use in science and engineering,
- “Collocated”, “P0”, “A-Grids”, “ijk” grids, non-conforming (for AMR) meshes

# Many flavors of adaptivity

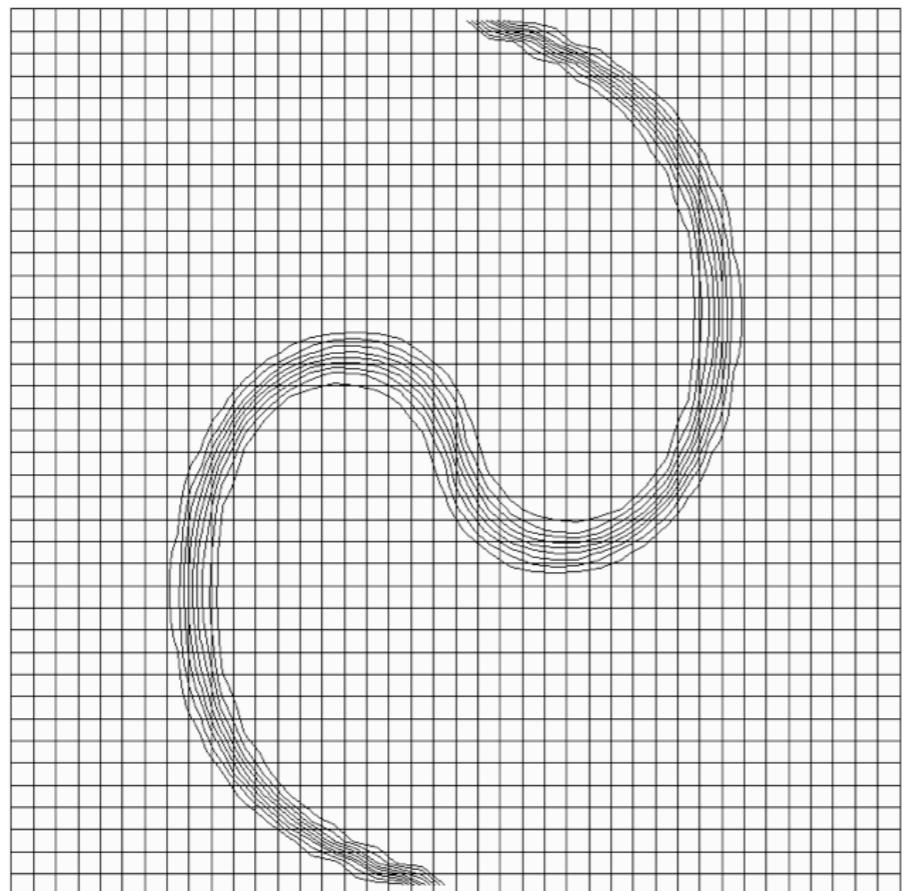
- Block-structured AMR (Berger, Oliger, Colella, ...)
- Tree-based adaptivity (Popinet, Tessyier, ...)
- Finite-element adaptivity includes both h-refinement (increase mesh resolution) and p-refinement (increase order of accuracy)



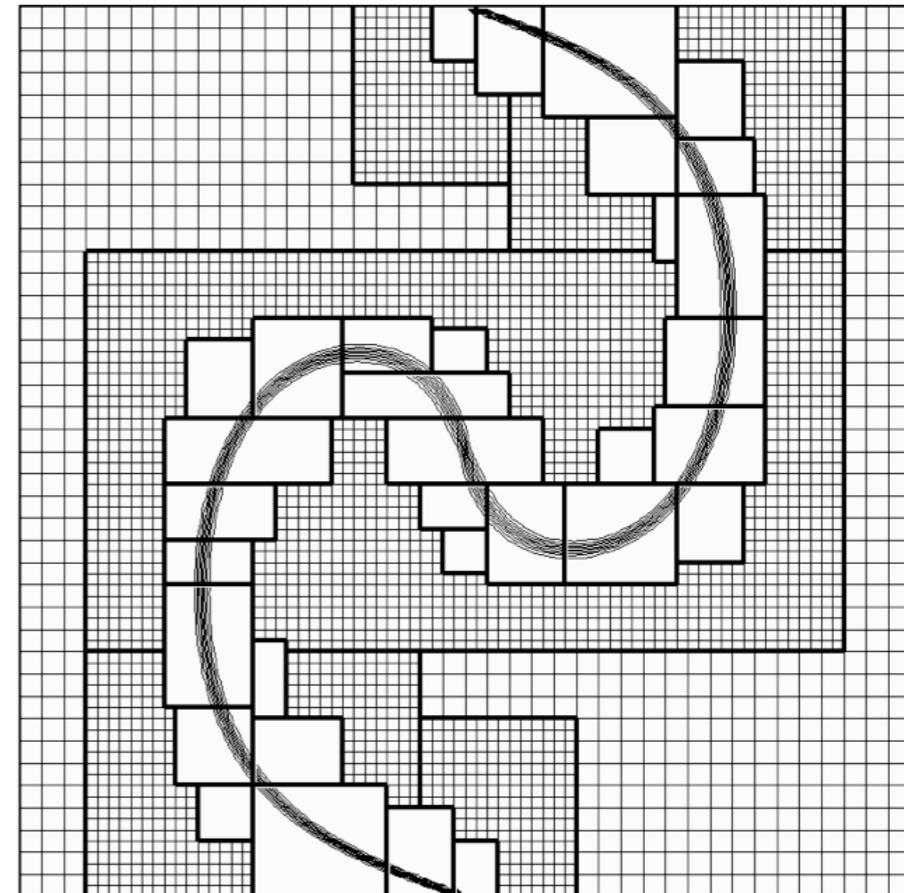
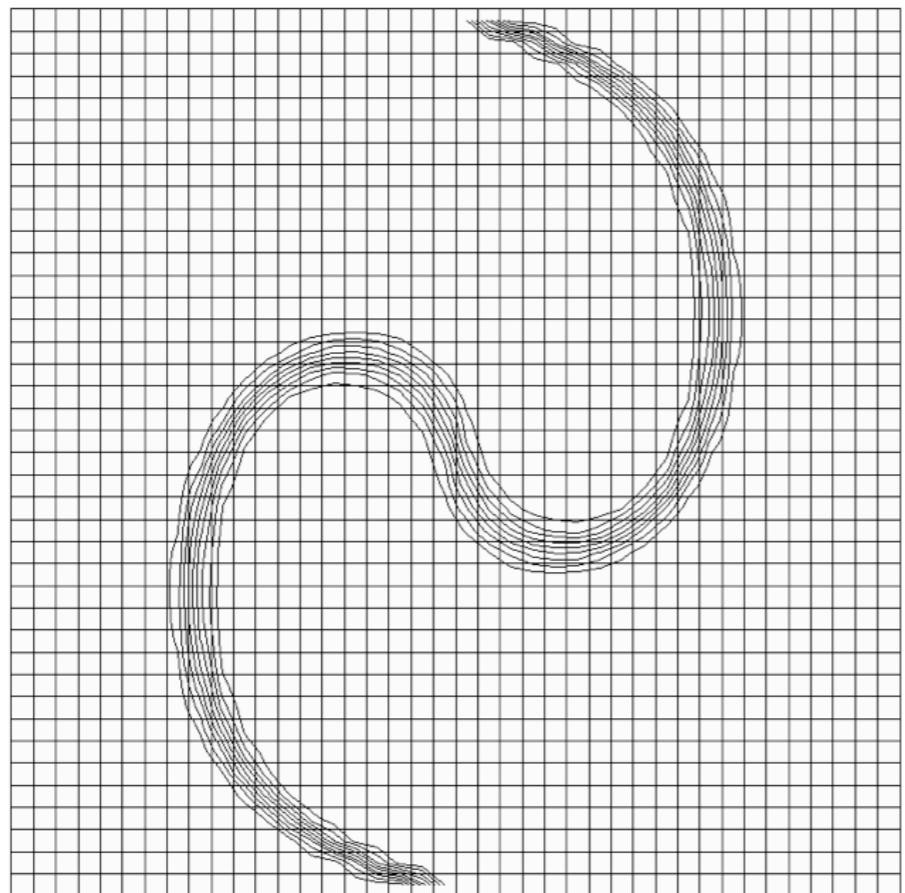
Tree-based adaptivity :

- *Gerris* (S. Popinet, NIWA, NZ),
- *Ramses* (R. Tessyier) and many other codes (including several in astrophysics)

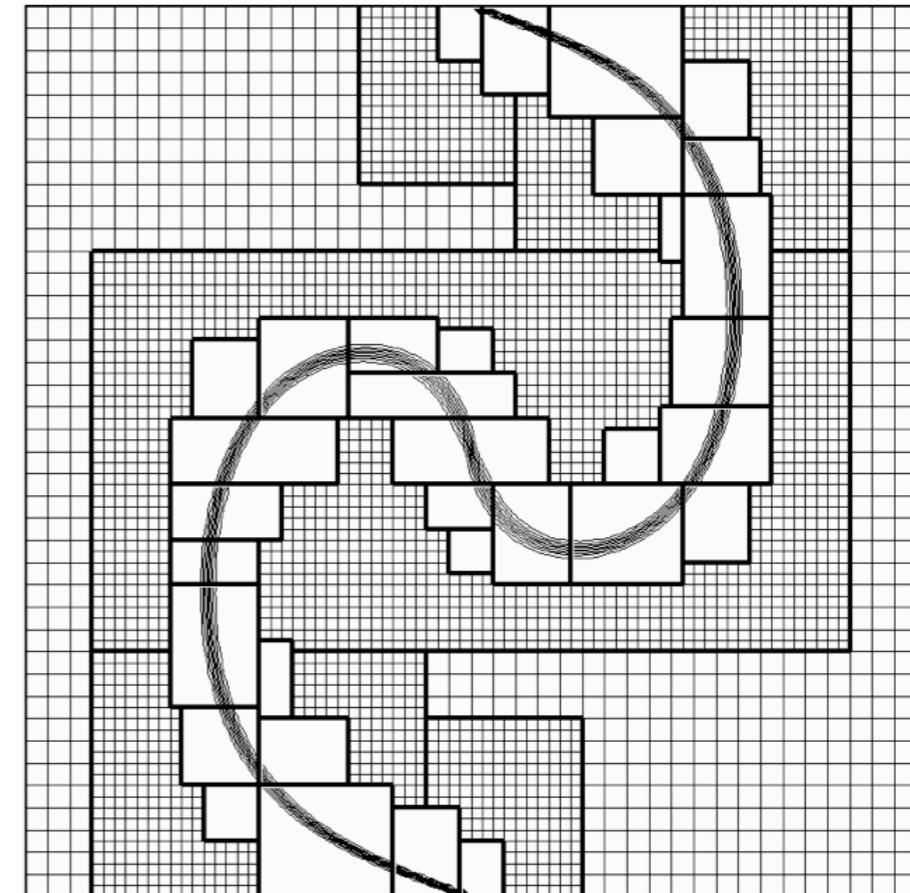
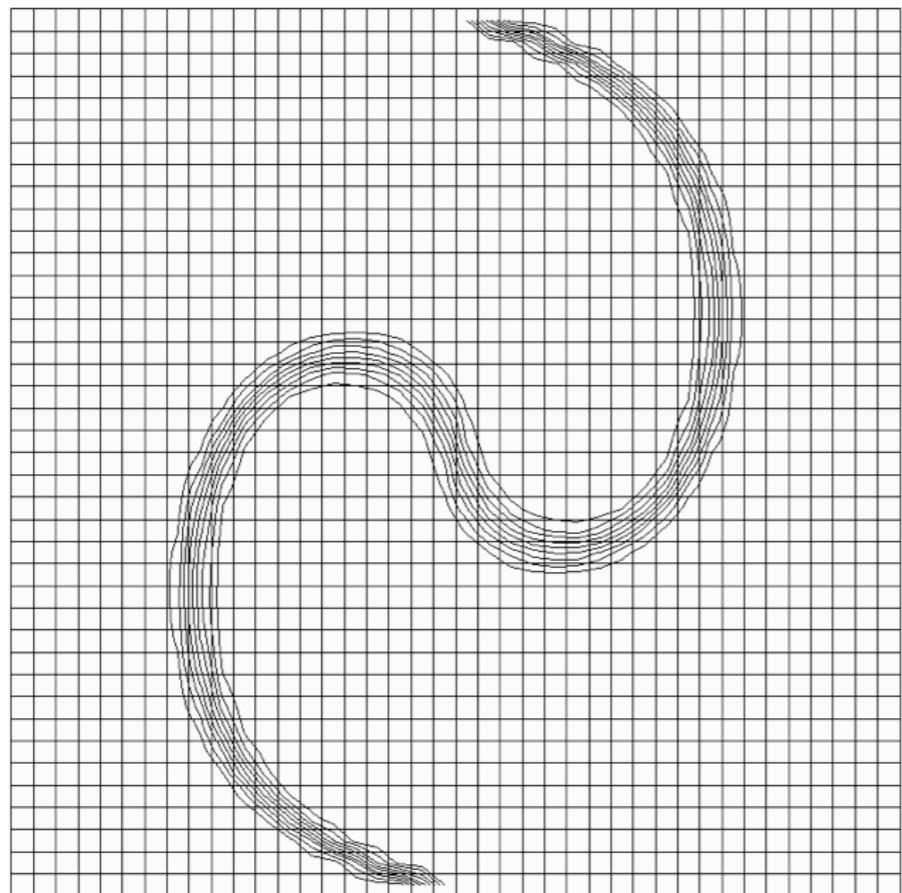
# Block structured AMR (Berger, Oliger 1984)



# Block structured AMR (Berger, Oliger 1984)

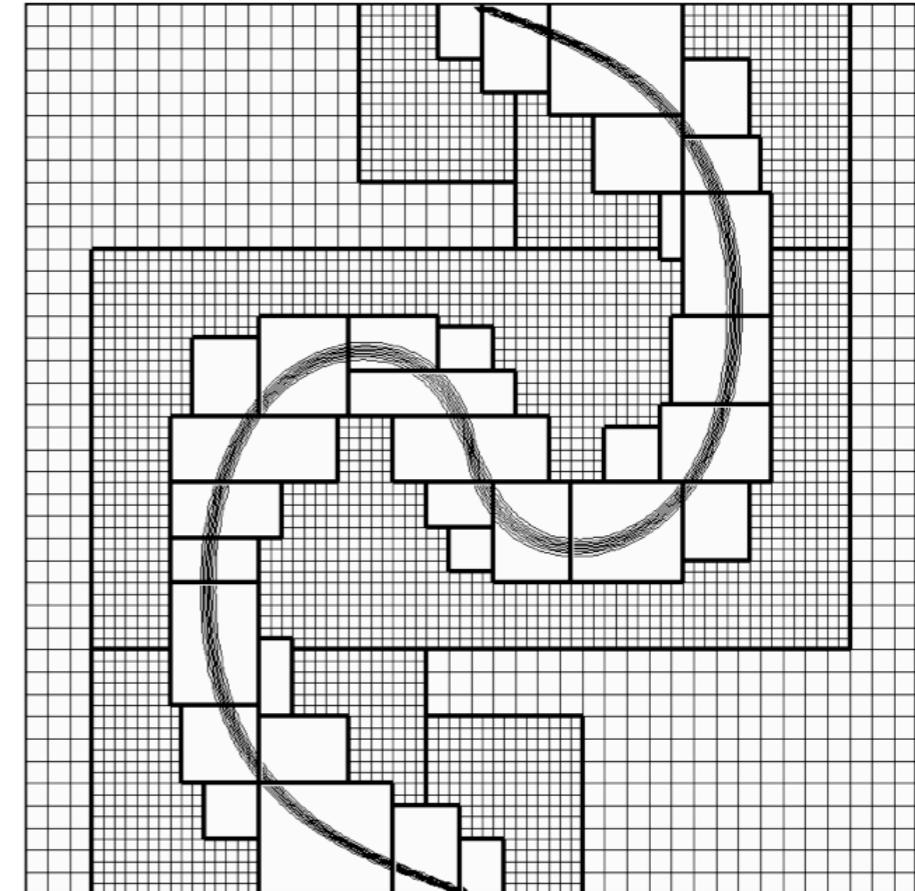
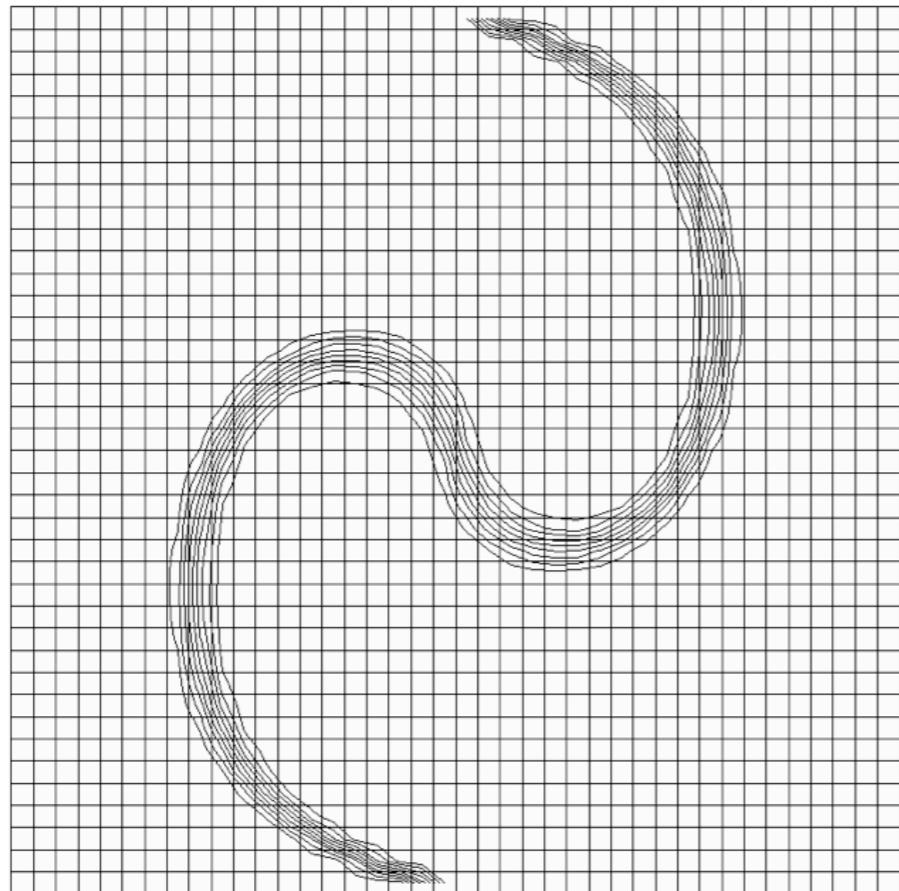


# Block structured AMR (Berger, Oliger 1984)



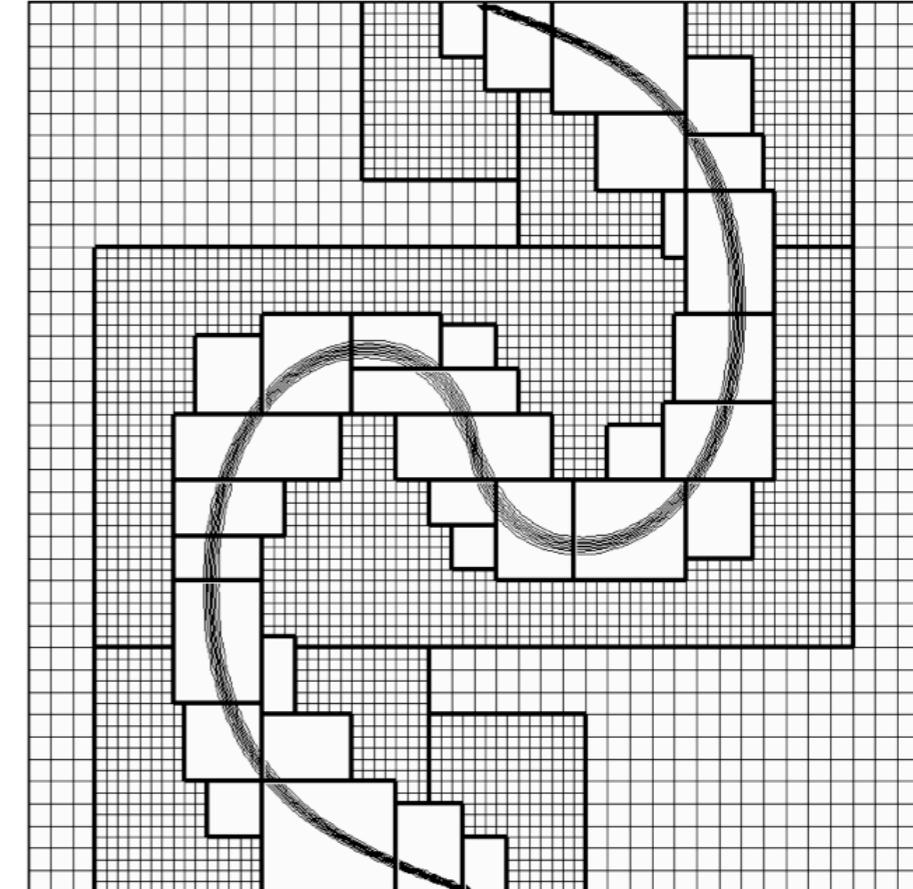
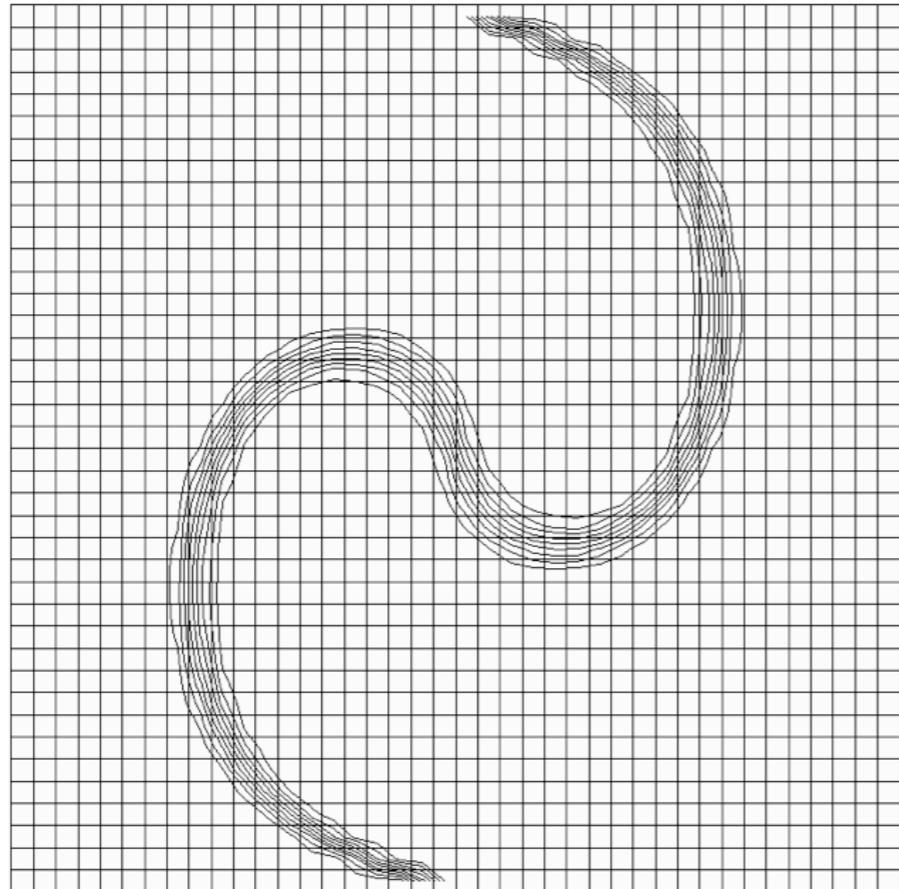
- Originally designed to improve shock capturing methods

# Block structured AMR (Berger, Oliger 1984)



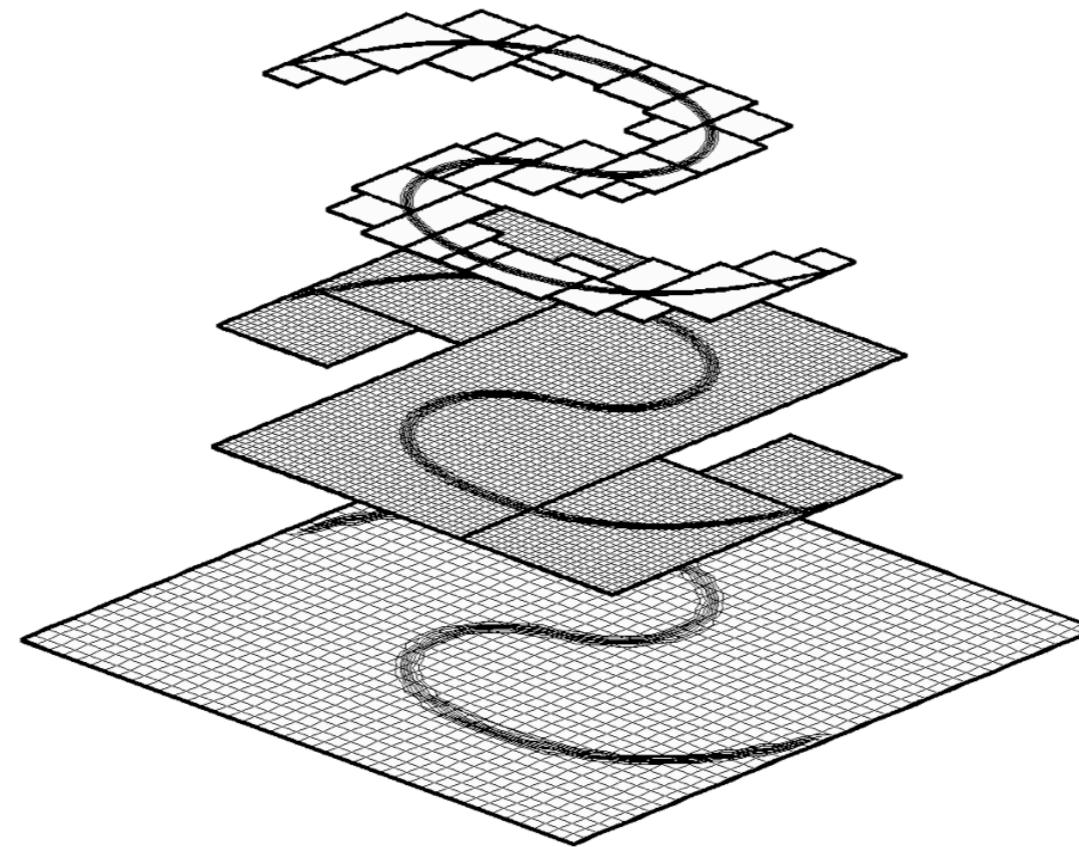
- Originally designed to improve shock capturing methods
- Gained widespread use in many application areas

# Block structured AMR (Berger, Oliger 1984)

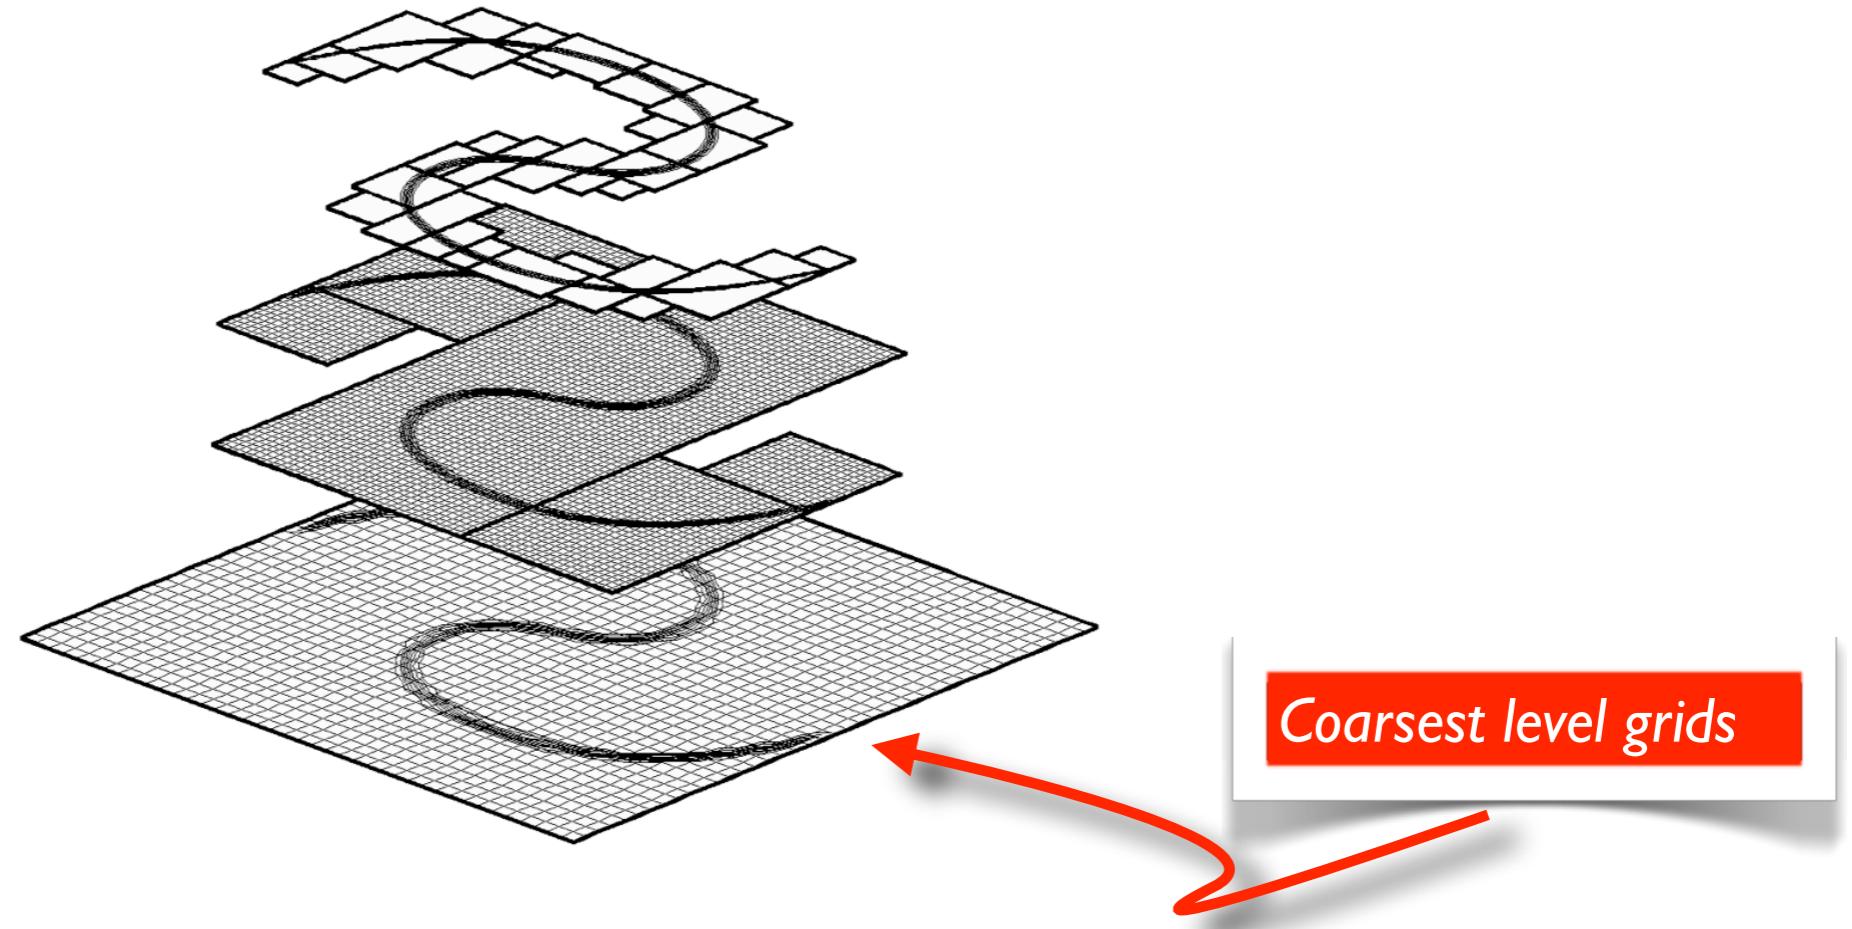


- Originally designed to improve shock capturing methods
- Gained widespread use in many application areas
- Colella, Bell, LeVeque, Almgren, Deiterding, and many others have developed methods and solvers

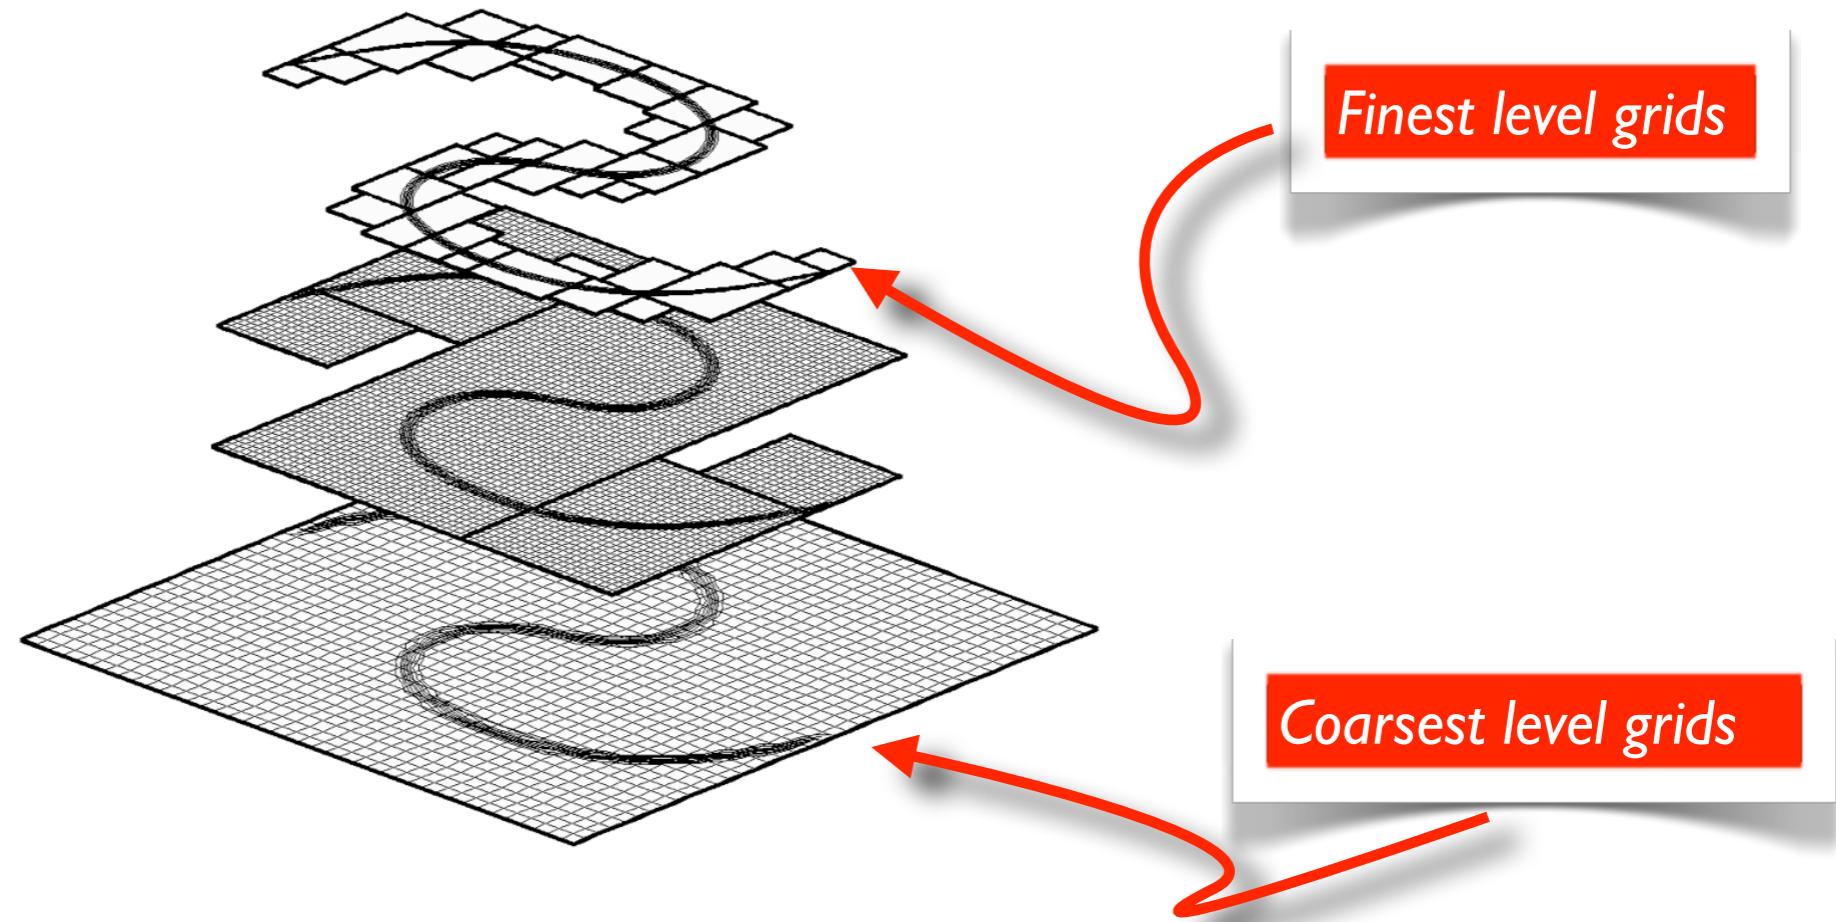
# Block structured AMR (ala Berger and Oliger)



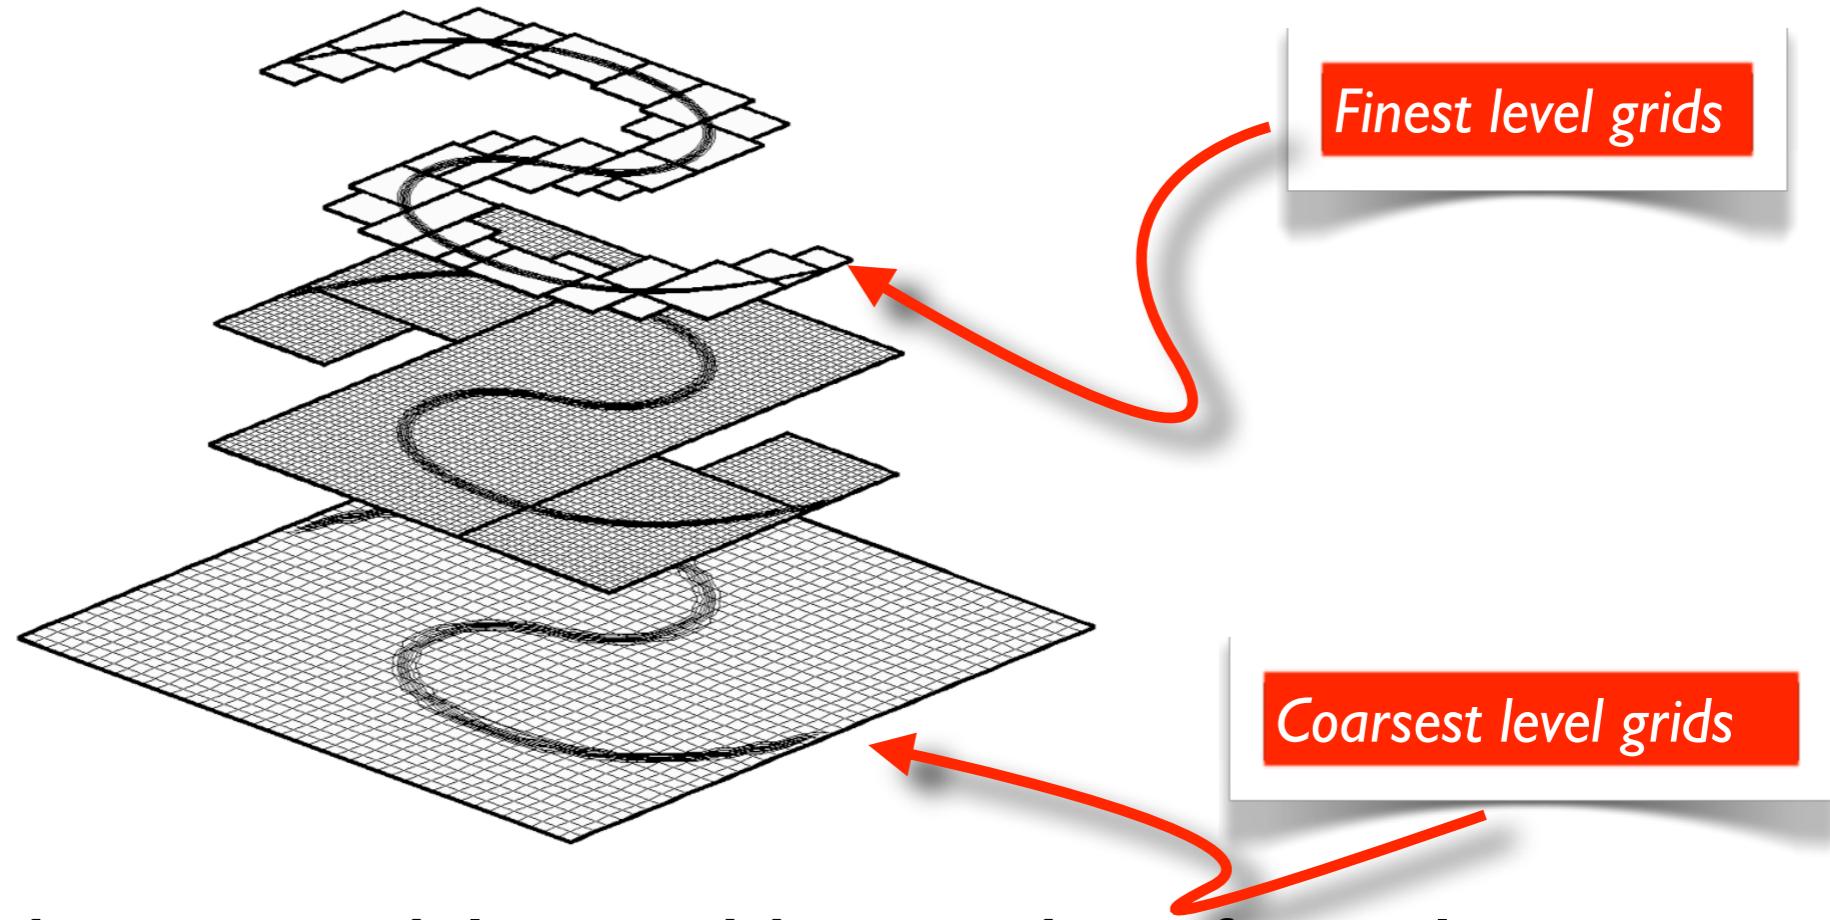
# Block structured AMR (ala Berger and Olliger)



# Block structured AMR (ala Berger and Olliger)

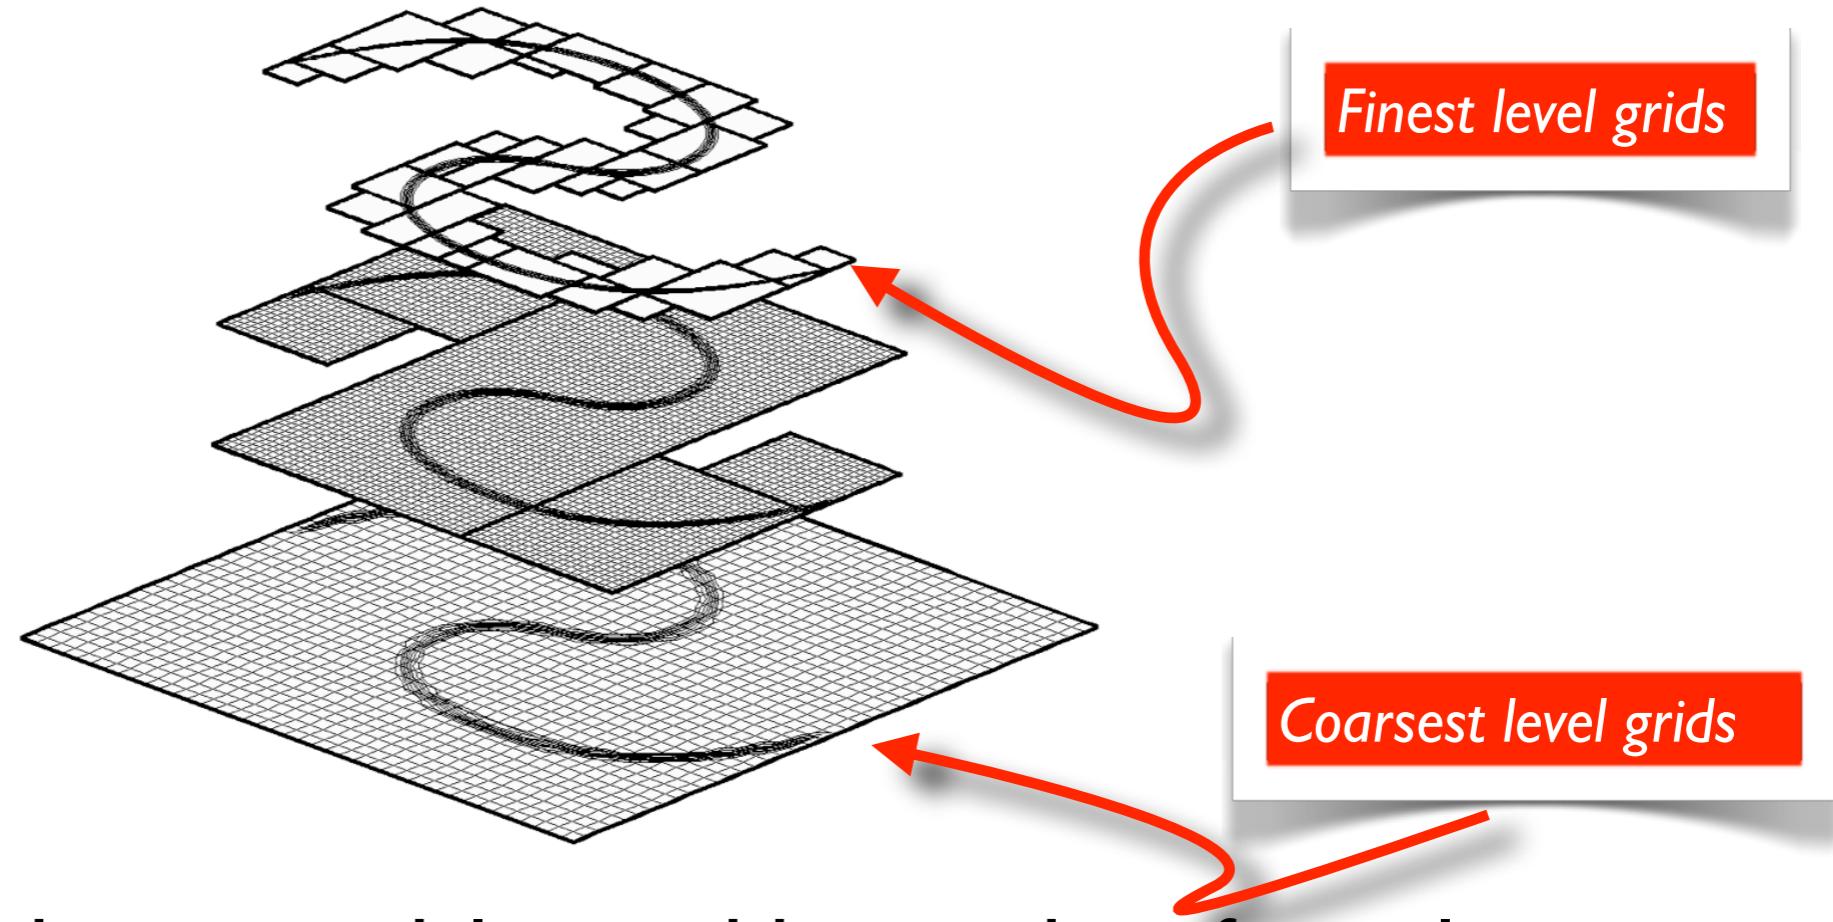


# Block structured AMR (ala Berger and Olliger)



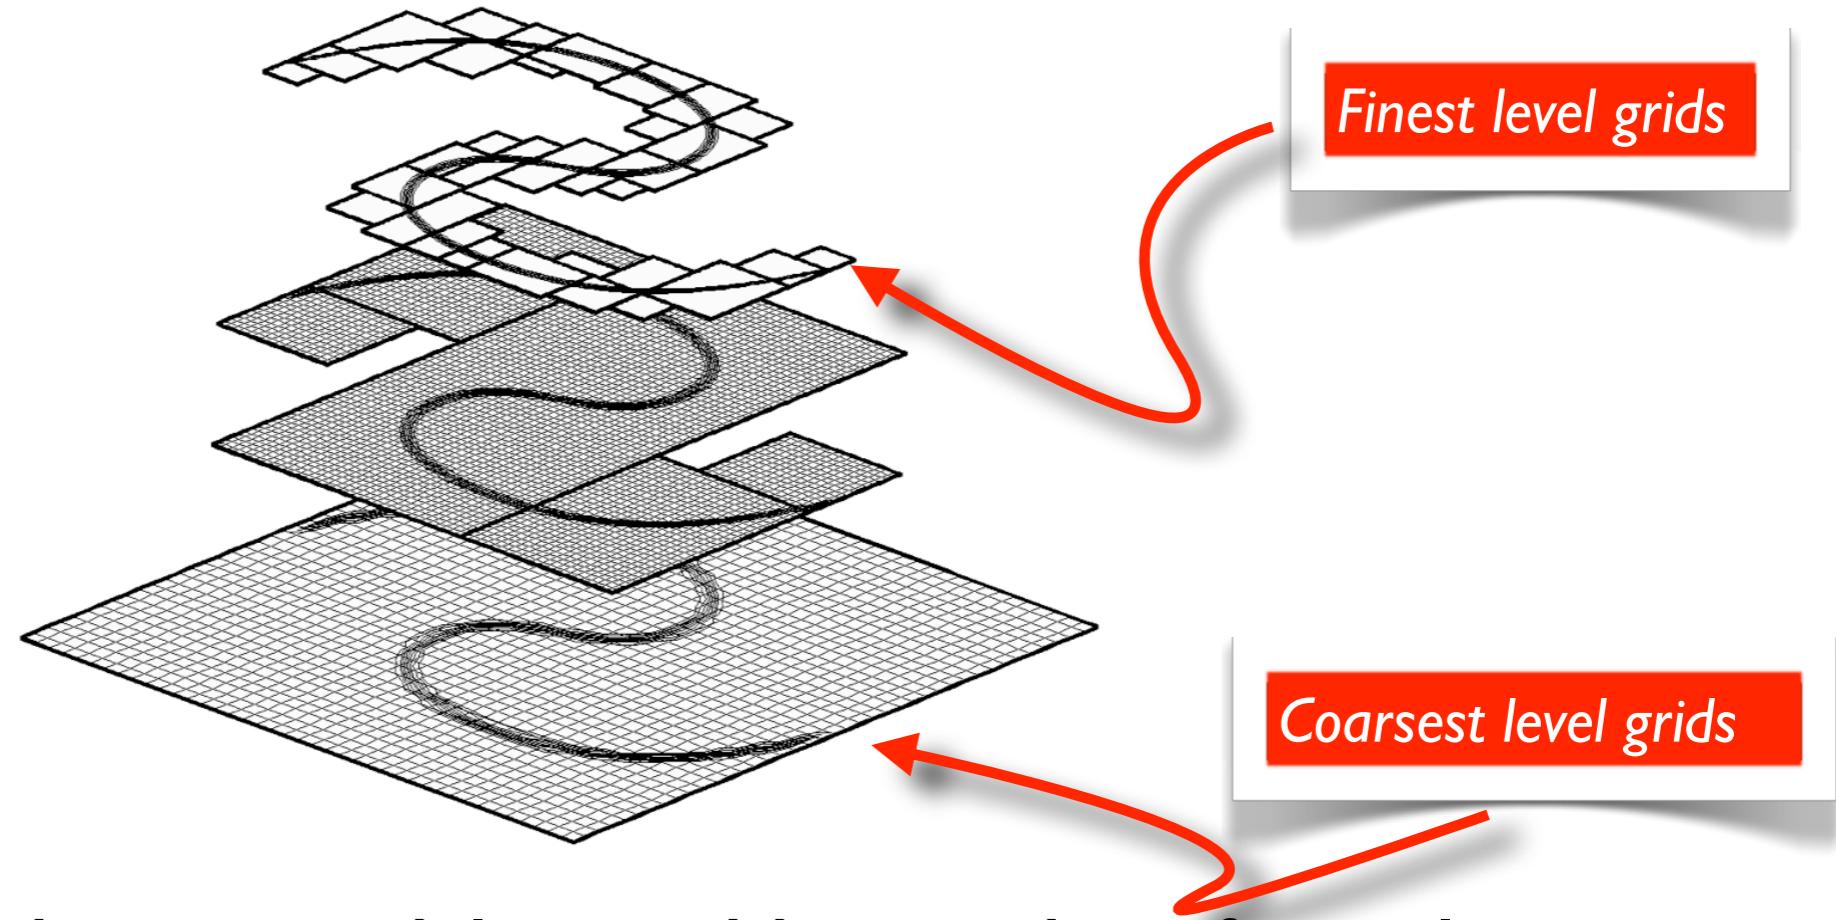
- Data is stored in nested, layered hierarchy of overlapping, logically Cartesian grids,

# Block structured AMR (ala Berger and Oliger)



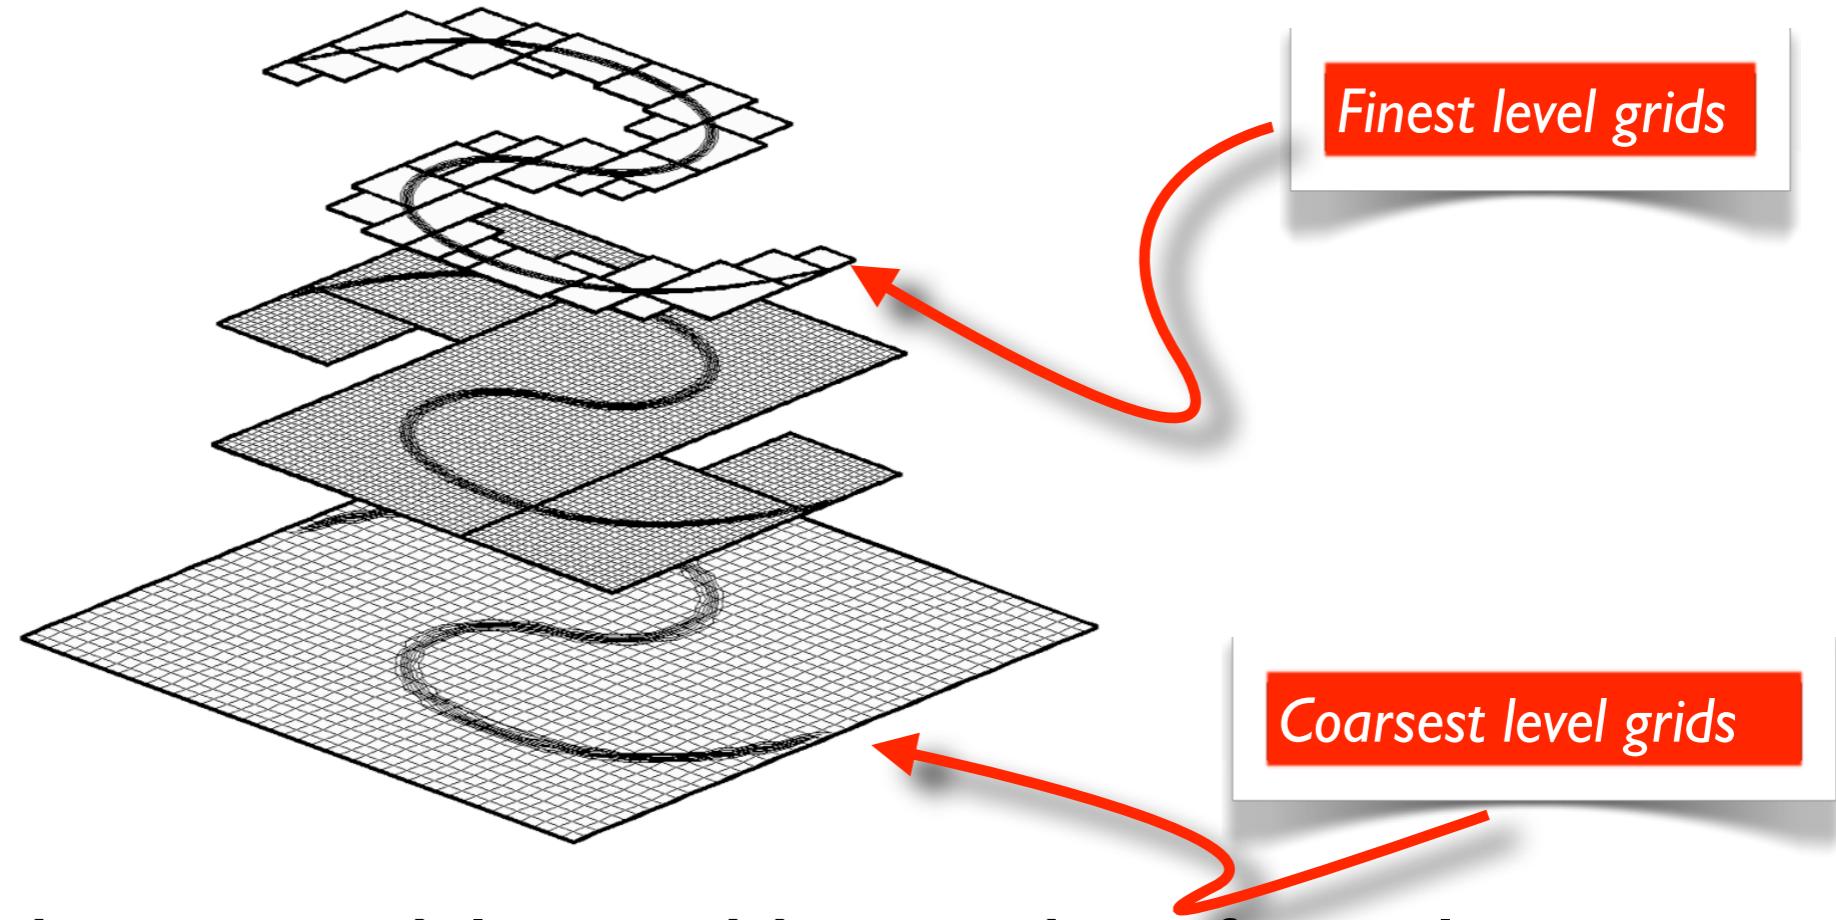
- Data is stored in nested, layered hierarchy of overlapping, logically Cartesian grids,
- Multi-rate time stepping based on mesh size,

# Block structured AMR (ala Berger and Oliger)



- Data is stored in nested, layered hierarchy of overlapping, logically Cartesian grids,
- Multi-rate time stepping based on mesh size,
- Grids are dynamically refined and de-refined to adapt to the solution features of interest.

# Block structured AMR (ala Berger and Olliger)



- Data is stored in nested, layered hierarchy of overlapping, logically Cartesian grids,
- Multi-rate time stepping based on mesh size,
- Grids are dynamically refined and de-refined to adapt to the solution features of interest.
- Communication between grids is done via ghost cells.

# Block structured AMR rules

# Block structured AMR rules

- Coarse grid and fine grid boundaries are aligned

# Block structured AMR rules

- Coarse grid and fine grid boundaries are aligned
- Finer meshes are properly nested into coarser ones

# Block structured AMR rules

- Coarse grid and fine grid boundaries are aligned
- Finer meshes are properly nested into coarser ones
- Ghost cell values are obtained from coarse grid or neighboring fine grids, if available

# Block structured AMR rules

- Coarse grid and fine grid boundaries are aligned
- Finer meshes are properly nested into coarser ones
- Ghost cell values are obtained from coarse grid or neighboring fine grids, if available
- Do not allow grids which overlap multiple levels of refinement

# Block structured AMR rules

- Coarse grid and fine grid boundaries are aligned
- Finer meshes are properly nested into coarser ones
- Ghost cell values are obtained from coarse grid or neighboring fine grids, if available
- Do not allow grids which overlap multiple levels of refinement
- Averaging fine grid solution to coarse grid; interpolate coarse grid solution to the fine grid.

# Block structured AMR rules

- Coarse grid and fine grid boundaries are aligned
- Finer meshes are properly nested into coarser ones
- Ghost cell values are obtained from coarse grid or neighboring fine grids, if available
- Do not allow grids which overlap multiple levels of refinement
- Averaging fine grid solution to coarse grid; interpolate coarse grid solution to the fine grid.
- Buffer cells keep solution features on finest level

# Block structured AMR rules

- Coarse grid and fine grid boundaries are aligned
- Finer meshes are properly nested into coarser ones
- Ghost cell values are obtained from coarse grid or neighboring fine grids, if available
- Do not allow grids which overlap multiple levels of refinement
- Averaging fine grid solution to coarse grid; interpolate coarse grid solution to the fine grid.
- Buffer cells keep solution features on finest level
- Grid clustering algorithm balances number of grids and refinement efficiency.

# Block structured AMR philosophy

# Block structured AMR philosophy

- Take advantage of existing Cartesian grid solvers whenever possible,

# Block structured AMR philosophy

- Take advantage of existing Cartesian grid solvers whenever possible,
- Avoid use of complicated stencils at coarse/fine grid interfaces by operating locally on patches whenever possible

# Block structured AMR philosophy

- Take advantage of existing Cartesian grid solvers whenever possible,
- Avoid use of complicated stencils at coarse/fine grid interfaces by operating locally on patches whenever possible
- Numerical solution on grid hierarchy should have the same order of accuracy as the single grid algorithm.

# Block structured AMR philosophy

- Take advantage of existing Cartesian grid solvers whenever possible,
- Avoid use of complicated stencils at coarse/fine grid interfaces by operating locally on patches whenever possible
- Numerical solution on grid hierarchy should have the same order of accuracy as the single grid algorithm.
- Conservation should be maintained if PDE is in conservative form

# Block structured AMR philosophy

- Take advantage of existing Cartesian grid solvers whenever possible,
- Avoid use of complicated stencils at coarse/fine grid interfaces by operating locally on patches whenever possible
- Numerical solution on grid hierarchy should have the same order of accuracy as the single grid algorithm.
- Conservation should be maintained if PDE is in conservative form
- Overhead in managing multiple grid levels should not impact performance significantly

# Block structured AMR philosophy

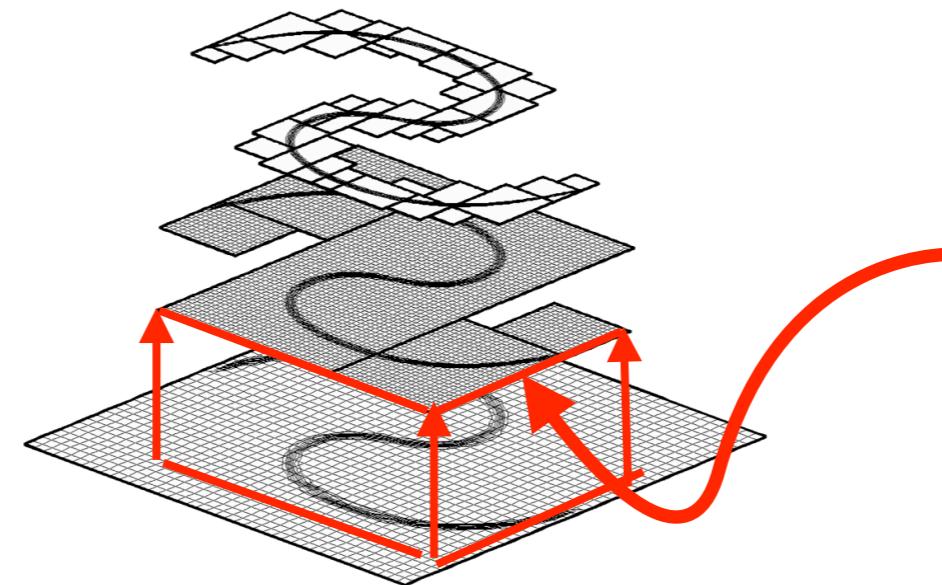
- Take advantage of existing Cartesian grid solvers whenever possible,
- Avoid use of complicated stencils at coarse/fine grid interfaces by operating locally on patches whenever possible
- Numerical solution on grid hierarchy should have the same order of accuracy as the single grid algorithm.
- Conservation should be maintained if PDE is in conservative form
- Overhead in managing multiple grid levels should not impact performance significantly
- Subcycle finer grid time steps (multi-rate time stepping)

# Explicit, single step multi-rate time stepping

A single time step advance, assuming a refinement factor of  $R$ .

1. Advance at the coarsest level by time step  $\Delta t$
2. Interpolate coarse grid solution to fine grid ghost cells
3. Advance fine grid  $R$  time steps, by a time step  $\Delta t/R$
4. Average solution from fine grids to coarse grid,
5. Adjust coarse grid solution to assure flux continuity at the coarse/fine boundaries,
6. Tag cells for refinement and regrid

*Grids at the same level exchange ghost cell values directly*



*Fine grid boundary conditions interpolated in space and time from coarse grid*

# Software for AMR

*How can I add adaptivity to my existing (Cartesian) code?*

- It is much harder than it looks (and you don't really just “add” adaptivity)
- And there are several general purpose codes already available which can use your single grid solver.

# Block structured AMR codes

- General purpose (freely available) block-structured codes
  - **PARAMESH** (NASA/Goddard)
  - **SAMRAI** (Lawrence Livermore National Lab)
  - **BoxLib** (Lawrence Berkeley Lab)
  - **Chombo** (Lawrence Berkeley Lab)
  - **AMRClaw** (University of Washington/NYU)
- All are large frameworks, with many developers
- Mostly C++ and Fortran libraries (no GUIs) that started life as research codes.

*See my website for a list of several more application specific codes*

# Block structured AMR codes

“PARAMESH is a package of Fortran 90 subroutines designed to provide an application developer with an easy route to extend an existing serial code which uses logically Cartesian structured mesh into a parallel code with adaptive mesh refinement”

SAMRAI - “Object oriented C++ library developed to provide algorithmic and software support to large scale multiphysics problems relevant to the US Department of Energy (DOE)”

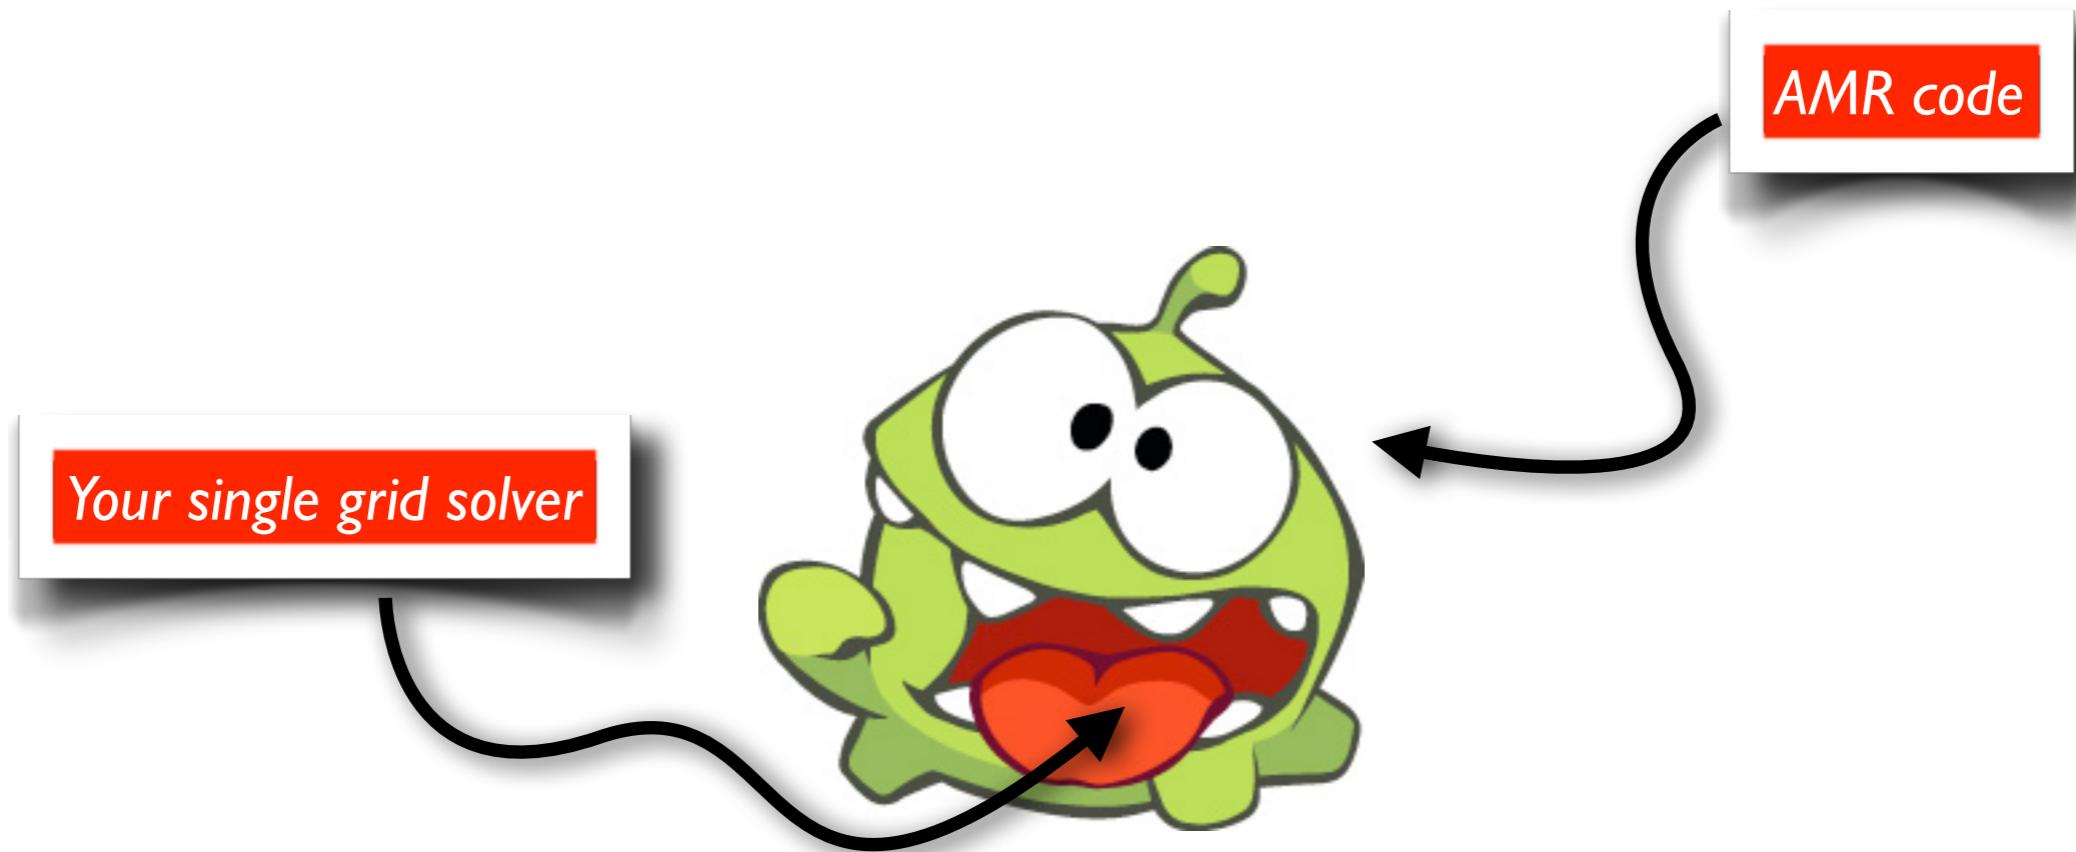
Boxlib - “These libraries provide the software infrastructure for the computational activities (combustion, astrophysics, porous media flow) at the Center for Computational Sciences and Engineering (CCSE) at Lawrence Berkeley Labs”

# Using block structured AMR codes

Mental model of how this might work :

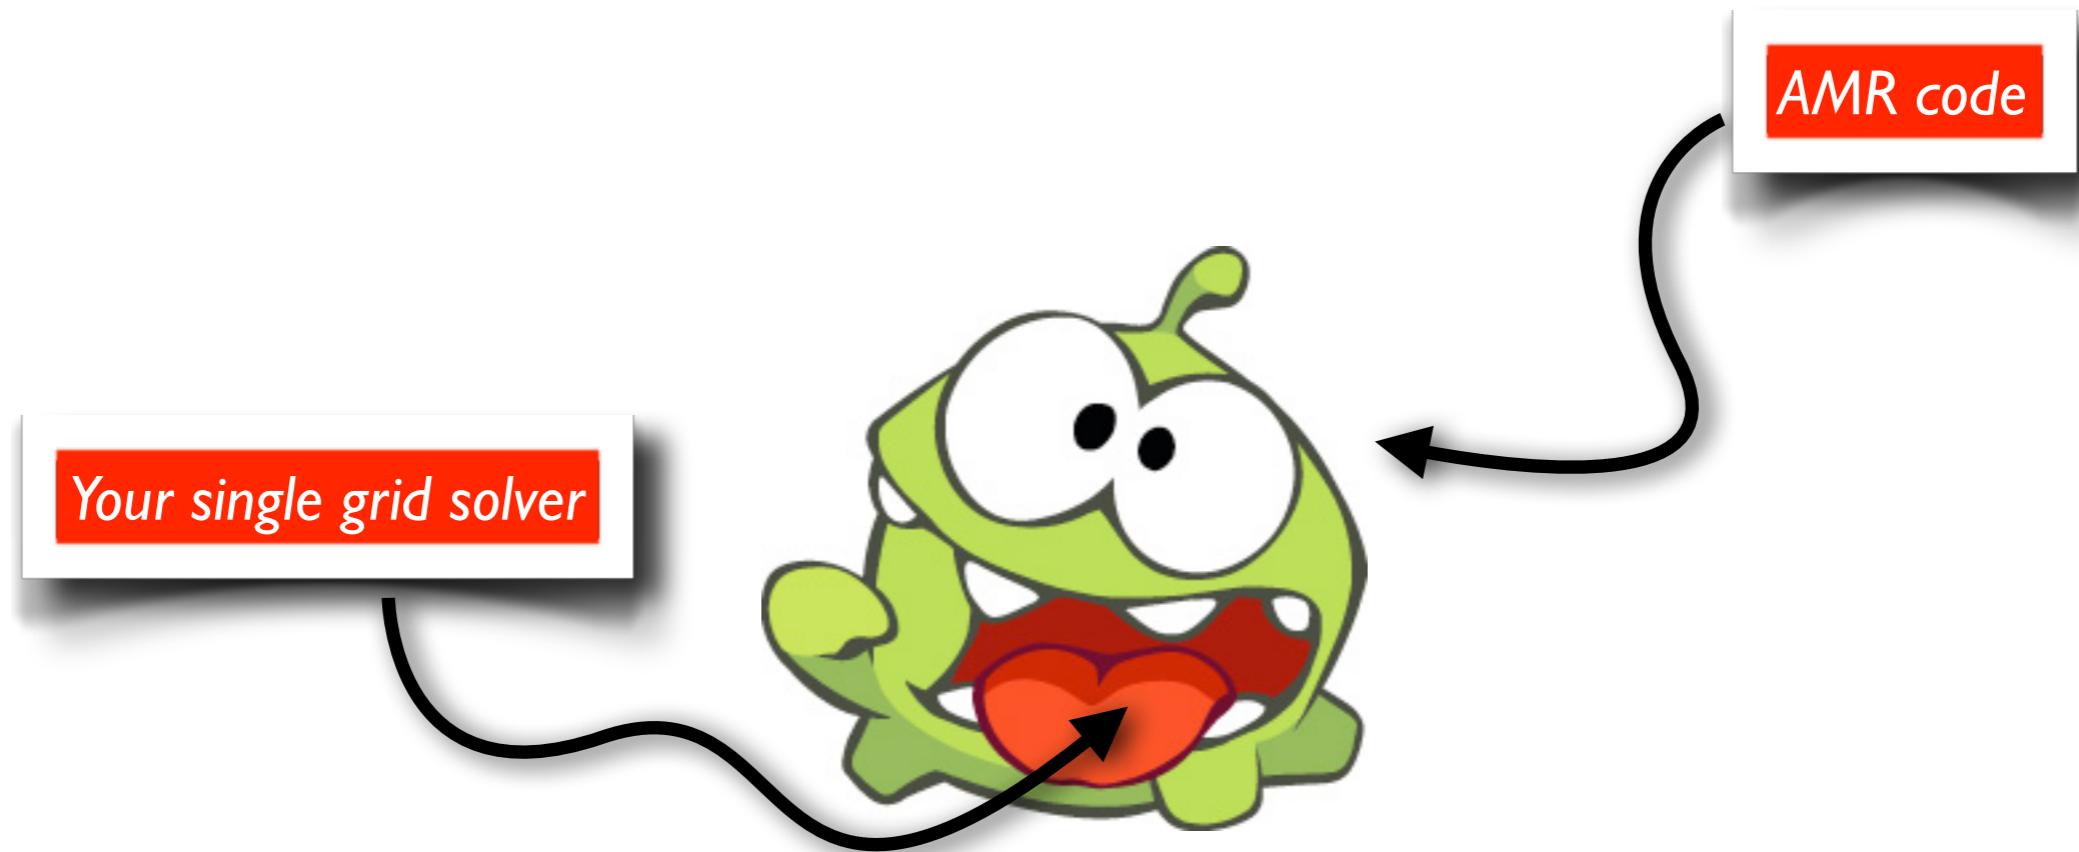
# Using block structured AMR codes

Mental model of how this might work :



# Using block structured AMR codes

Mental model of how this might work :



\* Idea for code name : OmNum

# Block structured AMR codes

## The Dream

# Block structured AMR codes

## The Dream

```
AMR.run(max_time, max_steps);
```

# Block structured AMR codes

## The Dream

```
AMR.run(max_time, max_steps);
```



*Your single grid solver  
is called from here.*

# Block structured AMR codes

## The Dream

```
AMR.run(max_time, max_steps);
```

*Your single grid solver  
is called from here.*



# Block structured AMR codes

## The Reality

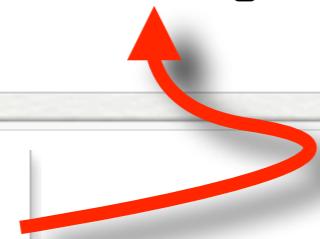
```
Tuple< RefCountedPtr<AMRLevelOpFactory<  
LevelData<FArrayBox> > >, SpaceDim> velTGAOpFactoryPtrs;  
  
for (int idir = 0; idir < SpaceDim; idir++)  
{velTGAOpFactoryPtrs[idir] =  
    RefCountedPtr<AMRLevelOpFactory<LevelData  
<FArrayBox> > >  
    ( (AMRLevelOpFactory<LevelData<FArrayBox> >*)  
    (new AMRPoissonOpFactory()) ); //.....
```

# Block structured AMR codes

## The Reality

```
Tuple< RefCountedPtr<AMRLevelOpFactory<  
LevelData<FArrayBox> > >, SpaceDim> velTGAOpFactoryPtrs;  
  
for (int idir = 0; idir < SpaceDim; idir++)  
{velTGAOpFactoryPtrs[idir] =  
    RefCountedPtr<AMRLevelOpFactory<LevelData  
<FArrayBox> > >  
    ( (AMRLevelOpFactory<LevelData<FArrayBox> >*)  
    (new AMRPoissonOpFactory()) ); //.....
```

Your single grid solver



# Block structured AMR codes

## The Reality

```
Tuple< RefCountedPtr<AMRLevelOpFactory<  
LevelData<FArrayBox> > >, SpaceDim> velTGAOpFactoryPtrs;  
  
for (int idir = 0; idir < SpaceDim; idir++)  
{velTGAOpFactoryPtrs[idir] =  
    RefCountedPtr<AMRLevelOpFactory<LevelData  
<FArrayBox> > >  
    ( (AMRLevelOpFactory<LevelData<FArrayBox> >*)  
    (new AMRPoissonOpFactory()) ); //.....
```

Your single grid solver

you



# Retro...

```
node(ndjhi,mptrnx) = node(ndjhi,mptr)
node(ndjhi,mptr)   = node(ndjlo,mptr) + nyl - 1
node(ndjlo,mptrnx) = node(ndjhi,mptr) + 1
node(ndihi,mptrnx) = node(ndihi,mptr)
node(ndilo,mptrnx) = node(ndilo,mptr)

rnode(cornxlo,mptrnx)    = cxlo
rnode(cornylo,mptrnx)    = cymid
rnode(cornyhi,mptrnx)    = cyhi
rnode(cornxhi,mptrnx)    = cxhi
node(nestlevel,mptrnx)   = node(nestlevel,mptr)
rnode(timemult,mptrnx)   = rnode(timemult,mptr)
go to 10
```

# Retro...

```
node(ndjhi,mptrnx) = node(ndjhi,mptr)
node(ndjhi,mptr)   = node(ndjlo,mptr) + nyl - 1
node(ndjlo,mptrnx) = node(ndjhi,mptr) + 1
node(ndihi,mptrnx) = node(ndihi,mptr)
node(ndilo,mptrnx) = node(ndilo,mptr)

rnode(cornxlo,mptrnx)    = cxlo
rnode(cornylo,mptrnx)    = cymid
rnode(cornyhi,mptrnx)    = cyhi
rnode(cornxhi,mptrnx)    = cxhi
node(nestlevel,mptrnx)   = node(nestlevel,mptr)
rnode(timemult,mptrnx)   = rnode(timemult,mptr)
go to 10
```



# Why are AMR codes challenging to develop?

# Why are AMR codes challenging to develop?

- Heterogeneous data structures associated with storing hierarchy of grids,

# Why are AMR codes challenging to develop?

- Heterogeneous data structures associated with storing hierarchy of grids,
- Dynamically creating and destroying grids :

# Why are AMR codes challenging to develop?

- Heterogeneous data structures associated with storing hierarchy of grids,
- Dynamically creating and destroying grids :
  - Need a “factory” paradigm to create new grids and any auxiliary data arrays (material properties, metric terms, bathymetry, etc) that go with each new grid,

# Why are AMR codes challenging to develop?

- Heterogeneous data structures associated with storing hierarchy of grids,
- Dynamically creating and destroying grids :
  - Need a “factory” paradigm to create new grids and any auxiliary data arrays (material properties, metric terms, bathymetry, etc) that go with each new grid,
- Periodic boundary conditions and multi-block domains,

# Why are AMR codes challenging to develop?

- Heterogeneous data structures associated with storing hierarchy of grids,
- Dynamically creating and destroying grids :
  - Need a “factory” paradigm to create new grids and any auxiliary data arrays (material properties, metric terms, bathymetry, etc) that go with each new grid,
- Periodic boundary conditions and multi-block domains,
- Providing support for multi-rate time stepping,

# Why are AMR codes challenging to develop?

- Heterogeneous data structures associated with storing hierarchy of grids,
- Dynamically creating and destroying grids :
  - Need a “factory” paradigm to create new grids and any auxiliary data arrays (material properties, metric terms, bathymetry, etc) that go with each new grid,
- Periodic boundary conditions and multi-block domains,
- Providing support for multi-rate time stepping,
- Storing information at grid boundaries needed to couple the solution on grids

# Why are AMR codes challenging to develop?

- Heterogeneous data structures associated with storing hierarchy of grids,
- Dynamically creating and destroying grids :
  - Need a “factory” paradigm to create new grids and any auxiliary data arrays (material properties, metric terms, bathymetry, etc) that go with each new grid,
- Periodic boundary conditions and multi-block domains,
- Providing support for multi-rate time stepping,
- Storing information at grid boundaries needed to couple the solution on grids
- Parallel load balancing for both the solution step and refinement step

# Why can AMR codes be difficult to use?

# Why can AMR codes be difficult to use?

- Time stepping beyond just single step explicit schemes  
(IMEX, SSP, RK, ...)

# Why can AMR codes be difficult to use?

- Time stepping beyond just single step explicit schemes (IMEX, SSP, RK, ...)
- Understanding how overall time stepping interacts with dynamic grid creation, destruction and management.

# Why can AMR codes be difficult to use?

- Time stepping beyond just single step explicit schemes (IMEX, SSP, RK, ...)
- Understanding how overall time stepping interacts with dynamic grid creation, destruction and management.
- Implicit solvers,

# Why can AMR codes be difficult to use?

- Time stepping beyond just single step explicit schemes (IMEX, SSP, RK, ...)
- Understanding how overall time stepping interacts with dynamic grid creation, destruction and management.
- Implicit solvers,
- Coupling of multiple solvers (e.g. advection + diffusion + elliptic solvers + source terms),

# Why can AMR codes be difficult to use?

- Time stepping beyond just single step explicit schemes (IMEX, SSP, RK, ...)
- Understanding how overall time stepping interacts with dynamic grid creation, destruction and management.
- Implicit solvers,
- Coupling of multiple solvers (e.g. advection + diffusion + elliptic solvers + source terms),
- Multi-physics,

# Why can AMR codes be difficult to use?

- Time stepping beyond just single step explicit schemes (IMEX, SSP, RK, ...)
- Understanding how overall time stepping interacts with dynamic grid creation, destruction and management.
- Implicit solvers,
- Coupling of multiple solvers (e.g. advection + diffusion + elliptic solvers + source terms),
- Multi-physics,
- Data visualization

# Why can AMR codes be difficult to use?

- Time stepping beyond just single step explicit schemes (IMEX, SSP, RK, ...)
- Understanding how overall time stepping interacts with dynamic grid creation, destruction and management.
- Implicit solvers,
- Coupling of multiple solvers (e.g. advection + diffusion + elliptic solvers + source terms),
- Multi-physics,
- Data visualization
- Computing diagnostics on a nested grid hierarchy,

# Why can AMR codes be difficult to use?

- Time stepping beyond just single step explicit schemes (IMEX, SSP, RK, ...)
- Understanding how overall time stepping interacts with dynamic grid creation, destruction and management.
- Implicit solvers,
- Coupling of multiple solvers (e.g. advection + diffusion + elliptic solvers + source terms),
- Multi-physics,
- Data visualization
- Computing diagnostics on a nested grid hierarchy,
- Error estimation, tuning for efficient use of grids, ...

# But what if you have ideas about ...

- Multi-stage, multi-step, IMEX, SSP, parallel-in-time, exponential integrators, and other time stepping schemes in an adaptive setting,
- Accuracy of multi-rate schemes for PDEs with mixed elliptic/parabolic/hyperbolic terms.
- Elliptic and parabolic solvers (iterative? direct? Explicit? Fast multipole?)
- Parallelism in the AMR setting?
- Error estimation
- Higher order accuracy
- Complex physics

*Should you write yet-another-AMR code?*

# ForestClaw : A hybrid approach to AMR

# ForestClaw : A hybrid approach to AMR

Use an existing parallel quad/octree library to do the grid management

# ForestClaw : A hybrid approach to AMR

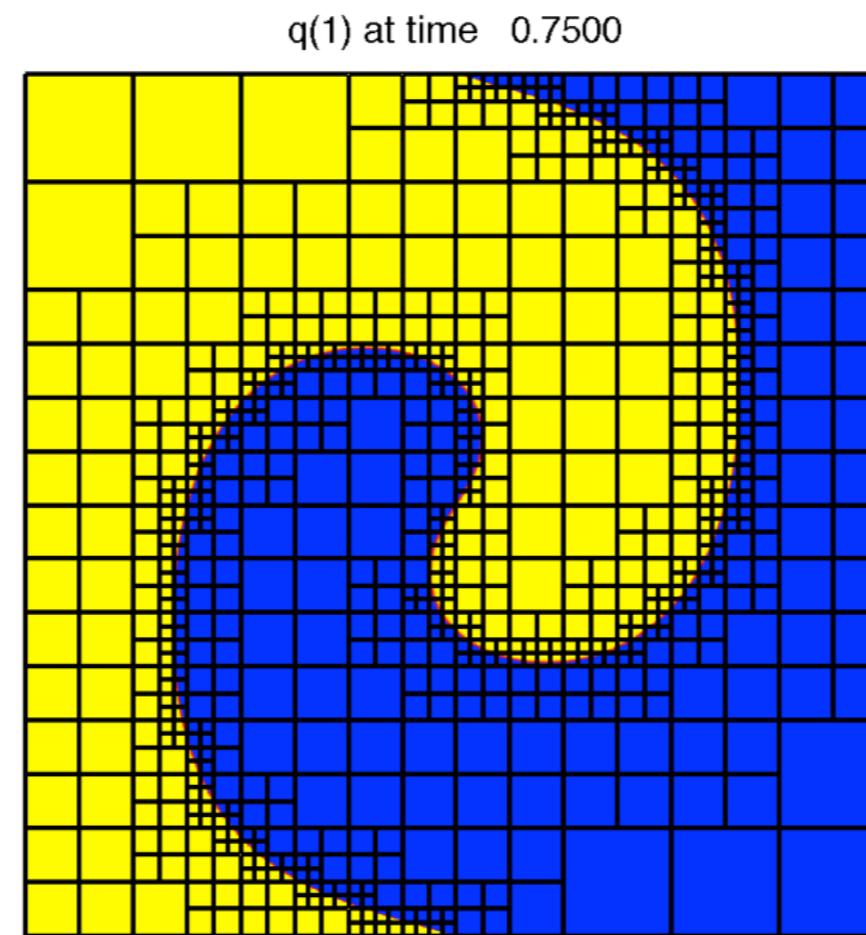
Use an existing parallel quad/octree library to do the grid management

- Store fixed sized non-overlapping grids as leaves in a tree

# ForestClaw : A hybrid approach to AMR

Use an existing parallel quad/octree library to do the grid management

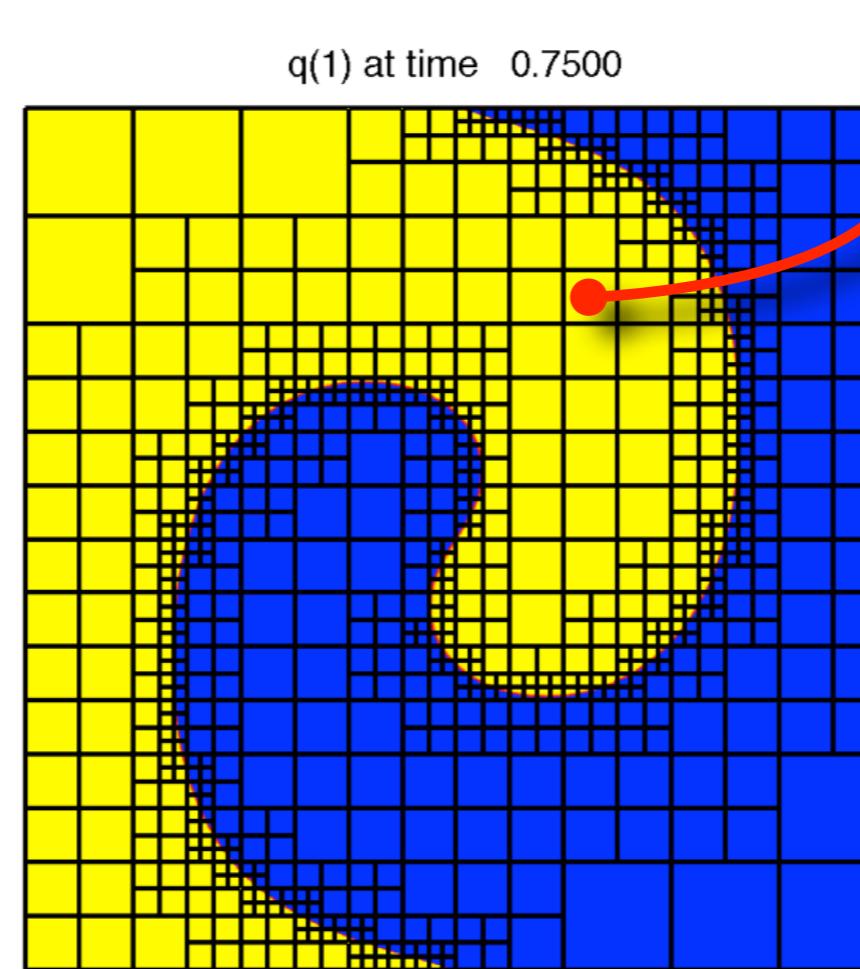
- Store fixed sized non-overlapping grids as leaves in a tree



# ForestClaw : A hybrid approach to AMR

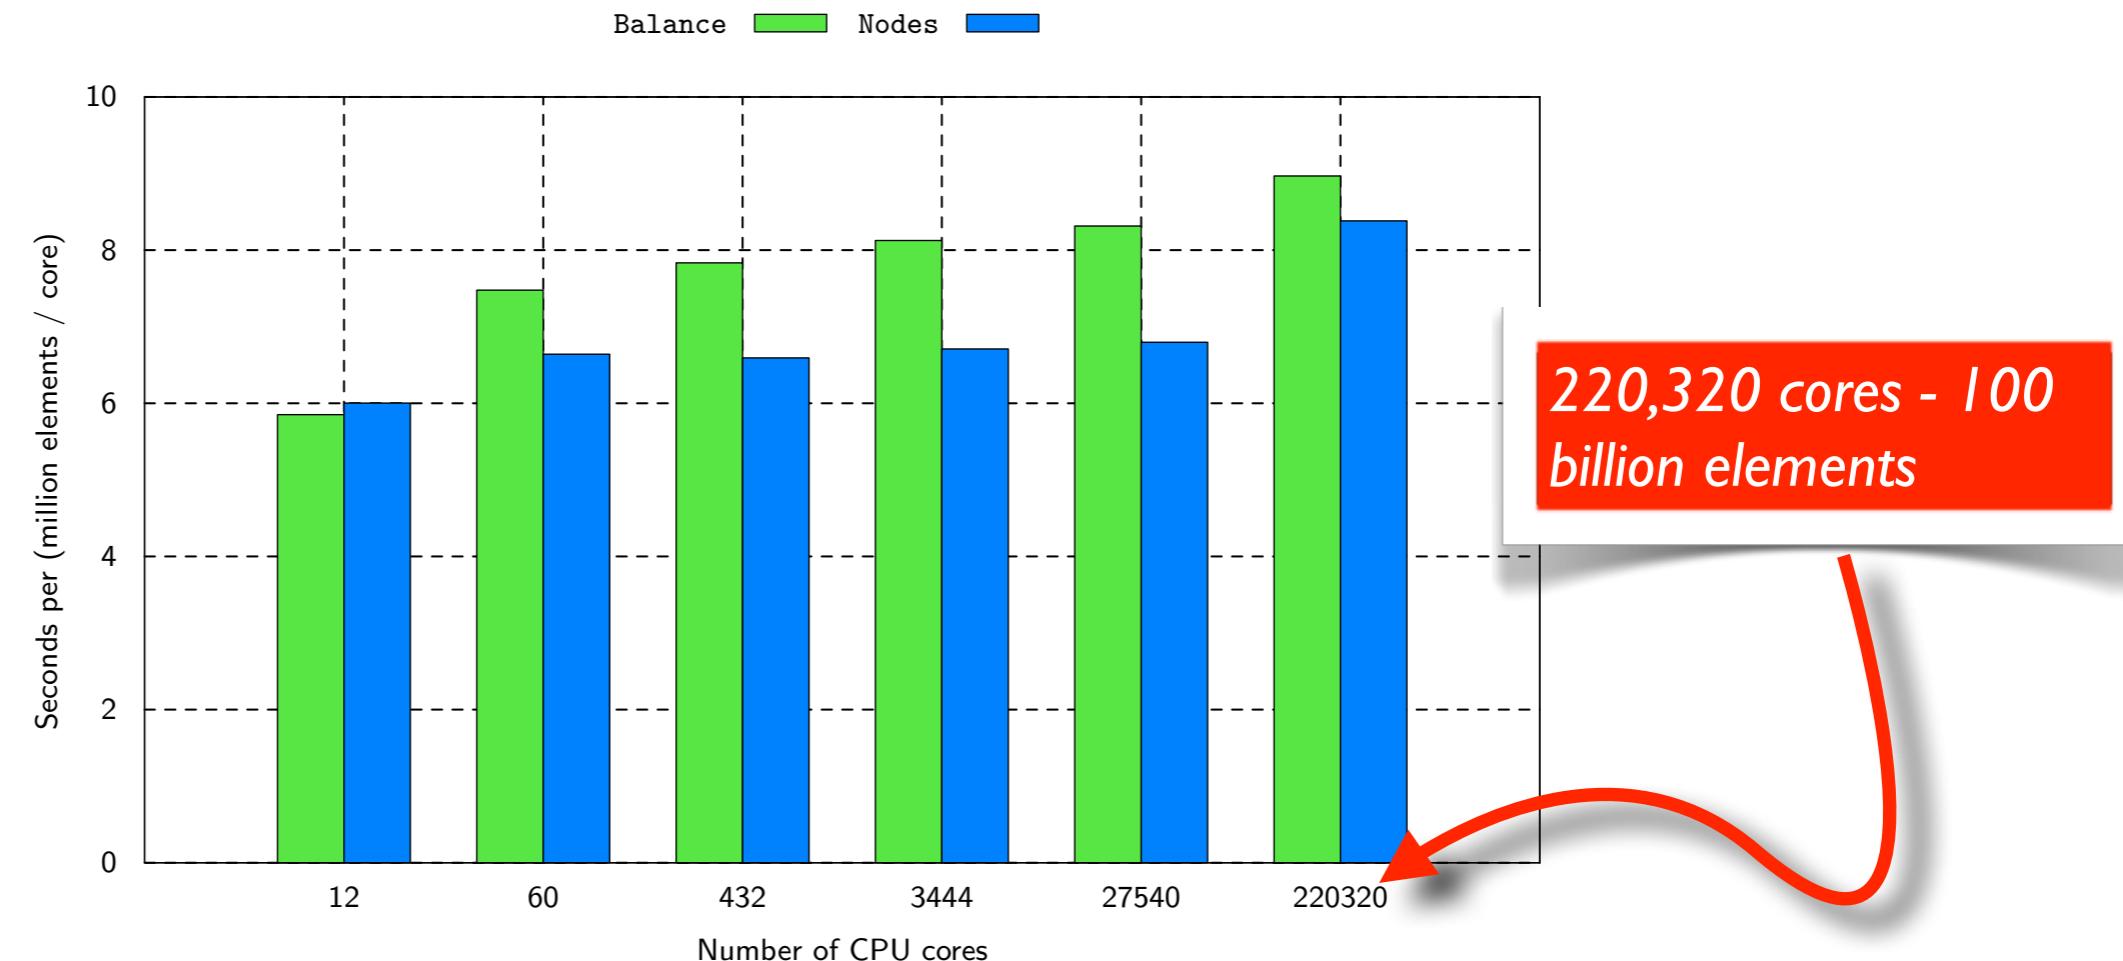
Use an existing parallel quad/octree library to do the grid management

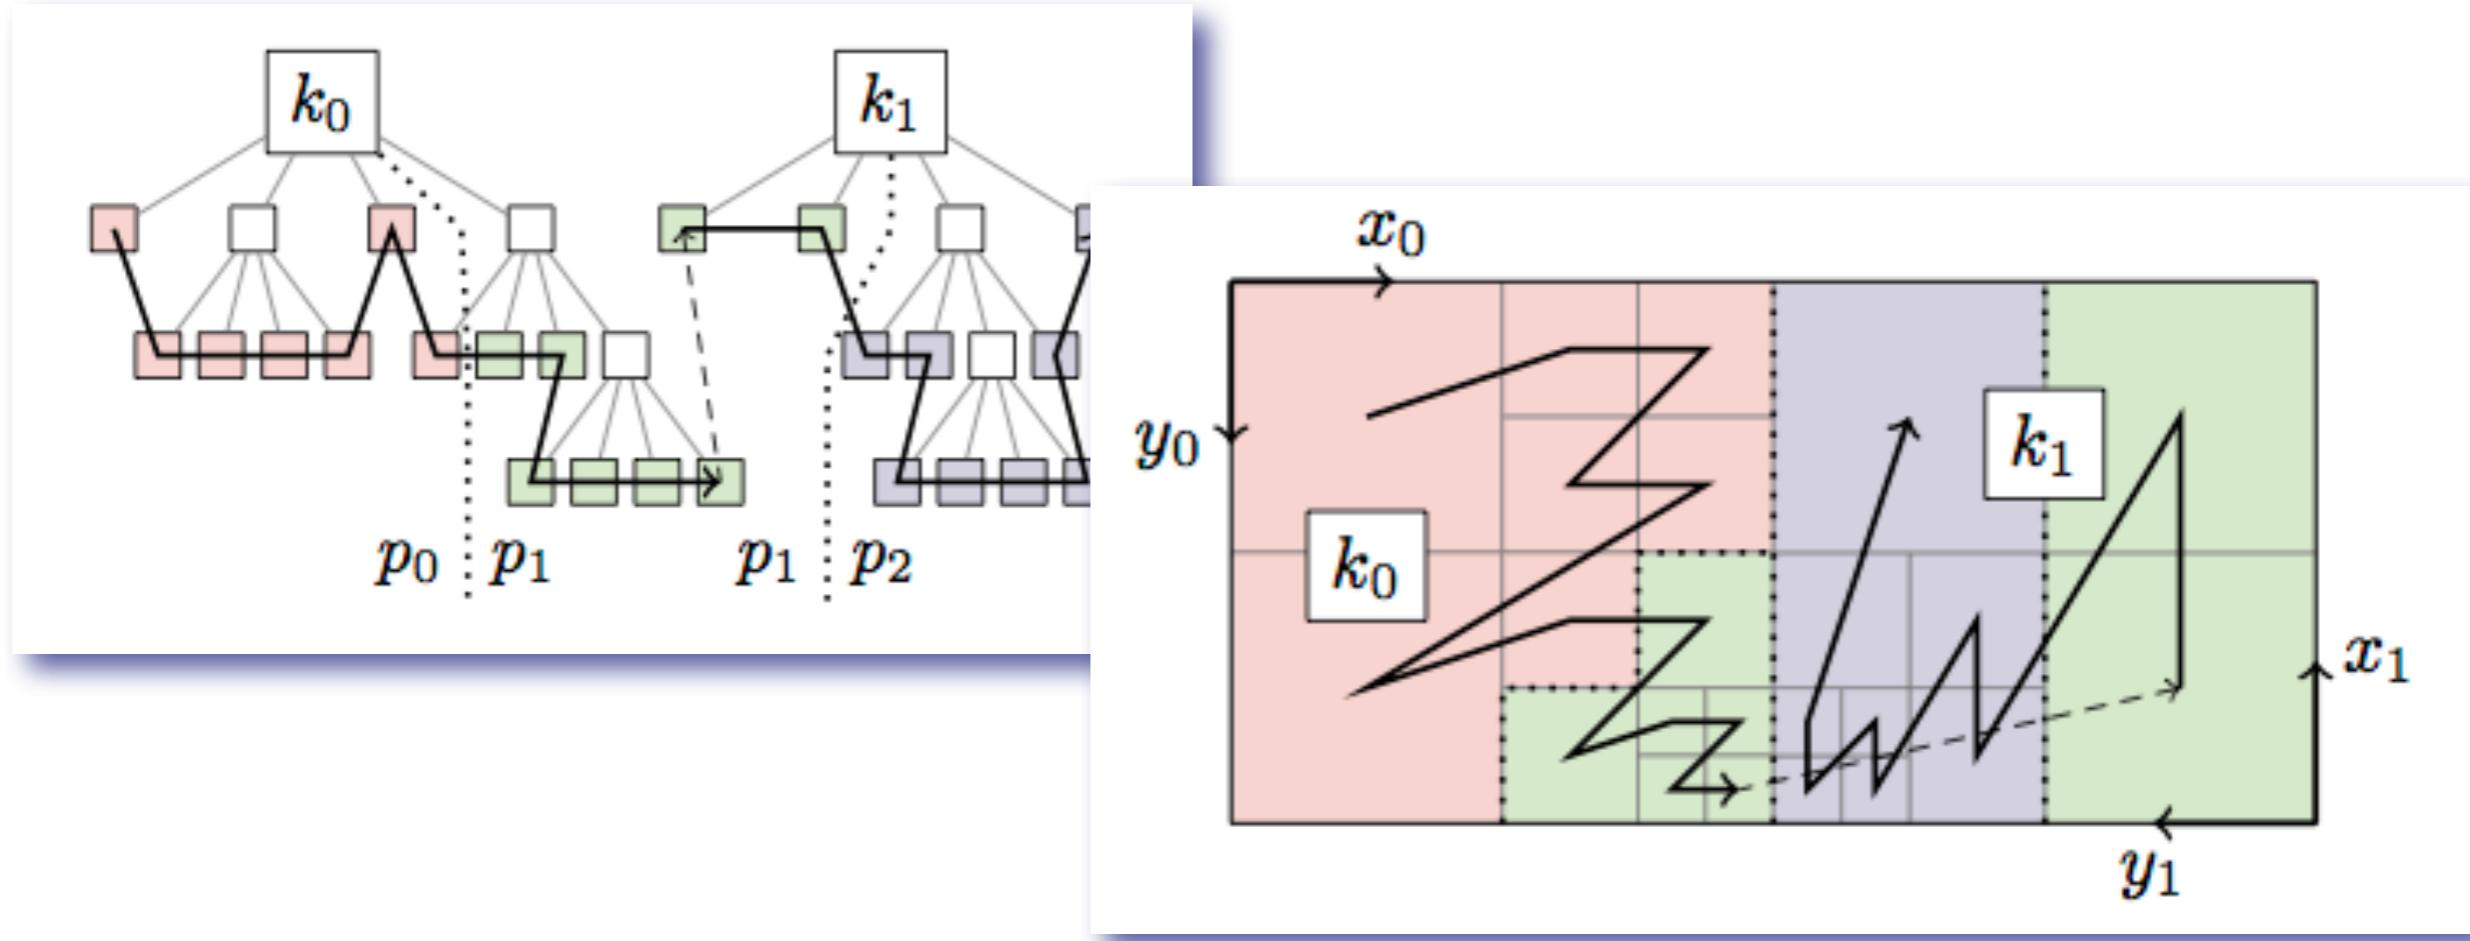
- Store fixed sized non-overlapping grids as leaves in a tree



*Tree structure based on refining quadrants, but each quadrant is itself a grid*

- Parallel, multi-block code for managing a forest of adaptive quad- or octrees.
- Highly scalable on realistic applications of interest
- Developed by Carsten Burstedde (Univ. of Bonn), with Wilcox, Ghattas and others

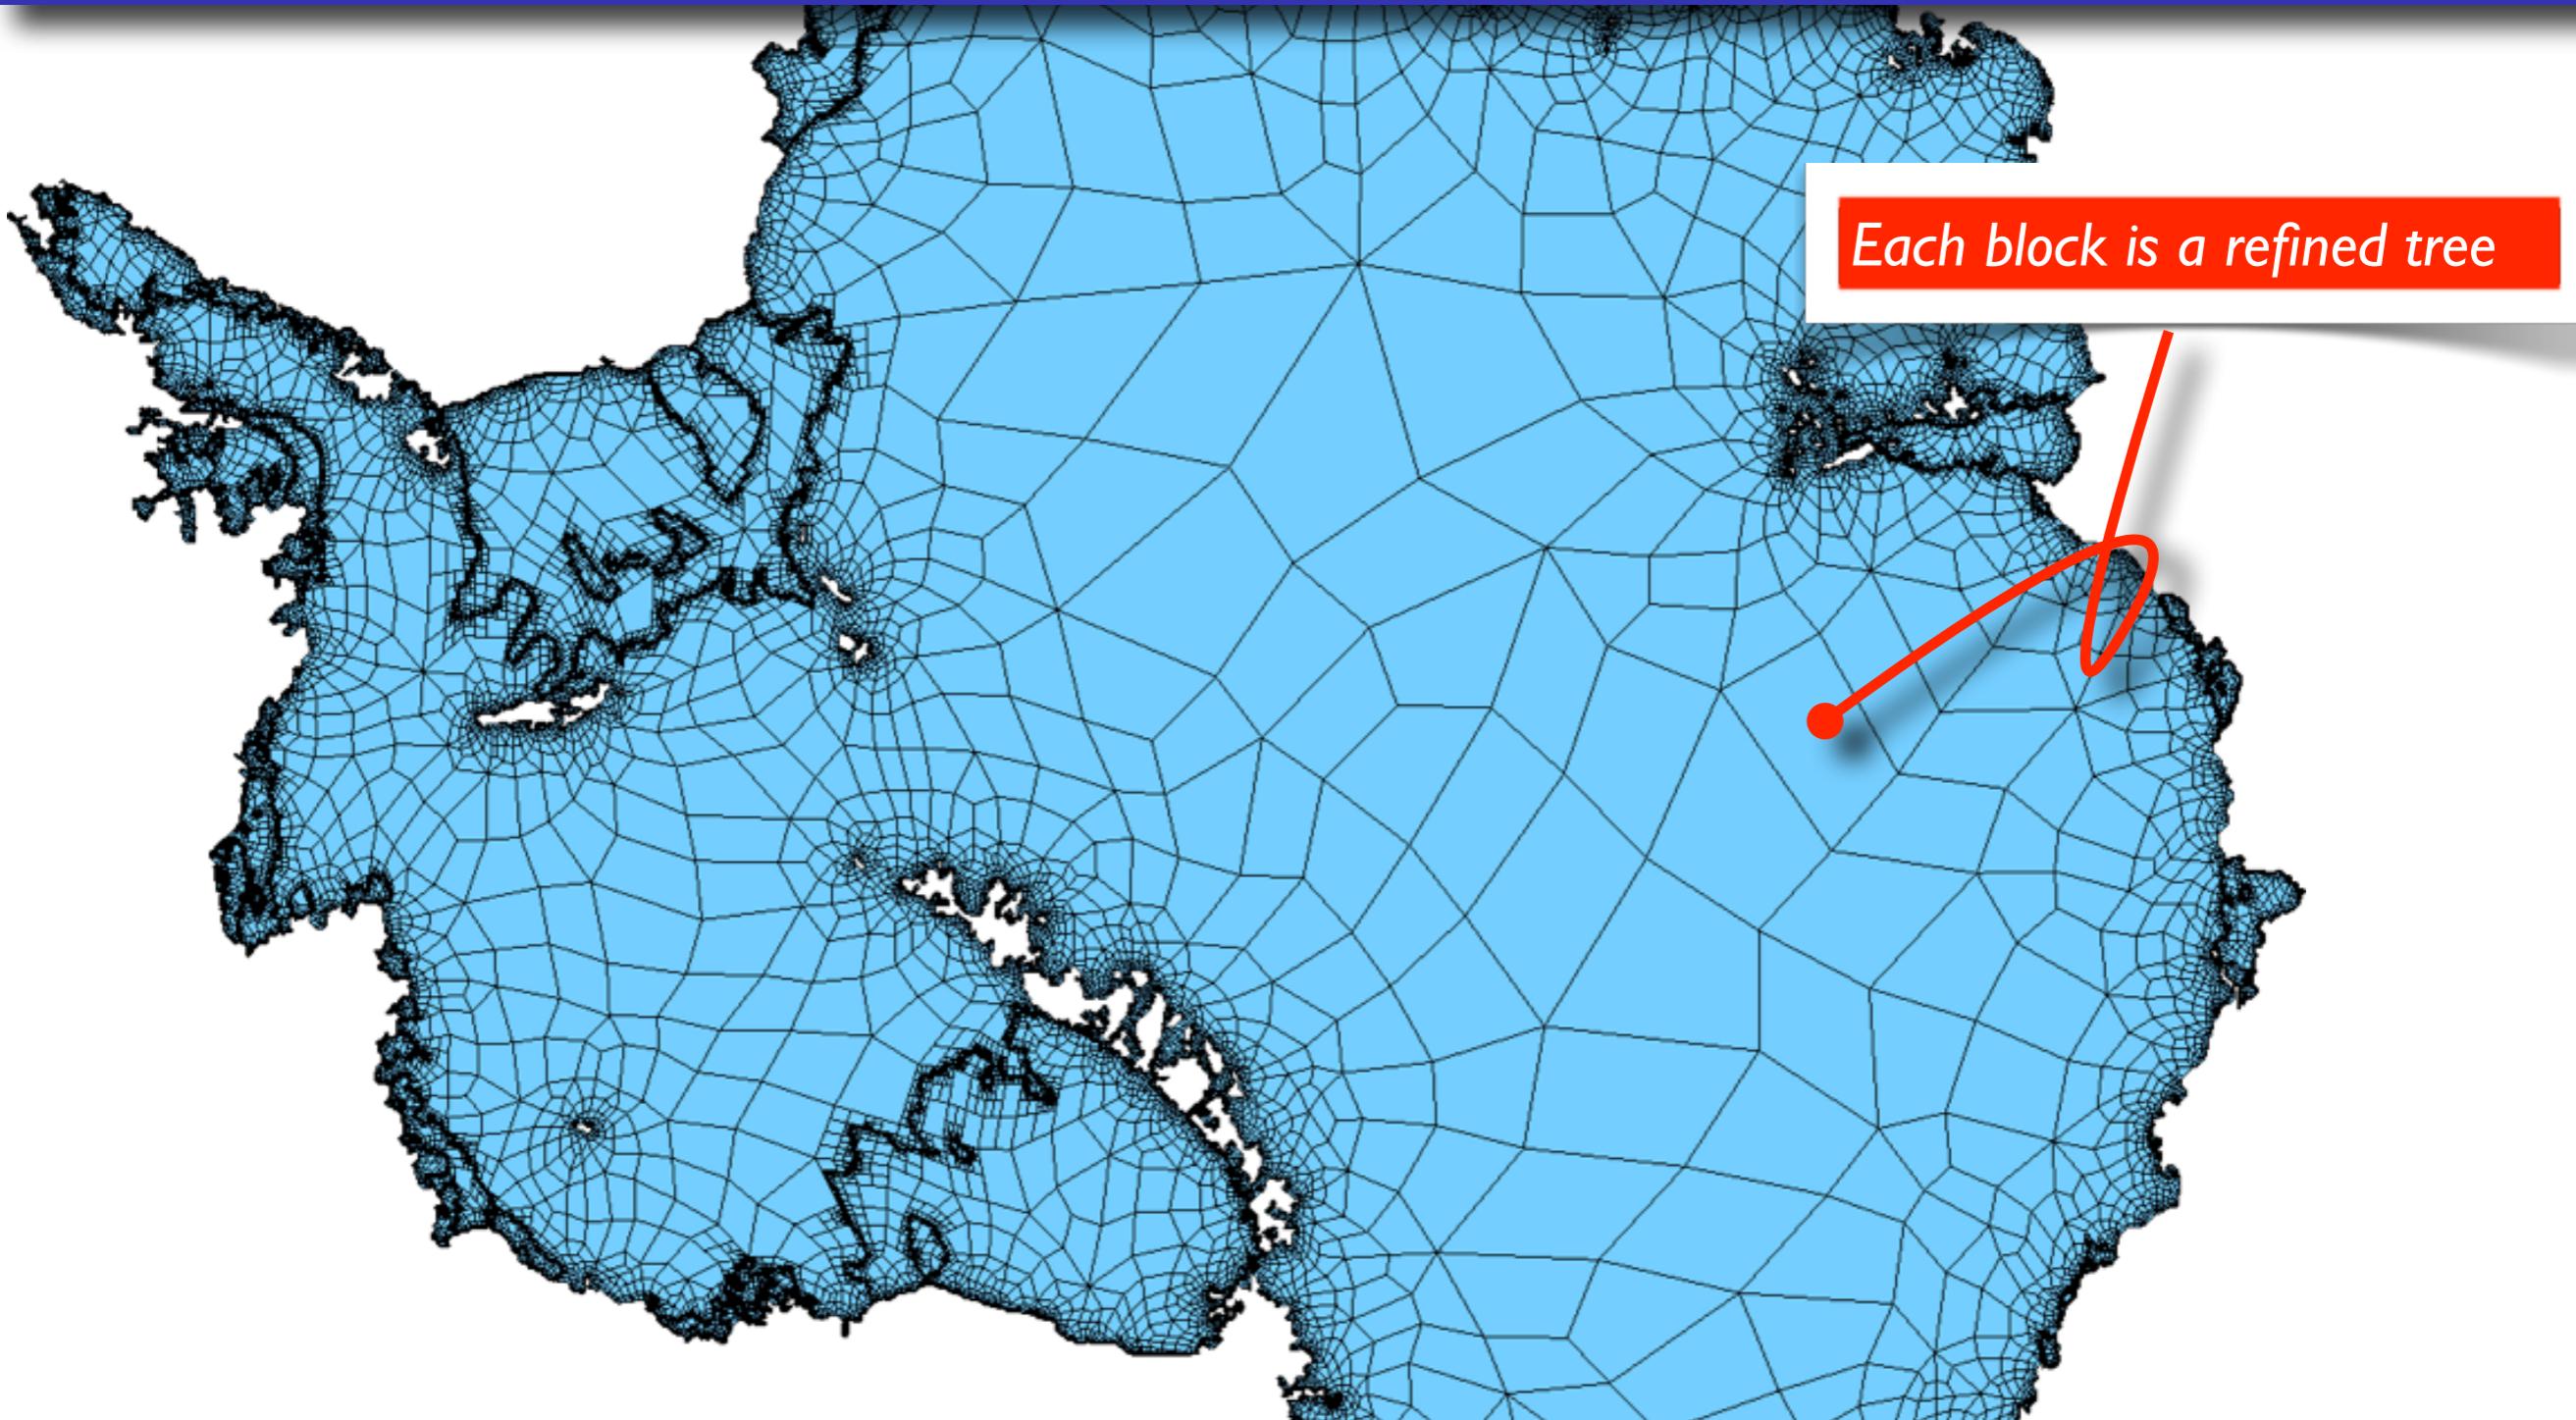




*High scalability is achieved while preserving data locality by using space-filling curves.*

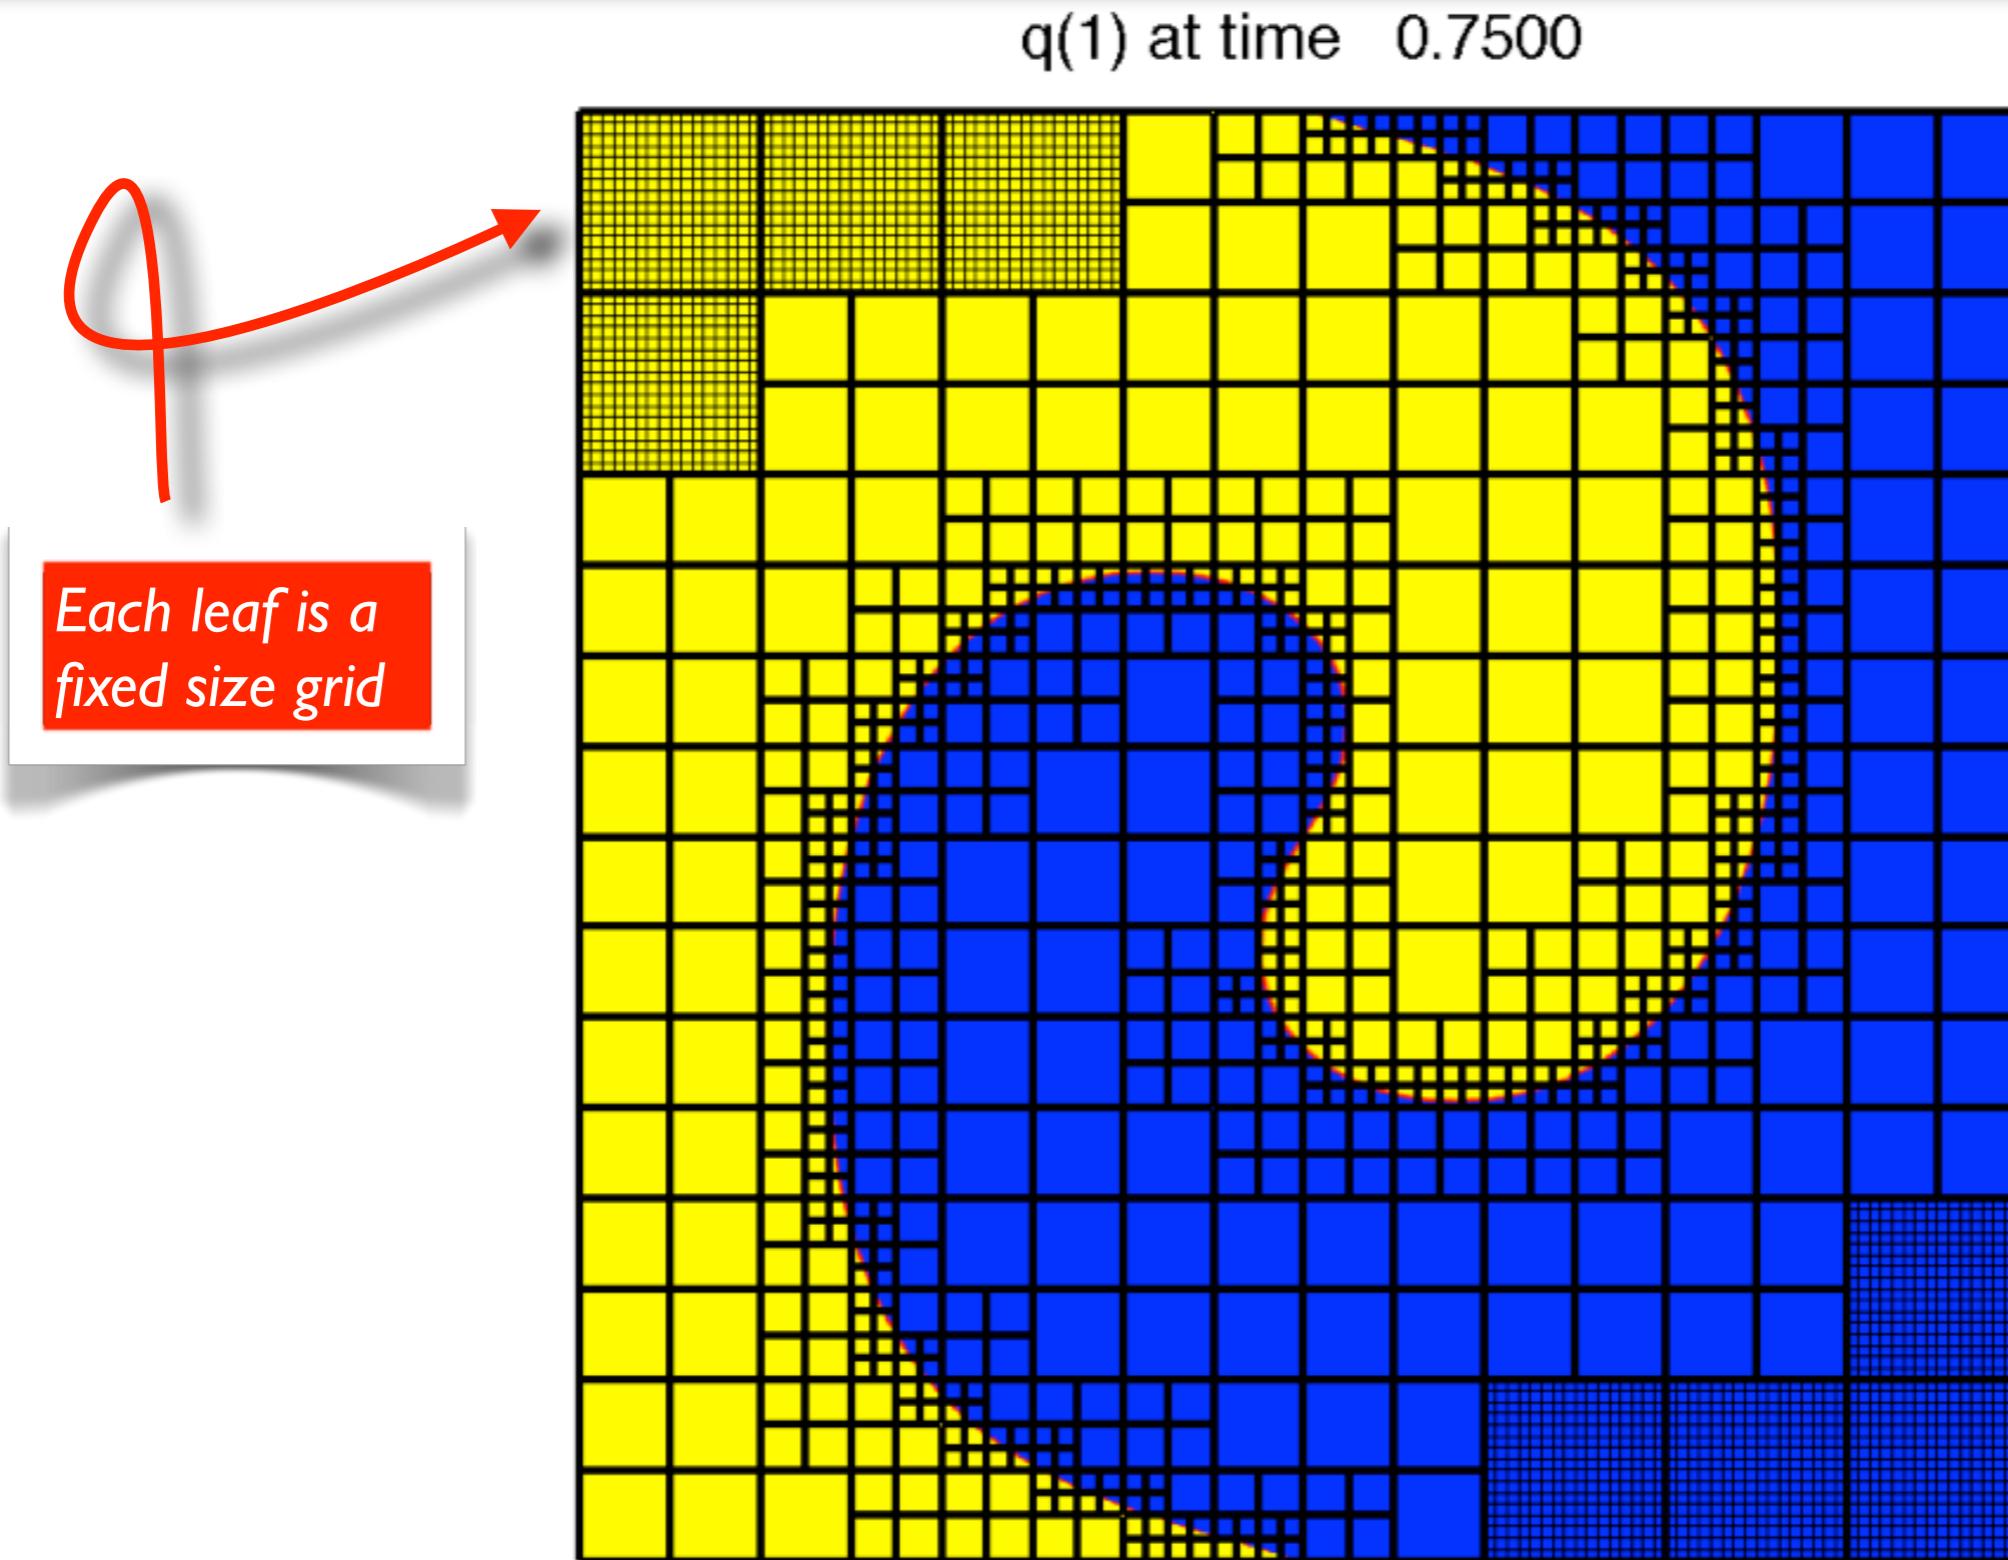
Carsten Burstedde, Lucas C. Wilcox, and Omar Ghattas, “p4est: Scalable Algorithms for Parallel Adaptive Mesh Refinement on Forests of Octrees”, SISC (2011)

# Multi-block support in p4est

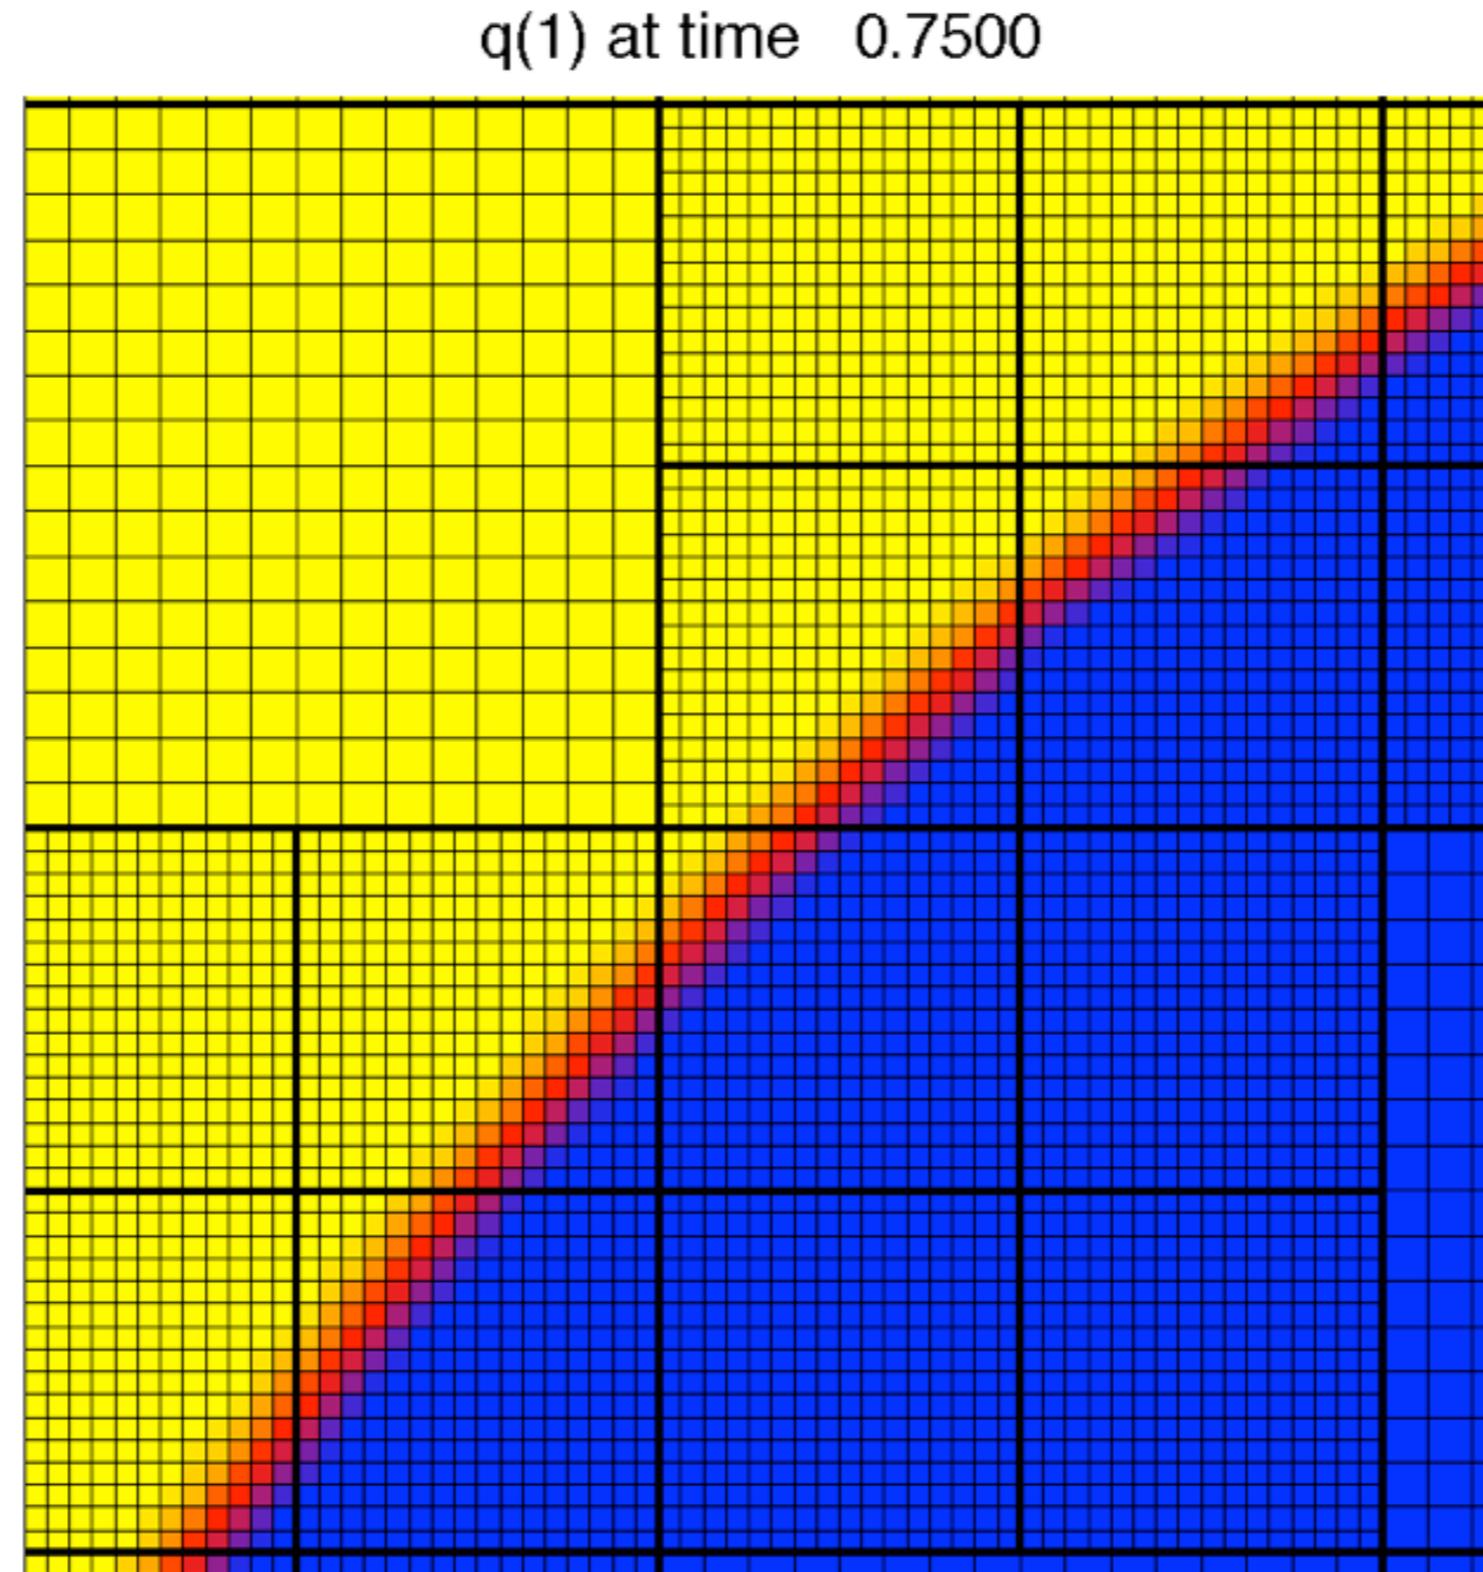


Antarctic ice sheet modeling (Tobin Isaac, C Burstedde)

# A hybrid approach to AMR

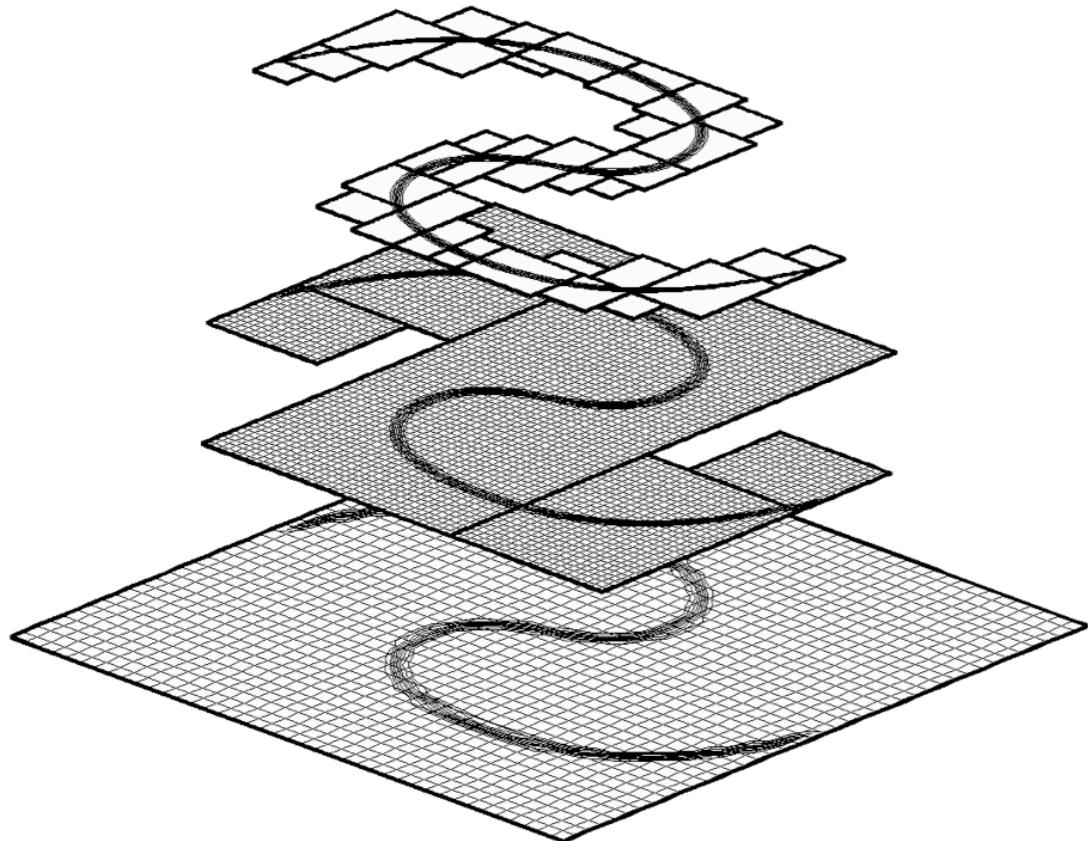


# A hybrid approach to AMR



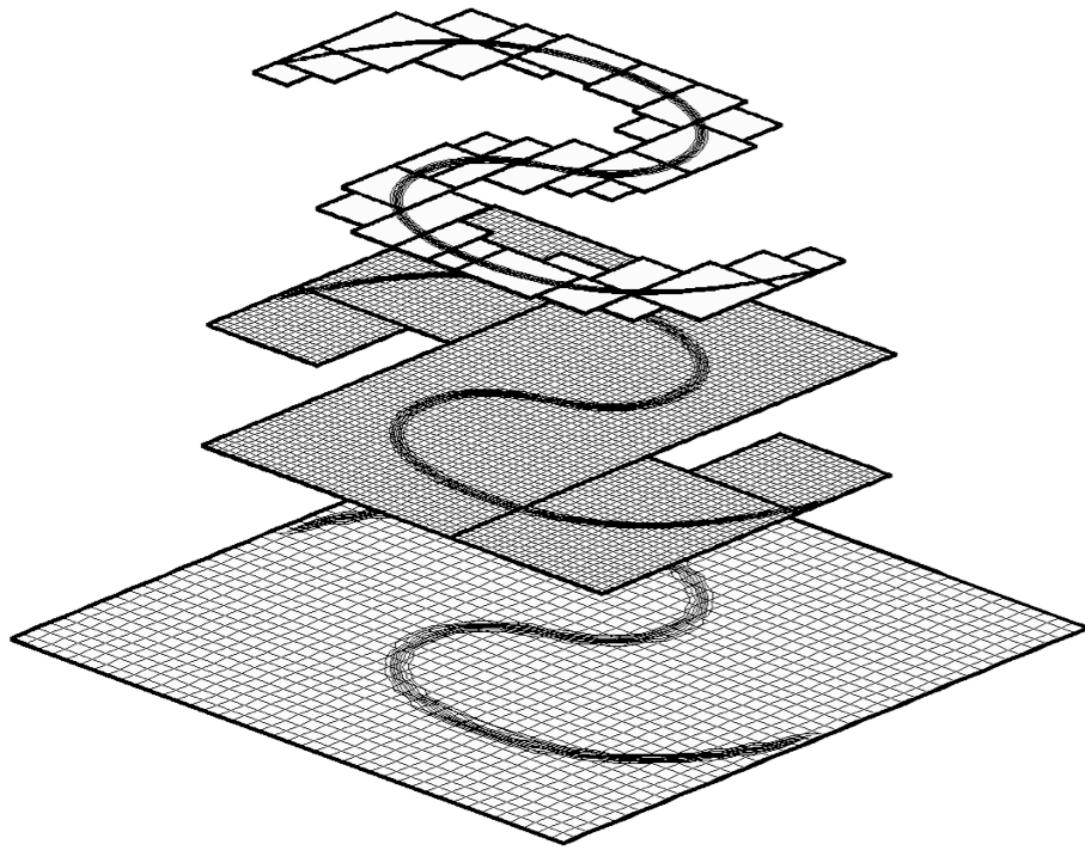
# A hybrid approach to AMR

# A hybrid approach to AMR

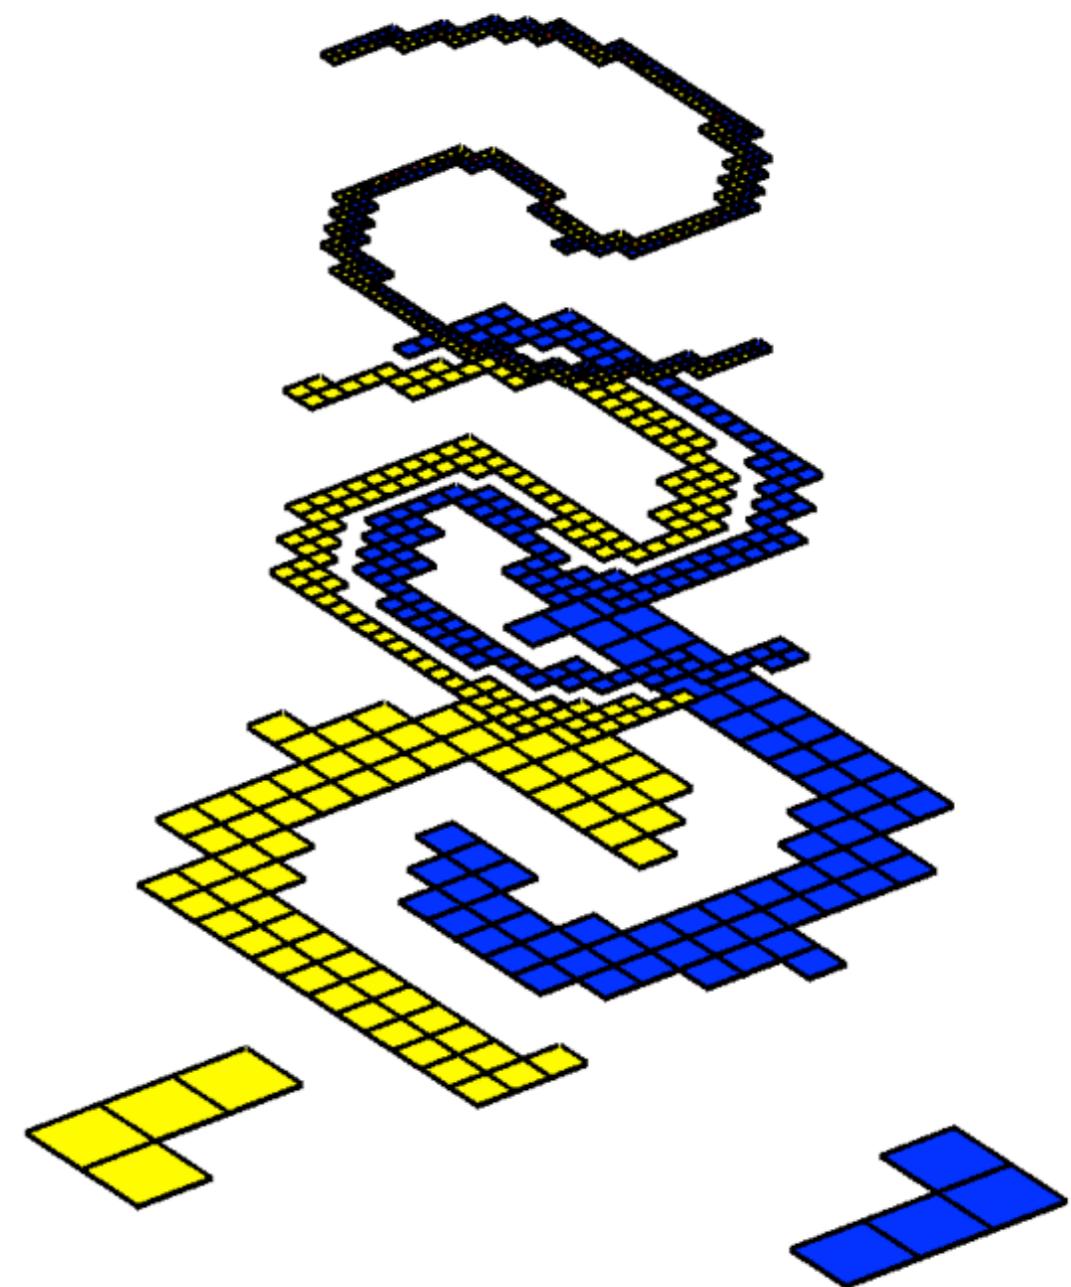


Berger-Oliger approach

# A hybrid approach to AMR

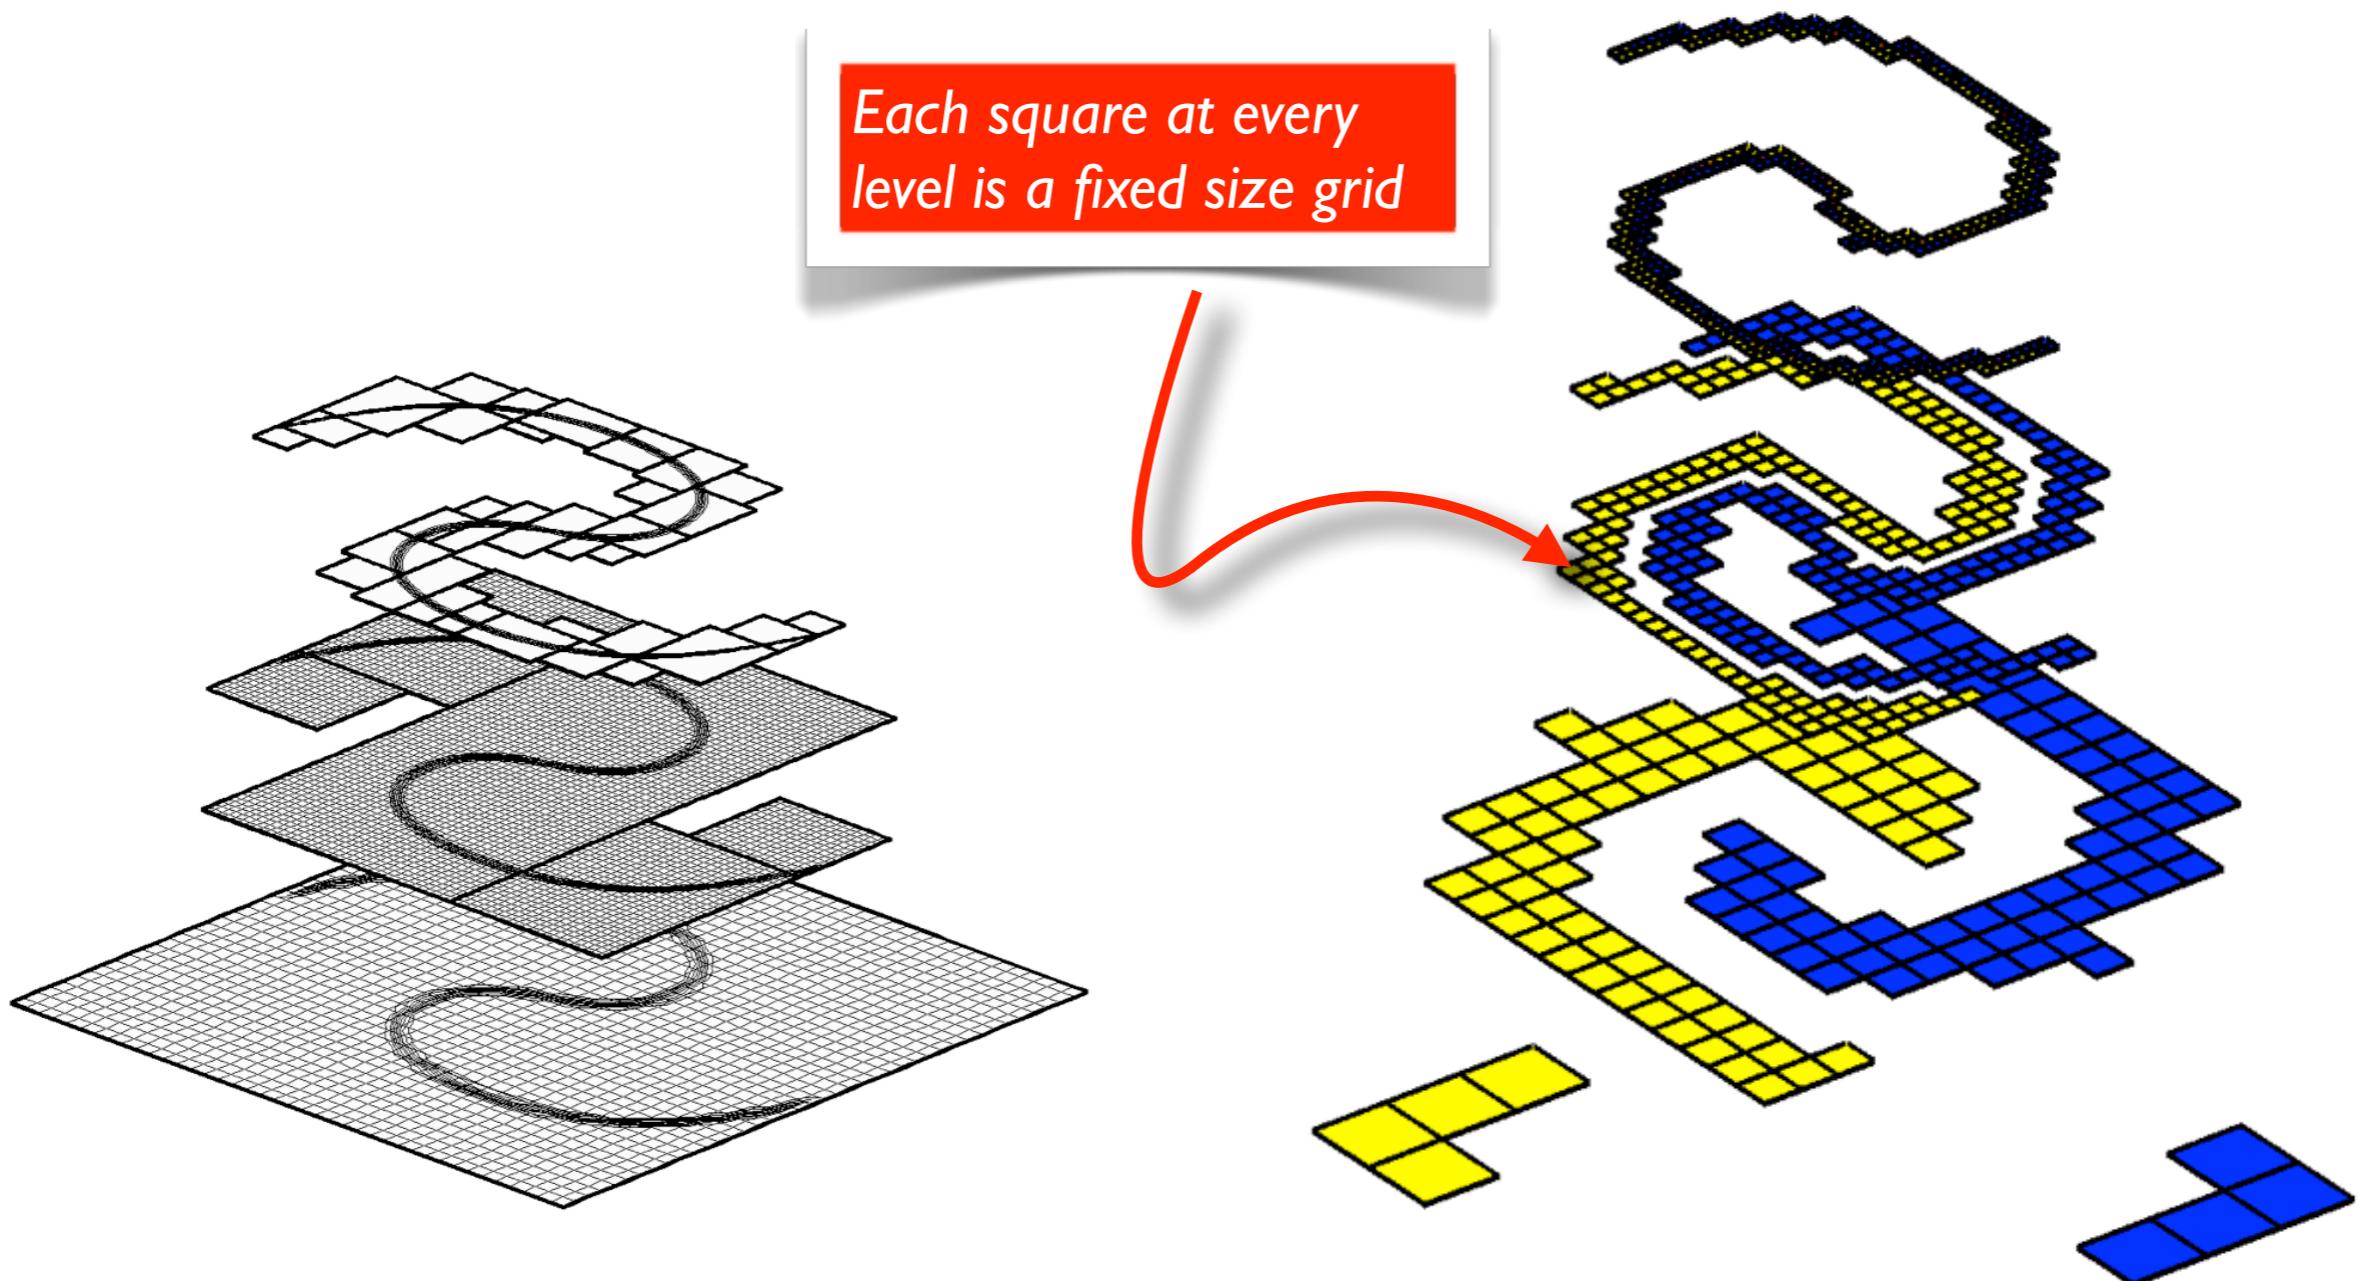


Berger-Oliger approach



ForestClaw

# A hybrid approach to AMR



Berger-Oliger approach

ForestClaw

# How is this easier for the developer?

# How is this easier for the developer?

- Management of “boxes” (leaves) containing grids handled behind the scenes,

# How is this easier for the developer?

- Management of “boxes” (leaves) containing grids handled behind the scenes,
- Refinement based on tagging leaves is handled trivially : one grid is refined into four grids.

# How is this easier for the developer?

- Management of “boxes” (leaves) containing grids handled behind the scenes,
- Refinement based on tagging leaves is handled trivially : one grid is refined into four grids.
- Interpolation/averaging and ghost cell exchange all done using few prescribed patterns of grid intersections - no need to store indices of locations of fine grid in the coarse grid

# How is this easier for the developer?

- Management of “boxes” (leaves) containing grids handled behind the scenes,
- Refinement based on tagging leaves is handled trivially : one grid is refined into four grids.
- Interpolation/averaging and ghost cell exchange all done using few prescribed patterns of grid intersections - no need to store indices of locations of fine grid in the coarse grid
- Parallelization is handled at the tree level, in a separate library,

# How is this easier for the developer?

- Management of “boxes” (leaves) containing grids handled behind the scenes,
- Refinement based on tagging leaves is handled trivially : one grid is refined into four grids.
- Interpolation/averaging and ghost cell exchange all done using few prescribed patterns of grid intersections - no need to store indices of locations of fine grid in the coarse grid
- Parallelization is handled at the tree level, in a separate library,
- Multi-block domains are handled automatically

# How is this easier for the user?

# How is this easier for the user?

- Ghost cell exchanges only occur at the edges/faces of grids, not in the interior of coarse grids,

# How is this easier for the user?

- Ghost cell exchanges only occur at the edges/faces of grids, not in the interior of coarse grids,
- Predictable pattern for grid neighbors - one grid at the same level, two at a finer level, or “half” at a coarser level (in 2d),

# How is this easier for the user?

- Ghost cell exchanges only occur at the edges/faces of grids, not in the interior of coarse grids,
- Predictable pattern for grid neighbors - one grid at the same level, two at a finer level, or “half” at a coarser level (in 2d),
- Composite view of the grid hierarchy provides a cleaner mental map of what the solution looks like (no overlapping grids)

# How is this easier for the user?

- Ghost cell exchanges only occur at the edges/faces of grids, not in the interior of coarse grids,
- Predictable pattern for grid neighbors - one grid at the same level, two at a finer level, or “half” at a coarser level (in 2d),
- Composite view of the grid hierarchy provides a cleaner mental map of what the solution looks like (no overlapping grids)
- Generic multi-block interface allows user to construct complicated domains

# How is this easier for the user?

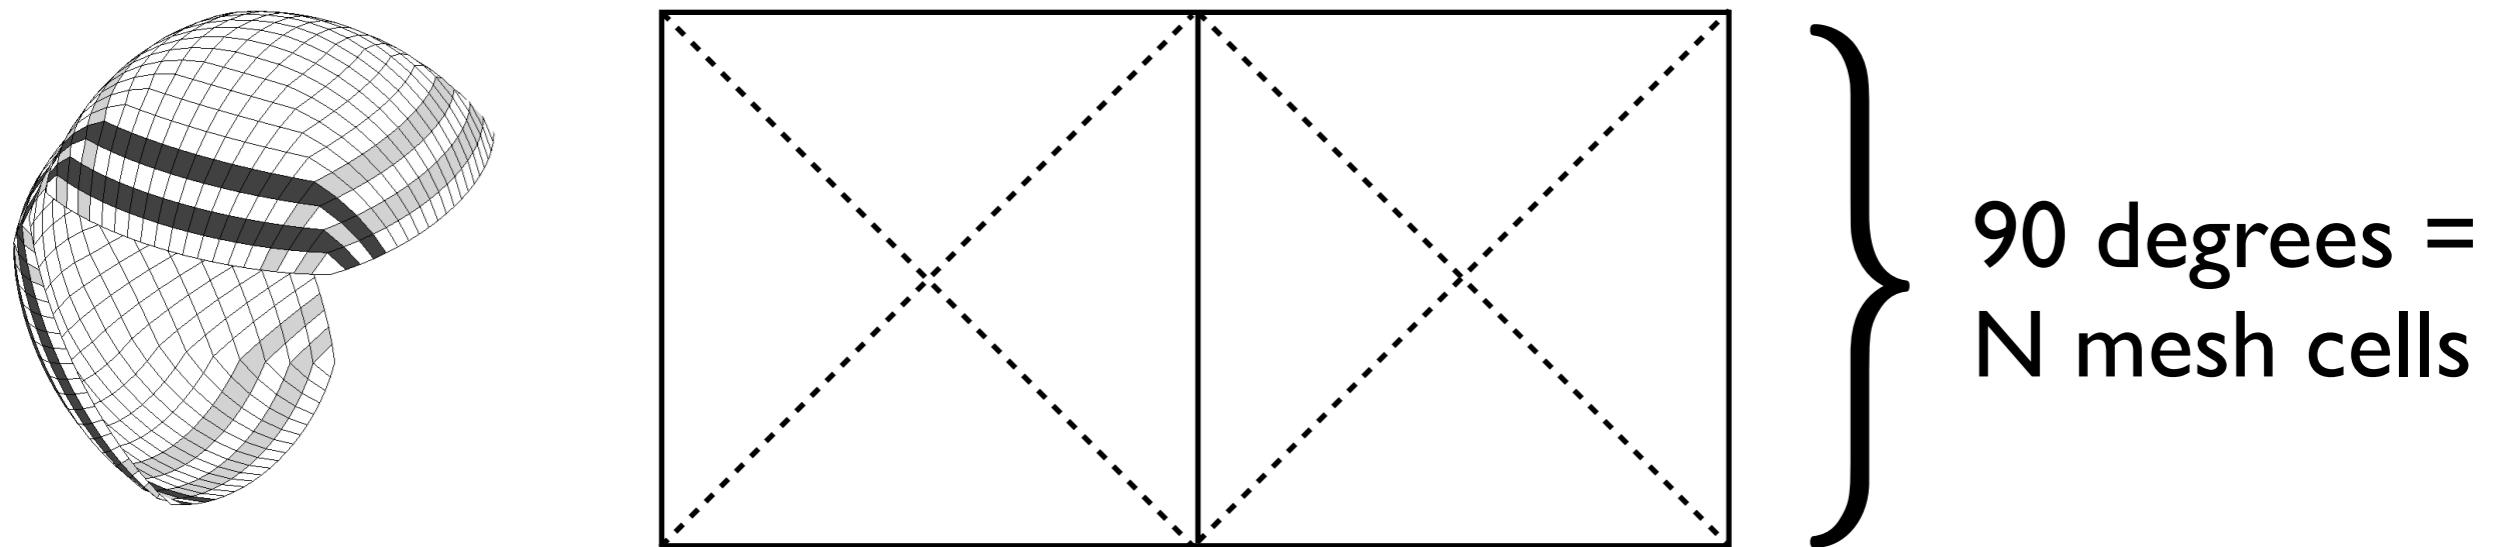
- Ghost cell exchanges only occur at the edges/faces of grids, not in the interior of coarse grids,
- Predictable pattern for grid neighbors - one grid at the same level, two at a finer level, or “half” at a coarser level (in 2d),
- Composite view of the grid hierarchy provides a cleaner mental map of what the solution looks like (no overlapping grids)
- Generic multi-block interface allows user to construct complicated domains
- Has the aesthetic appeal of a quad/octree refinement with the superior performance of the Berger-Oliger approach to AMR.

# Other AMR approaches using quad/octrees

- “Building Cubes Method” (Sasaki, Akahito, Yamazaki, ...)
- Parallel adaptive methods for weather prediction C. Jablonowski, Oehmke, Stout and others
- NIRVANA (U. Ziegler)
- Racoon II (J. Dreher)
- PARAMESH (NASA,/Drexel, MacNeice, Olson)
- Block-structured AMR codes (Chombo, Boxlib, AMRClaw, SAMRAI, AMROC, ...) could probably be run with fixed size grids and prescribed refinement regions

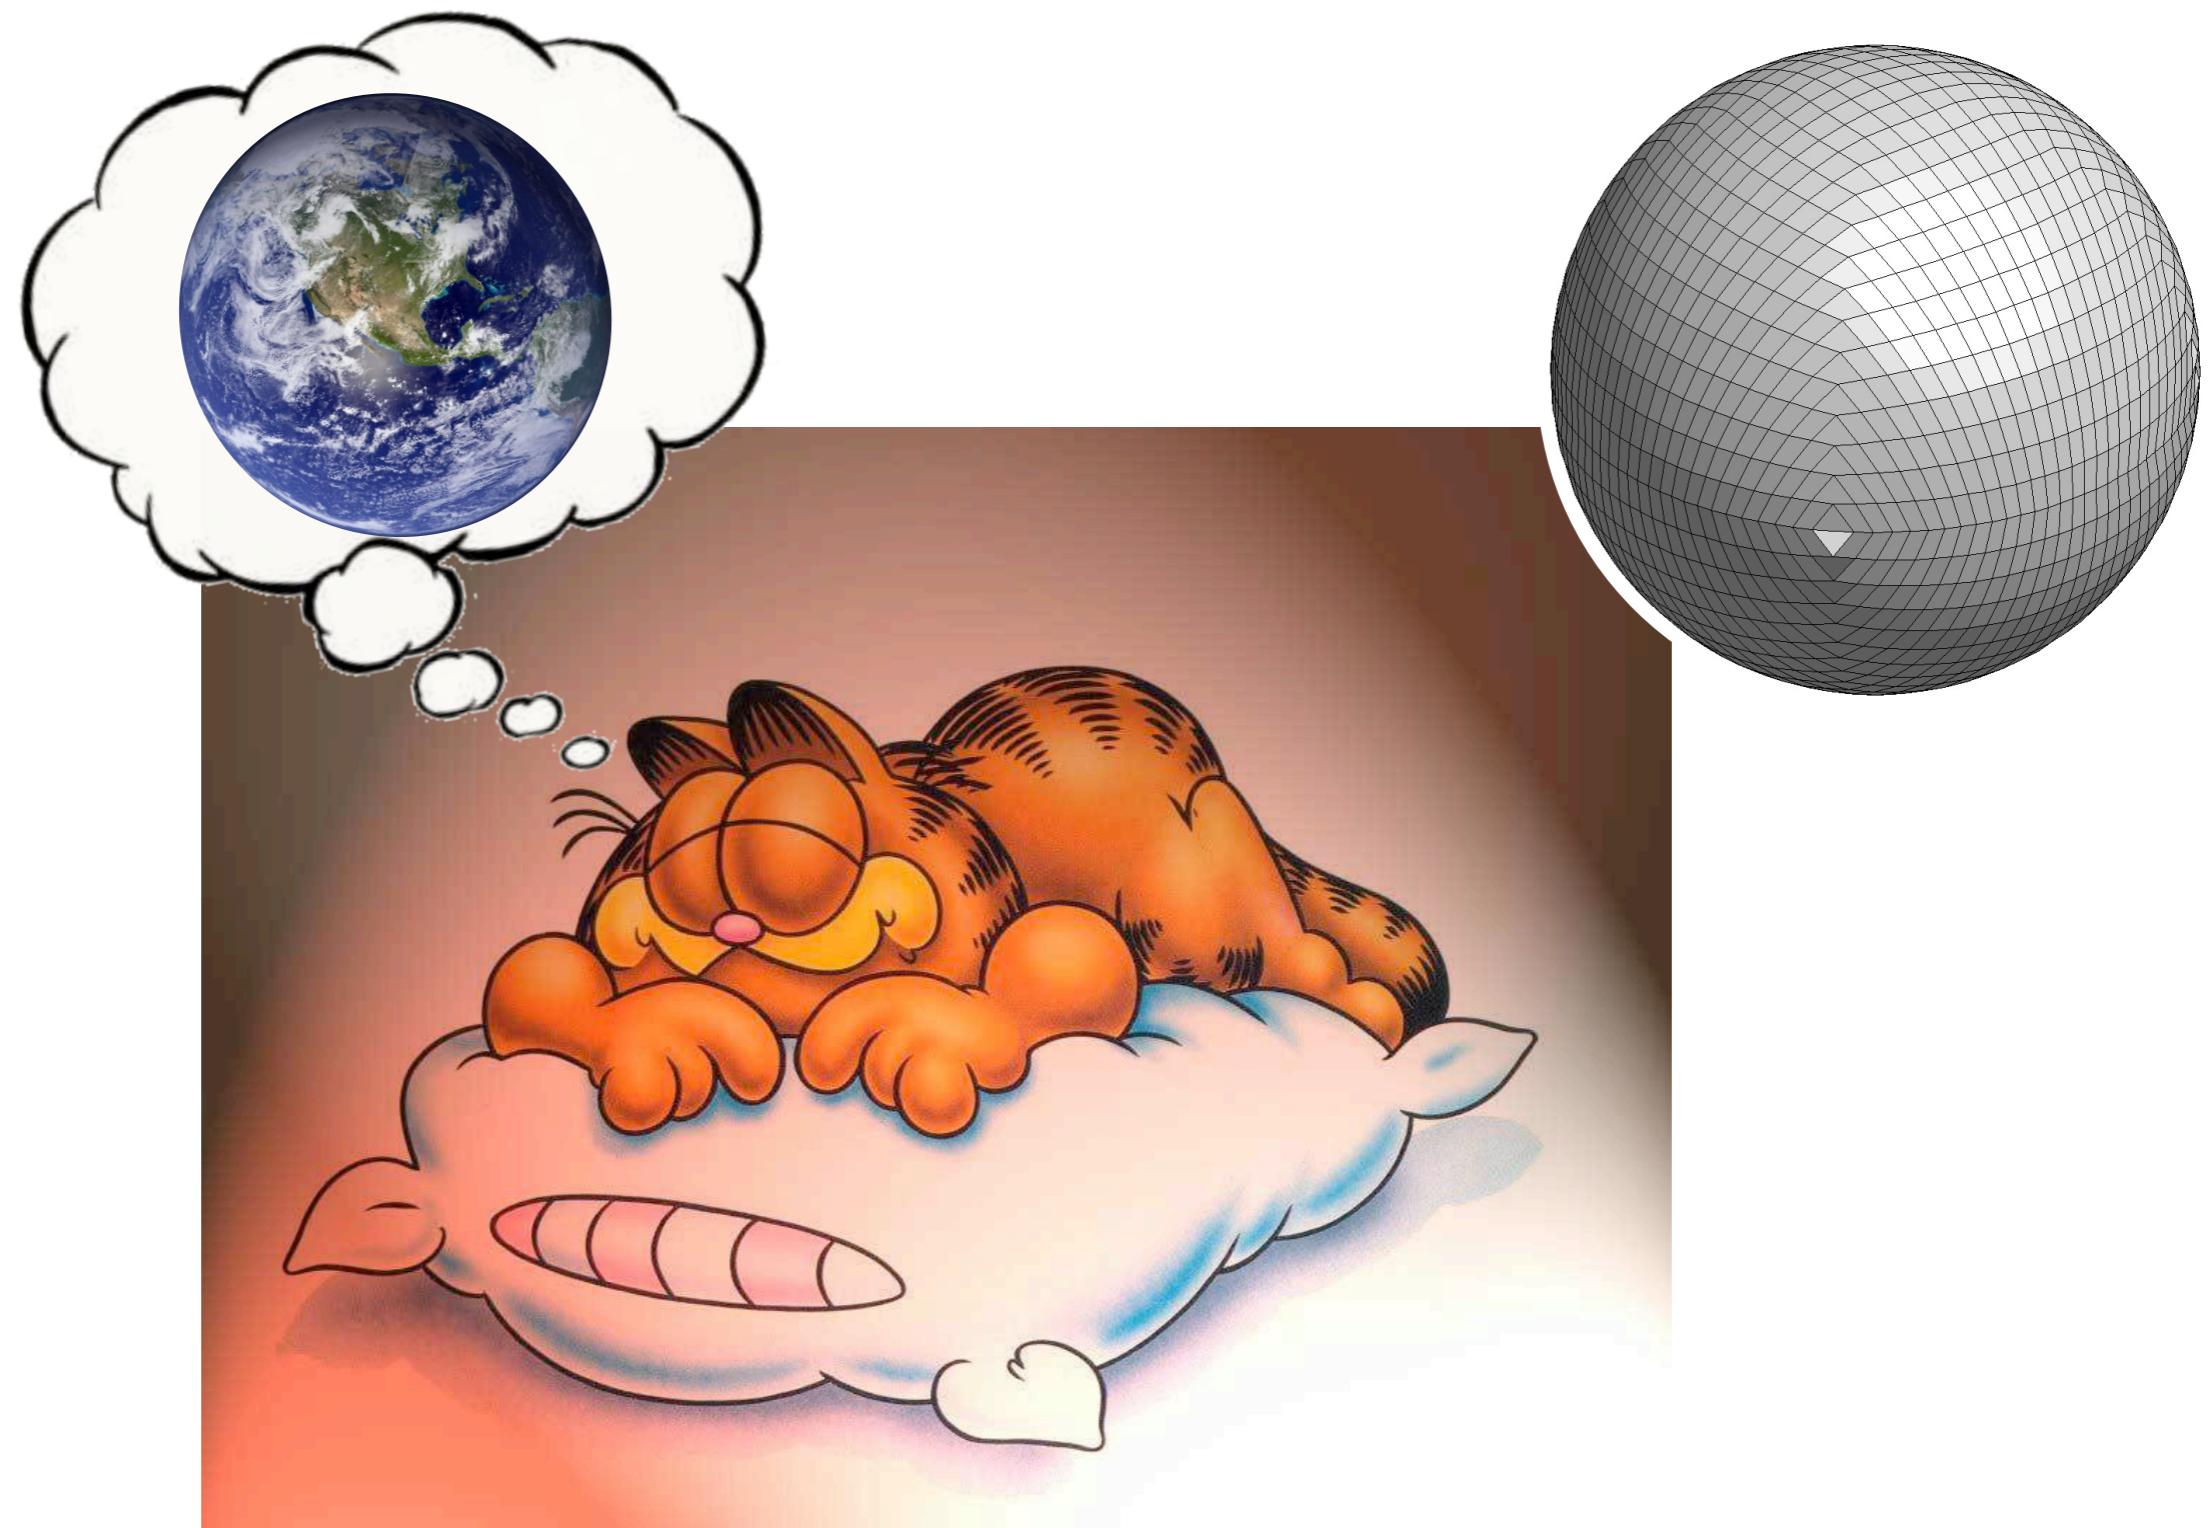
# A two-patch sphere grid?

- Our sphere grid is like the cubed-sphere grid, but with two patches
- We will refer to the grid resolution by the number of grid cells on a patch edge, which is approximately 90 degrees.

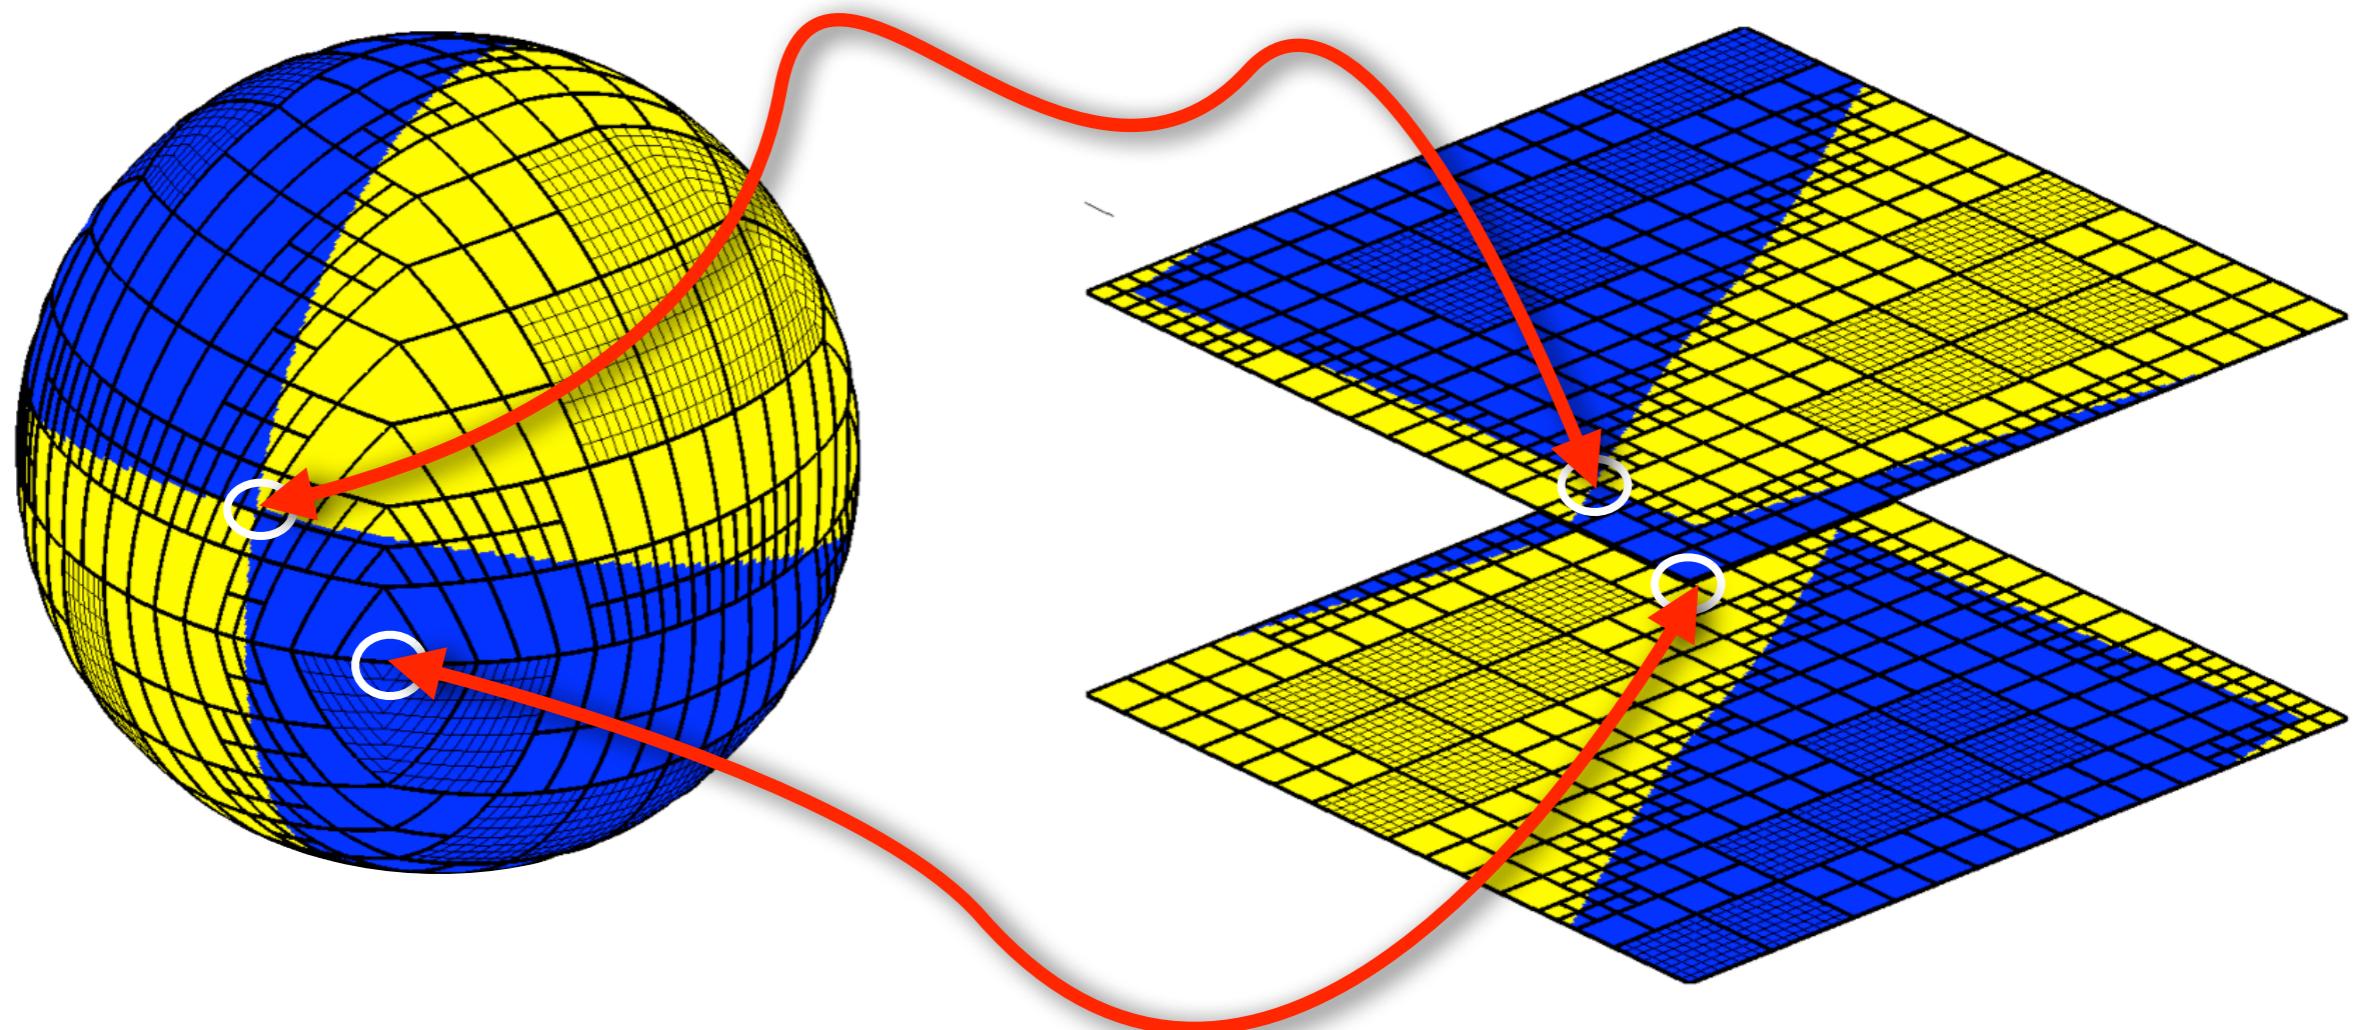


*Dashed and solid lines are discontinuities in the mapping*

# A Pillow grid?



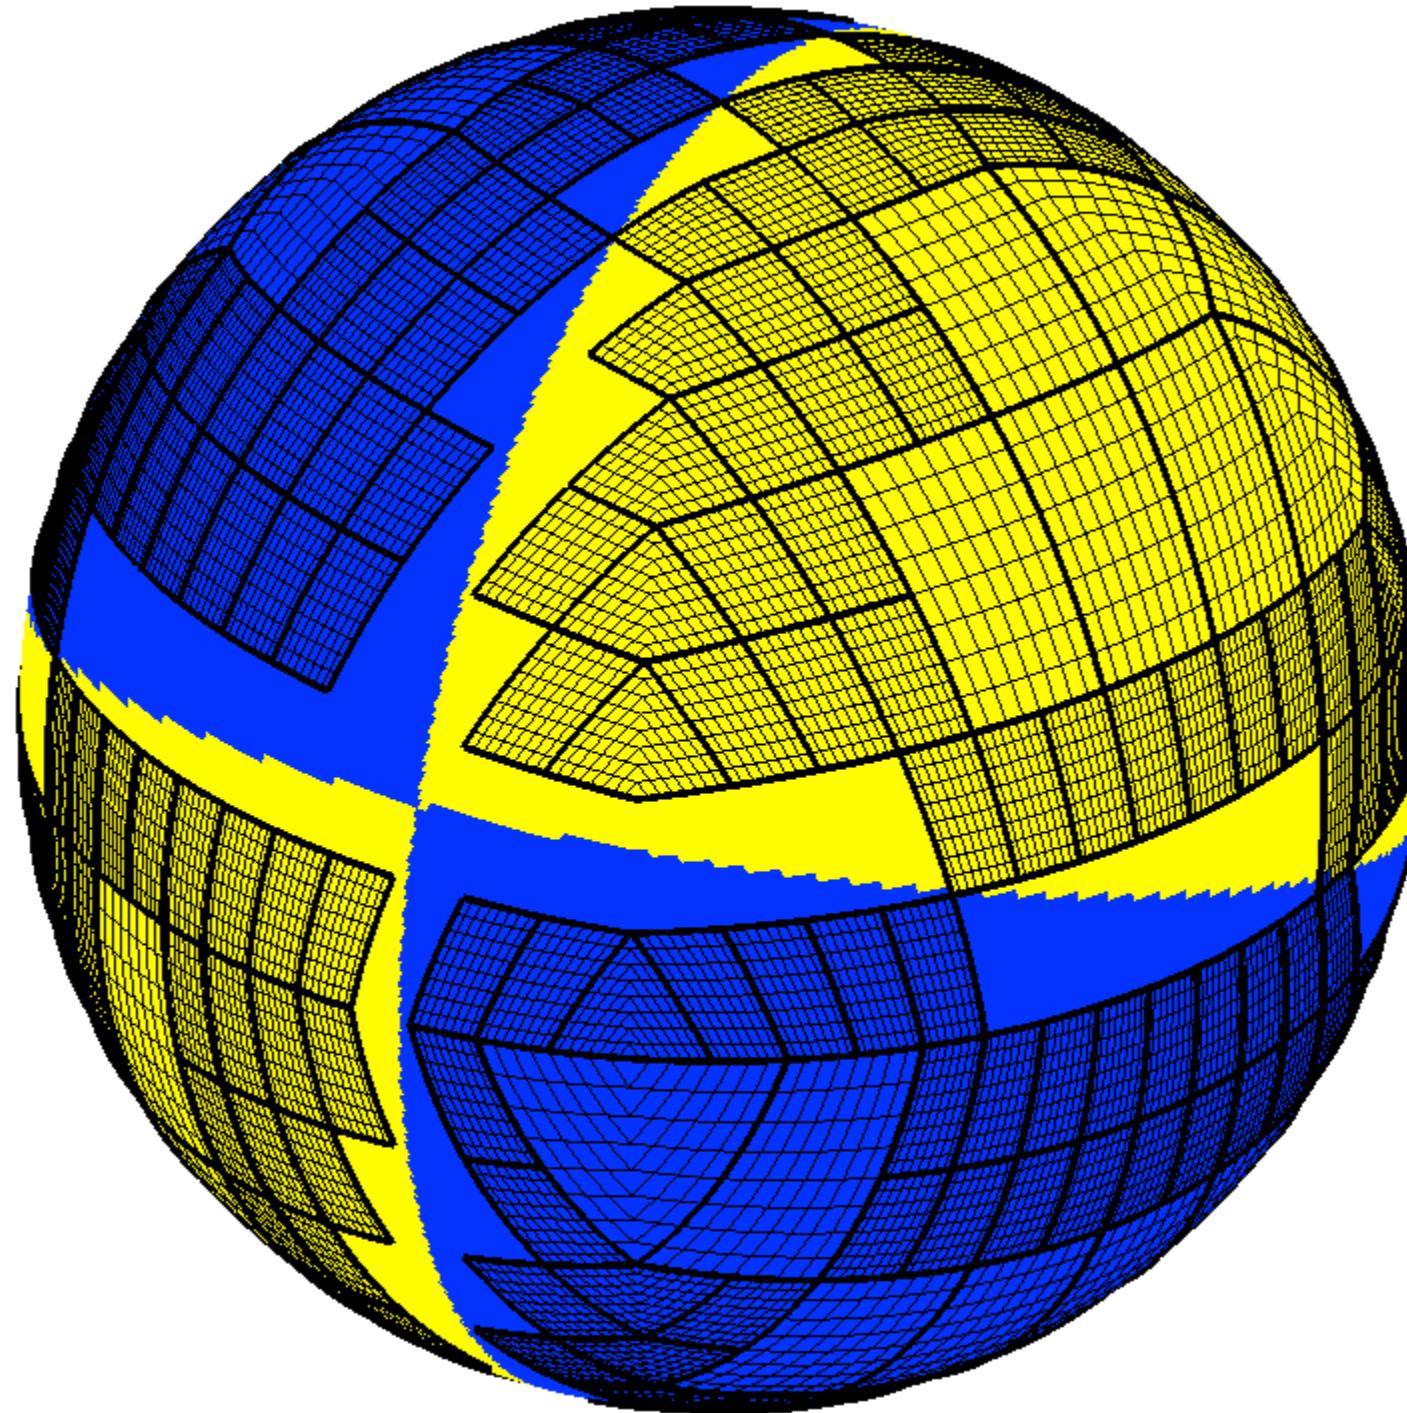
# Scalar transport on the sphere



Scalar advection using finite volume wave propagation  
algorithms (ClawPACK, R. J. LeVeque)

# Scalar advection

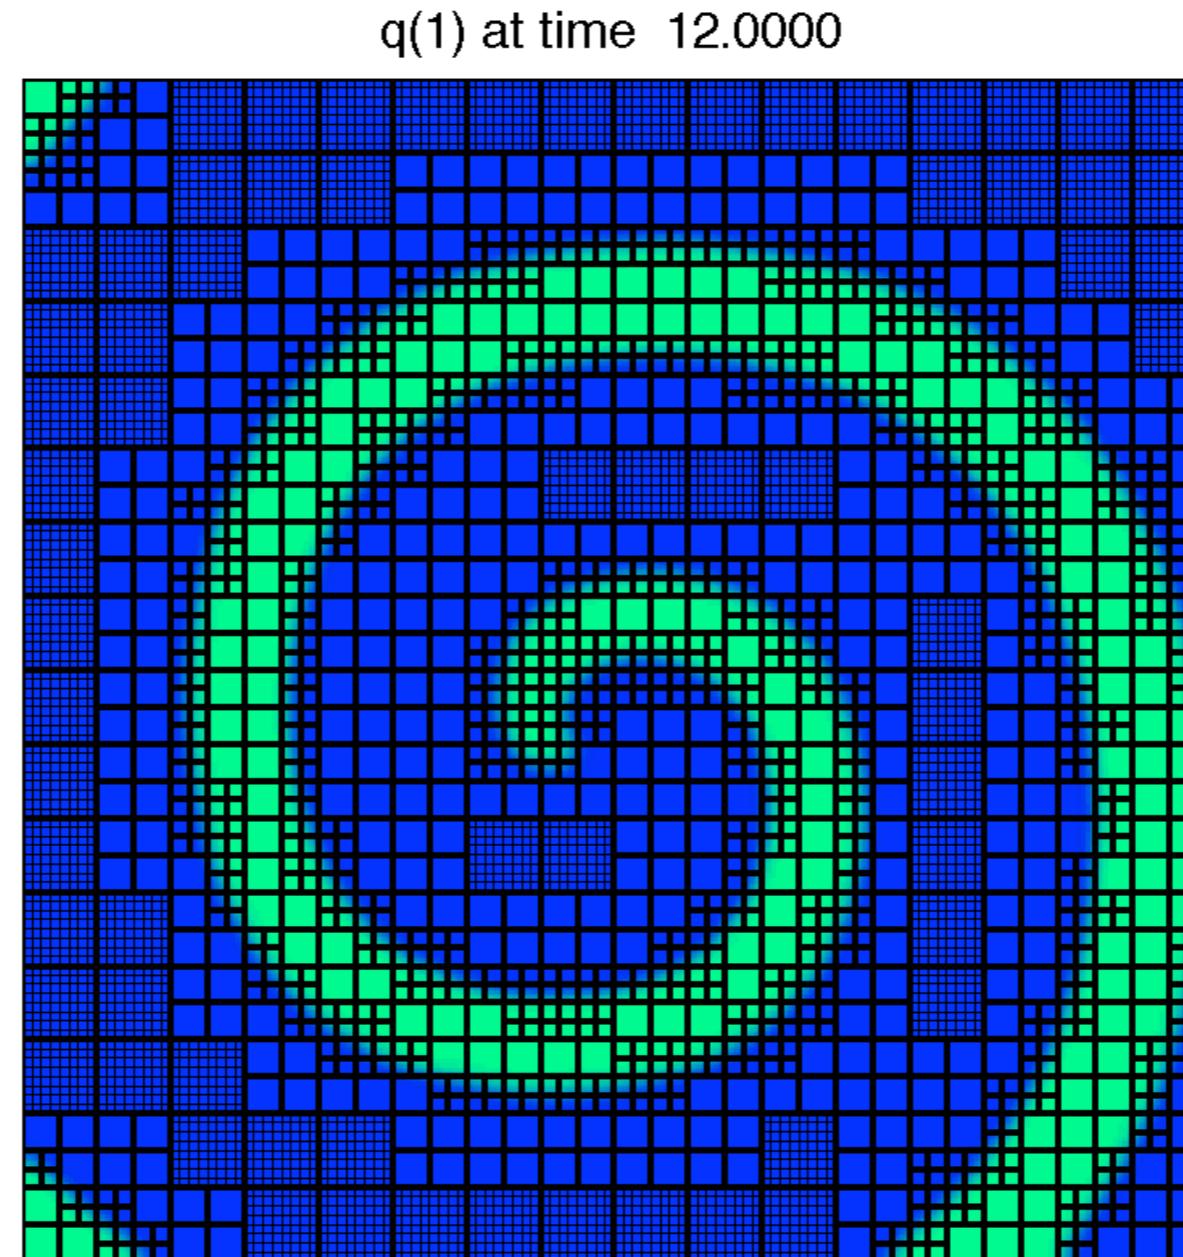
# Scalar advection



# Spiral waves (Barkely model)

Reaction-diffusion using an explicit Runge-Kutta Chebychev  
(RKC) time stepping

# Spiral waves (Barkely model)



Reaction-diffusion using an explicit Runge-Kutta Chebychev  
(RKC) time stepping

# Near future challenges

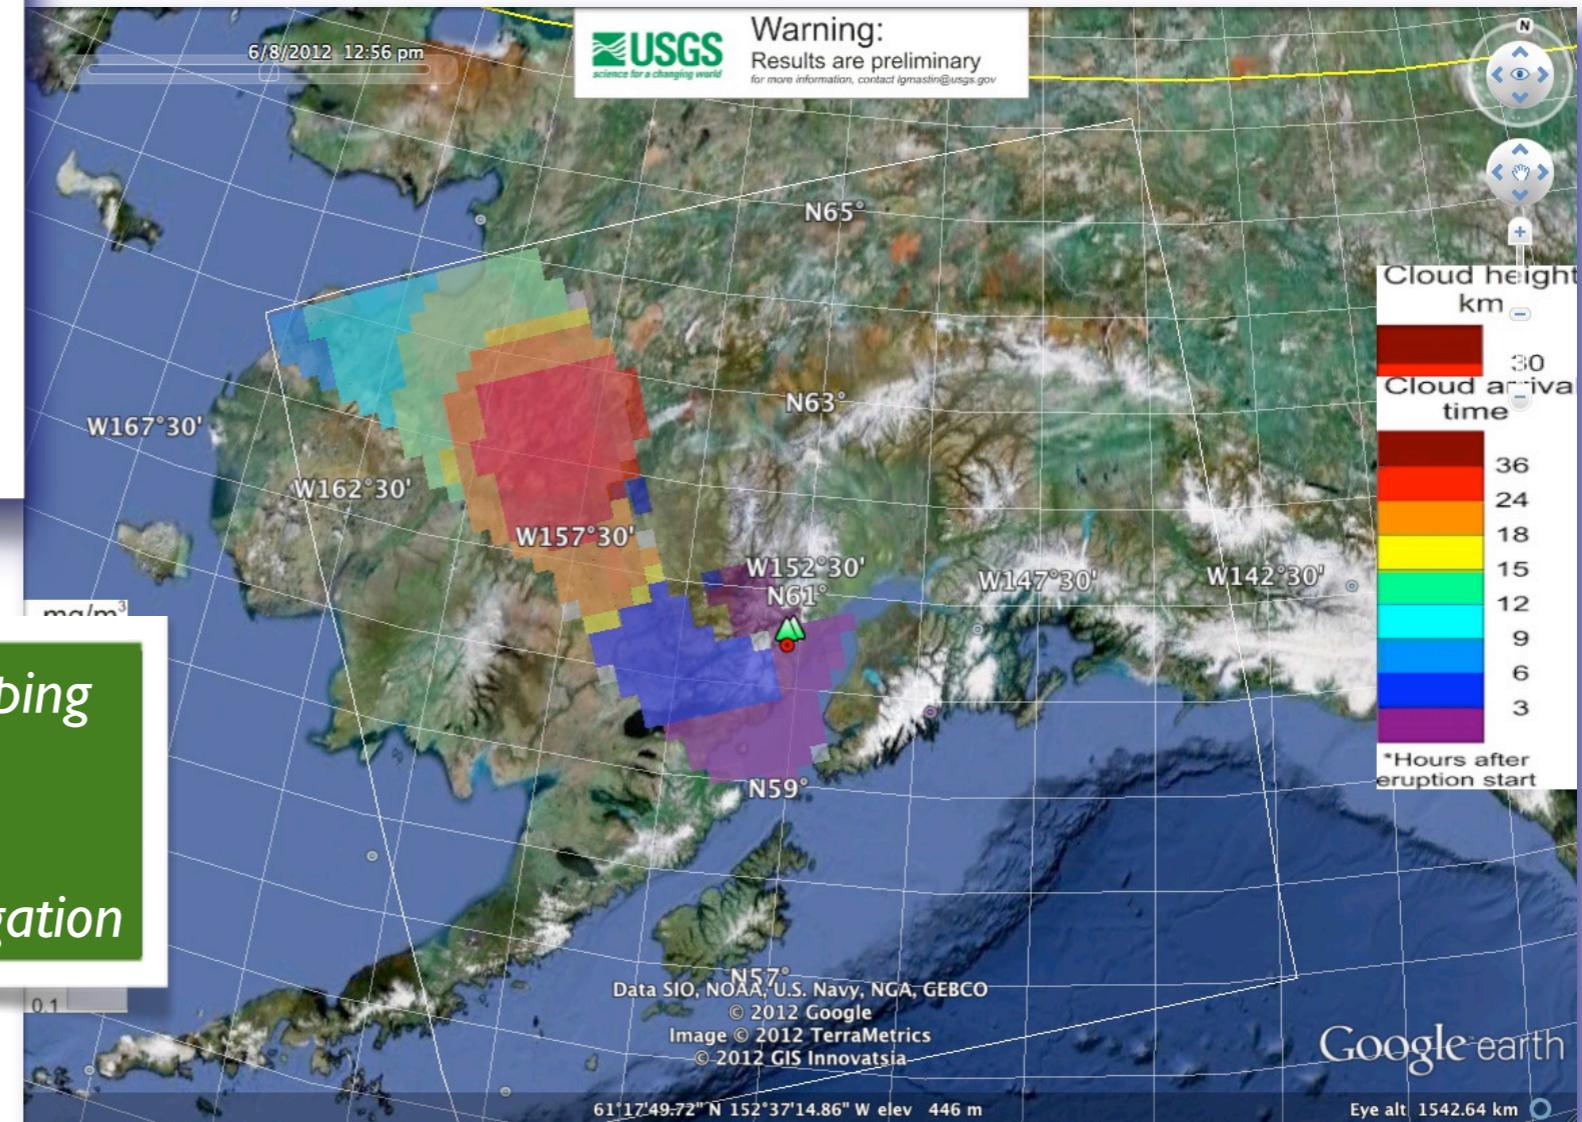
- Include anisotropic refinement for atmospheric applications by putting 3d grids into 2d quadtrees
- Parallel ghost cell exchanges (p4est is already fully parallelized; parallel exchanges between ghost cells need to be worked out)
- General handling of grid orientations in multi-block setting,
- New topologies, i.e. the cubed-sphere (already available in p4est)

see <http://www.forestclaw.org>

# Ash cloud modeling



Ash3d



- Split horizontal, vertical time stepping
- Fully conservative,
- Eulerian, finite volume
- Algorithms based on wave propagation

Ash3d :A finite-volume, conservative numerical model for ash transport and tephra deposition,  
Schwaiger, Denlinger, Mastin, JGR (2012)